

# Kerberos and IPSec

# Kerberos - Scalability

- In an enterprise network with  $N$  users, suppose that we want to authenticate each other, Alice, Bob, Charlie, David, Eva...
- Authentication using public keys (how many key pairs?)
  - $N$  users  $\Rightarrow$   $N$  key pairs (but needs PKI)
- Authentication using symmetric keys (how many keys?)
  - $N$  users requires (on the order of)  $N^2$  keys

**When  $N$  is large, symmetric key case does not scale**

Discussion: what should we do?

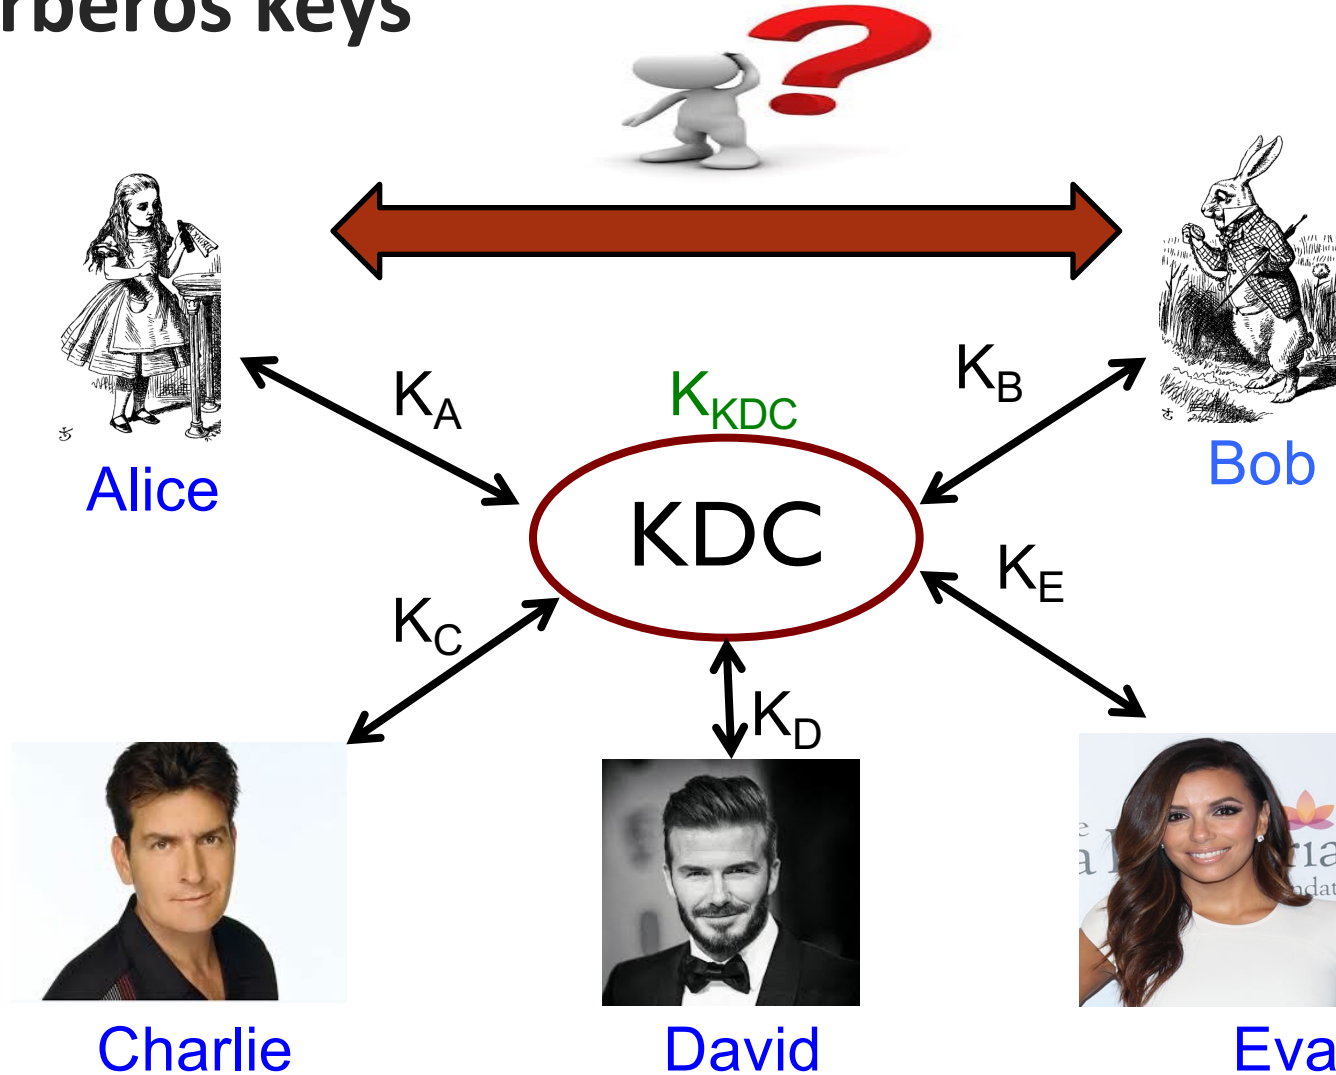
# Kerberos in Computer Security

- In security, Kerberos is an authentication protocol based on symmetric key crypto
  - Designed for LANs or corporate networks
  - Originated at MIT
  - Relies on a Trusted Third Party (TTP)
    - Security depends on TTP
  - Kerberos based on symmetric keys but only requires N keys for N users
  - No PKI is needed

# Kerberos KDC

- Kerberos **Key Distribution Center** or **KDC**
  - KDC acts as the TTP
  - TTP is trusted, so it must not be compromised
- KDC shares symmetric key  $K_A$  with Alice, key  $K_B$  with Bob, key  $K_C$  with Carol, etc.
- **And a master key  $K_{KDC}$  known *only* to KDC**
- KDC enables authentication, session keys
  - Session key for confidentiality and integrity

# Kerberos keys



# Kerberos Tickets

- KDC issue **tickets** containing info needed to access network resources
- KDC also issues **Ticket-Granting Tickets** or **TGTs** that are used to obtain tickets
- Each TGT contains
  - Session key
  - User's ID
  - Expiration time
- Every TGT is encrypted with  $K_{KDC}$ 
  - So, TGT can only be read by the KDC



# Three phases of Kerberos

- Phase I: Kerberized Login
- Phase II: Alice Requests “Ticket to Bob”
- Phase III: Alice Uses Ticket to Bob

# Phase I: Kerberized Login

- Alice enters her password
- Then Alice's computer does following:
  - Derives  $K_A$  from Alice's password
  - Uses  $K_A$  to get TGT for Alice from KDC
- Alice then uses her TGT (credentials) to securely access network resources
- **Plus:** Security is transparent to Alice
- **Minus:** KDC *must* be secure — it's trusted!




# Kerberized Login



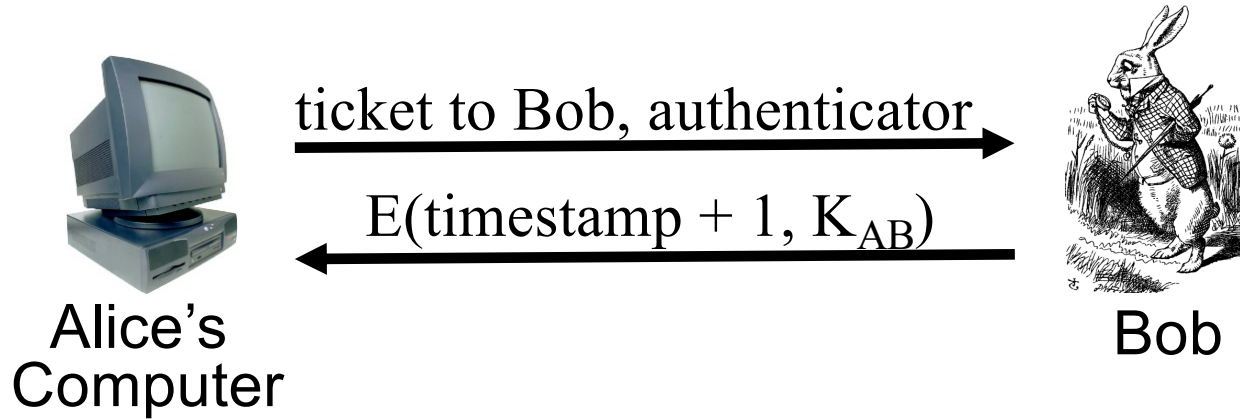
- Key  $K_A = h(\text{Alice's password})$
- KDC creates session key  $S_A$
- Alice's computer decrypts  $S_A$  and TGT
  - Then it forgets  $K_A$
- $TGT = E(\text{"Alice"}, S_A, K_{KDC})$

# Phase II: Alice Requests “Ticket to Bob”



- $\text{REQUEST} = (\text{TGT}, \text{authenticator})$ 
  - $\text{authenticator} = E(\text{timestamp}, S_A)$
- KDC gets  $S_A$  from TGT to verify timestamp (why time?)
  - $\text{TGT} = E(\text{"Alice"}, S_A, K_{\text{KDC}})$
- $\text{REPLY} = E(\text{"Bob"}, K_{AB}, \text{ticket to Bob}, S_A)$ 
  - $\text{ticket to Bob} = E(\text{"Alice"}, K_{AB}, K_B)$

# Phase III: Alice Uses Ticket to Bob



- ticket to Bob =  $E(\text{"Alice"}, K_{AB}, K_B)$
- authenticator =  $E(\text{timestamp}, K_{AB})$
- Bob decrypts "ticket to Bob" to get  $K_{AB}$  which he then uses to verify timestamp

# Kerberos Question 1

- When Alice logs in, KDC sends  $E(S_A, TGT, K_A)$  where  $TGT = E(\text{"Alice"}, S_A, K_{KDC})$


**Q:** Why is TGT encrypted with  $K_A$ ?

**A:** Extra work for no added security!



# Kerberos Question 2

- In Alice's "Kerberized" login to Bob, can Alice remain anonymous?
- Alice remain anonymous in REQUEST
- $\text{REQUEST} = (\text{TGT}, \text{authenticator})$ 
  - $\text{authenticator} = E(\text{timestamp}, S_A)$



# Kerberos Question 3

- Why is “ticket to Bob” sent to Alice?
  - Why doesn't KDC send it directly to Bob?
    - Bob needs to remember  $K_{AB}$  until Alice initiate communication, make it **stateful** which against the design of Kerberos
- $REPLY = E(\text{“Bob”}, K_{AB}, \text{ticket to Bob}, S_A)$ 
  - ticket to Bob =  $E(\text{“Alice”}, K_{AB}, K_B)$




Alice

Talk to Bob



Computer




KDC

# Kerberos Question 4

- Why not have KDC remember session key instead of putting it in a TGT? If so, there is no need for TGT!
  - $TGT = E(\text{“Alice”}, S_A, K_{KDC})$
- **Answer: Stateless KDC is a major feature of Kerberos**

# IPSec

- IPSec lives at the network layer
- IPSec is transparent to applications





# Overview

- IPsec comprises **a suite of protocols** to ensure the **integrity**, **confidentiality**, and **authentication** of data communications over an IP network
- Designed with the goal to achieve security for all IP-related protocols
- But predominantly used in VPNs at this moment

# IKE and ESP/AH

- Two parts of IPSec: IKE and ESP/AH
- **IKE:** Internet Key Exchange
  - Mutual authentication
  - Establish session key
  - Two “phases” — like SSL session/connection
- **ESP/AH**
  - **ESP:** Encapsulating Security Payload — for encryption and/or integrity of IP packets
  - **AH:** Authentication Header — integrity only

# IKE

- IKE has 2 phases
  - Phase 1 — IKE security association (IKE SA)
  - Phase 2 — IPSec security association (IPSec SA)
- Phase 1 is comparable to SSL *session*
- Phase 2 is comparable to SSL *connection*
- Not an obvious need for two phases in IKE
- If multiple Phase 2's do not occur, then it is **more** costly to have two phases!

# IKE Phase 1

- Four different “key” options
  - Public key encryption (original version)
  - Public key encryption (improved version)
  - Public key signature
  - Symmetric key
- For each of these, two different “modes”
  - Main mode and aggressive mode
- **There are 8 versions of IKE Phase 1!**
- Over-engineered!

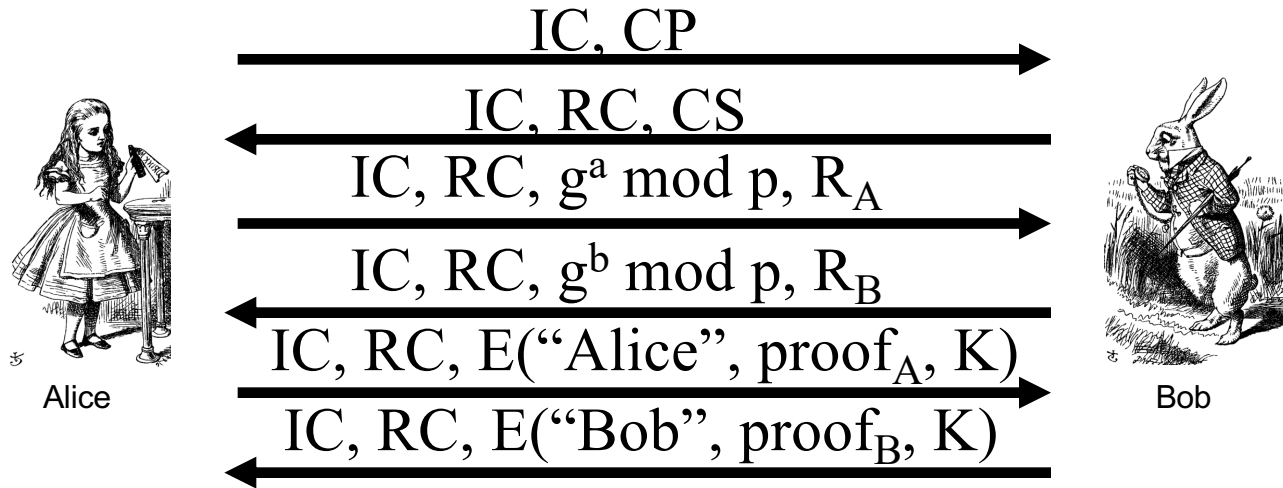
# IKE Phase 1

- We discuss 6 of 8 Phase 1 variants
  - Public key signatures (main & aggressive modes)
  - Symmetric key (main and aggressive modes)
  - Public key encryption (main and aggressive)
- Why public key encryption and public key signatures?
  - Always know your own private key
  - **May not** (initially) know other side's public key

# IKE Phase 1

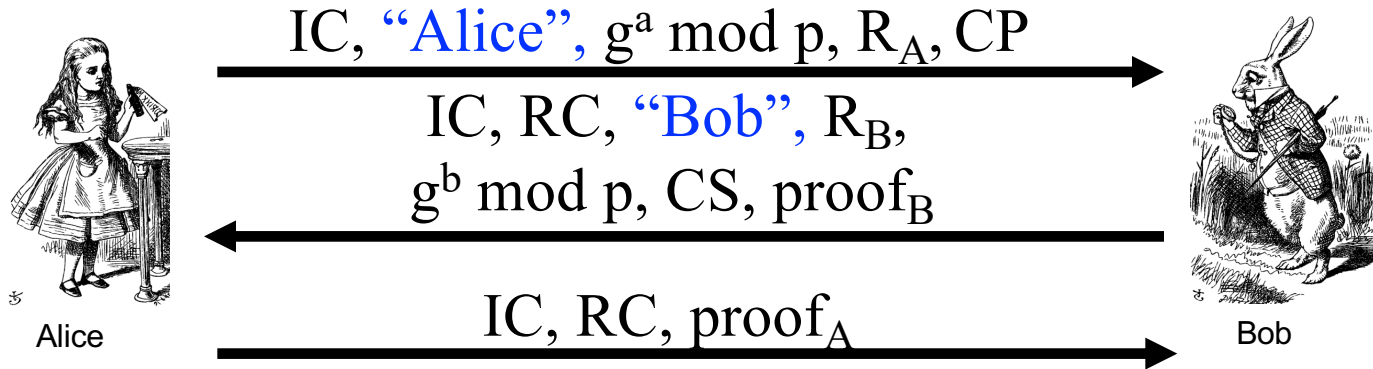
- Uses Diffie-Hellman to establish session key
- Let  $a$  be Alice's Diffie-Hellman exponent
- Let  $b$  be Bob's Diffie-Hellman exponent
- Let  $g$  be generator and  $p$  prime
- Recall that  $p$  and  $g$  are public

# IKE Phase 1: Digital Signature (Main Mode)



- CP = crypto proposed, CS = crypto selected
- IC = initiator “cookie”, RC = responder “cookie”
- $K = h(IC, RC, g^{ab} \bmod p, R_A, R_B)$
- $SKEYID = h(R_A, R_B, g^{ab} \bmod p)$
- $\text{proof}_A = [h(SKEYID, g^a \bmod p, g^b \bmod p, IC, RC, CP, \text{"Alice"})]_{\text{Alice}}$

# IKE Phase 1: Digital Signature (Aggressive Mode)



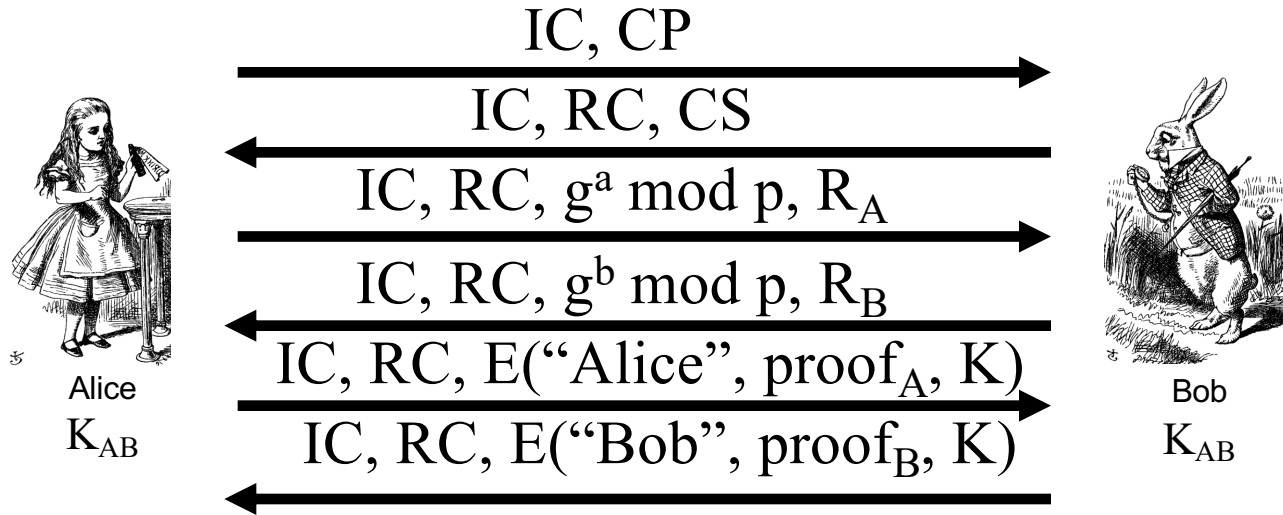
- Main difference from main mode
  - Not trying to protect identities
  - Cannot negotiate  $g$  or  $p$



# Main vs Aggressive Modes


- Main mode **MUST** be implemented
- Aggressive mode **SHOULD** be implemented
  - So, if aggressive mode is not implemented, “you should feel guilty about it”
- Might create interoperability issues

# IKE Phase 1: Symmetric Key (Main Mode)



- Same as signature mode except
  - $K_{AB}$  = symmetric key shared in advance
  - $K = h(IC, RC, g^{ab} \bmod p, R_A, R_B, K_{AB})$
  - $SKEYID = h(K, g^{ab} \bmod p)$
  - $\text{proof}_A = h(SKEYID, g^a \bmod p, g^b \bmod p, IC, RC, CP, \text{"Alice"})$

# Problems with Symmetric Key (Main Mode)



- Catch-22(dilemma)
  - Alice sends her ID in message 5
  - Alice's ID encrypted with  $K$
  - To find  $K$  Bob must know  $K_{AB}$ :  $K = h(IC, RC, g^{ab} \bmod p, R_A, R_B, K_{AB})$
  - To get  $K_{AB}$  Bob must know he's talking to Alice!
- Result: **Alice's ID must be IP address!**
- Useless mode
  - Alice must use a static IP
  - Failed to hide identity

# Quiz

- **What is SSH primarily used for?**
  - A) Transferring files between machines
  - **B) Securing remote computer connections**
  - C) Encrypting email communications
  - D) Browsing the web anonymously

# Quiz

- **What is the purpose of the SSH key pair in SSH connections?**
  - A) To encrypt the connection
  - **B) To verify the identity of the client to the server**
  - C) To increase the connection speed
  - D) To monitor the data transfer

# Quiz

- **What is a primary reason for a server not to authenticate clients in SSL/TLS communications?**
  - A) To simplify the server configuration
  - B) To reduce the computational load on the server
  - C) To increase the security of the server
  - **D) Both A and B are correct**

# Quiz

- **What is the primary purpose of SSL?**
  - A) Encrypting data transfers between a client and a server
  - B) Speeding up website performance
  - C) Providing stronger passwords
  - D) Filtering spam emails

# Quiz

- **SSL certificates are issued by entities known as:**
  - A) Internet service providers
  - **B) Certificate authorities**
  - C) Domain registrars
  - D) Web hosts



# Quiz

- Which protocol has largely replaced SSL for security purposes?
  - A) HTTPS
  - B) SFTP
  - C) TLS
  - D) SSH

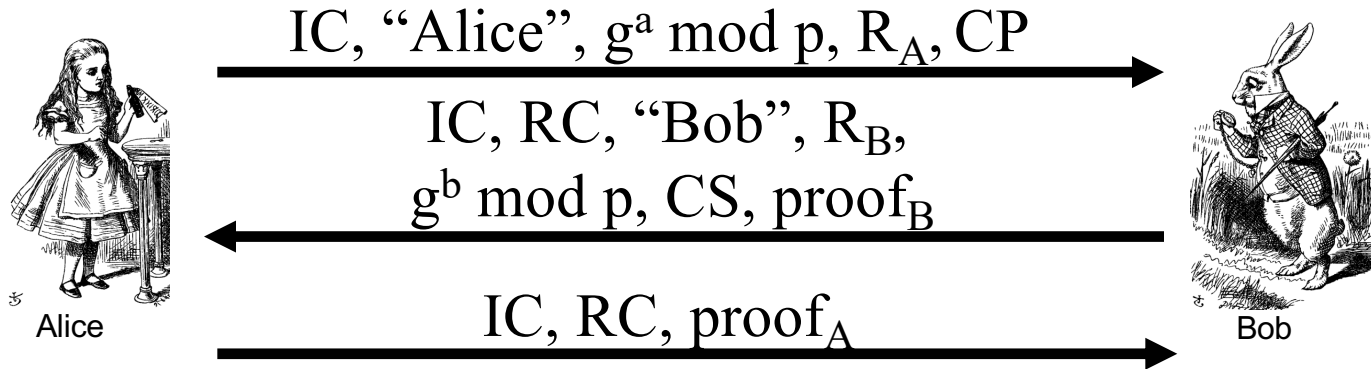
# Quiz

- Which of the following measures can mitigate the impact of the Heartbleed?
  - A) Changing user passwords
  - B) Updating to a patched version of OpenSSL
  - C) Installing antivirus software
  - D) Reducing the number of network connections

# Quiz

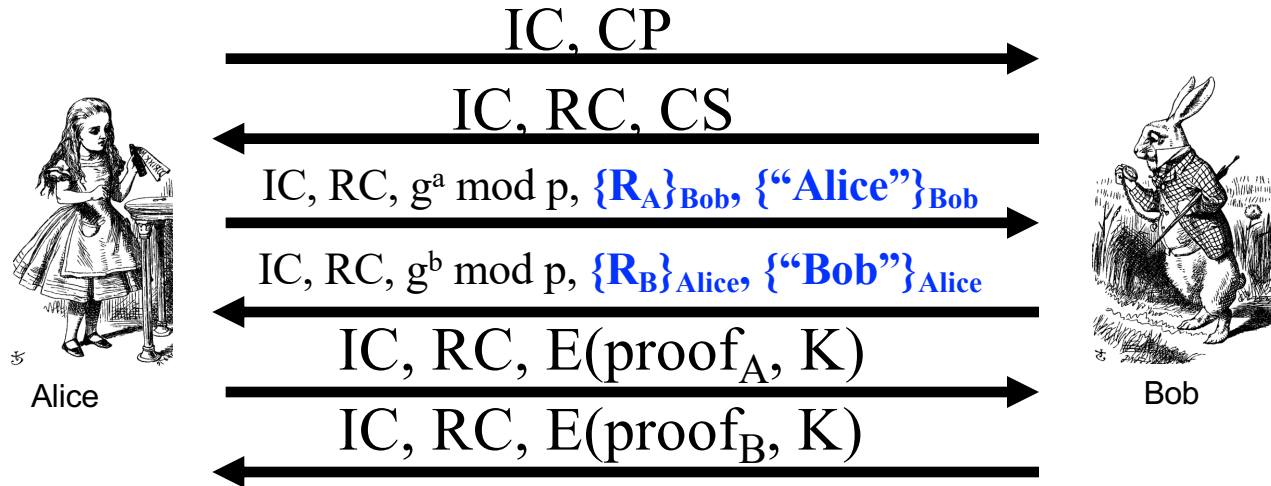
- **How does the Heartbleed bug expose sensitive data?**
  - A) By intercepting data during transmission
  - B) By allowing unauthorized access to databases
  - **C) By causing buffer over-reads in memory**
  - D) By corrupting data encryption

# IKE Phase 1: Symmetric Key (Aggressive Mode)



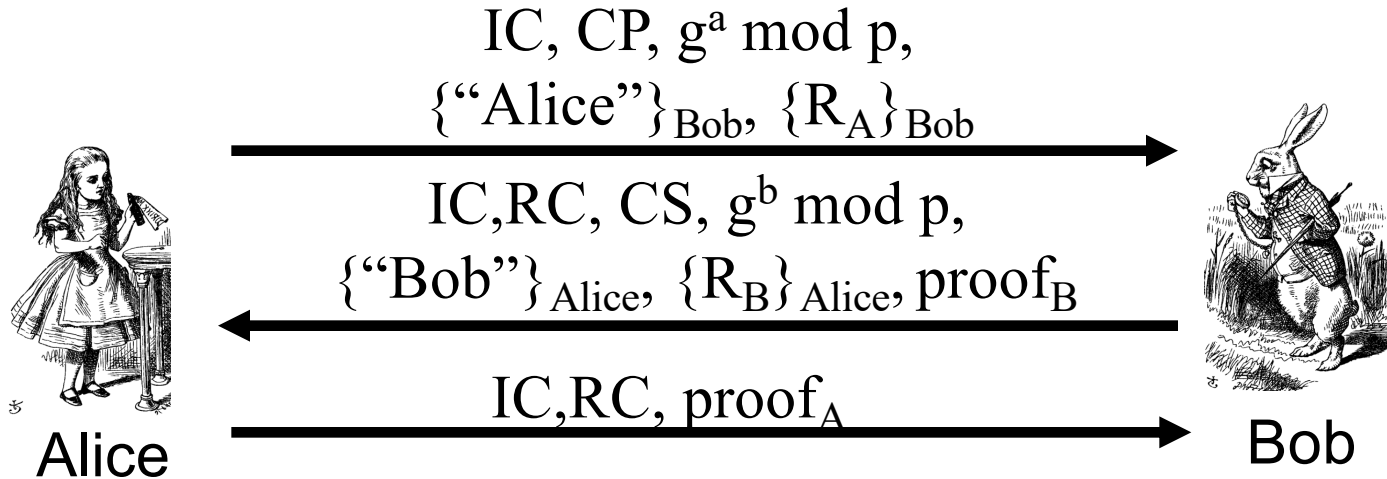
- Same format as digital signature aggressive mode
- Not trying to hide identities...
- As a result, does **not** have problems of main mode
- But does not (pretend to) hide identities

# IKE Phase 1: Public Key Encryption (Main Mode)



- CP = crypto proposed, CS = crypto selected
- IC = initiator “cookie”, RC = responder “cookie”
- $K = h(IC, RC, g^{ab} \bmod p, R_A, R_B)$
- $SKEYID = h(R_A, R_B, g^{ab} \bmod p)$
- $\text{proof}_A = h(SKEYID, g^a \bmod p, g^b \bmod p, IC, RC, CP, \text{"Alice"})$

# IKE Phase 1: Public Key Encryption (Aggressive Mode)

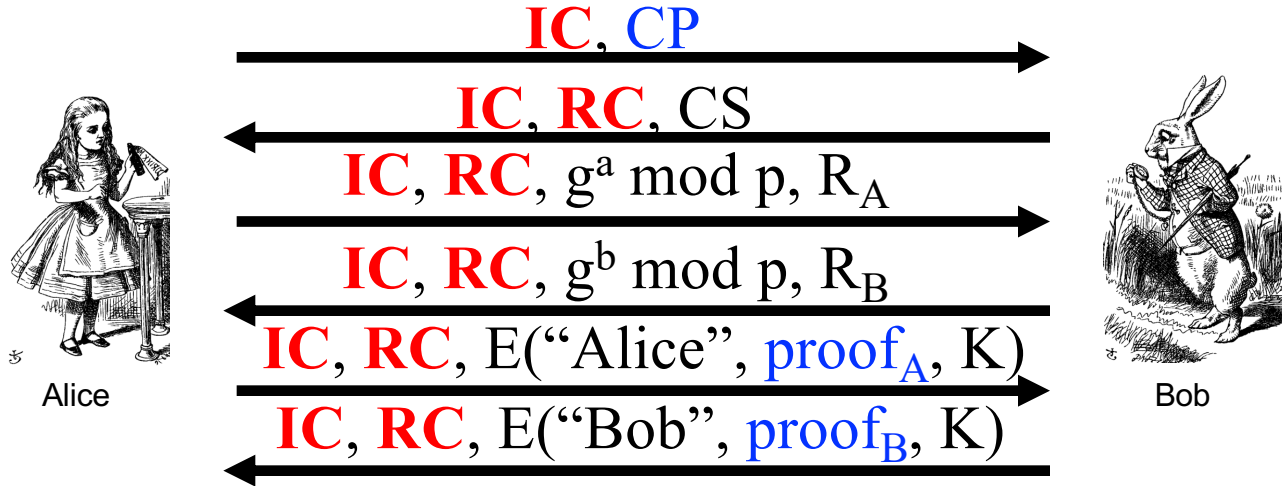


- $K, proof_A, proof_B$  computed as in main mode
- Note that **identities are hidden**
  - The **only aggressive mode to hide identities**
  - So, why have a main mode?
  - Negotiate  $g$  and  $p$  in main mode

# IKE Phase 1 Cookies

- IC and RC — cookies (or “anti-clogging tokens”) supposed to prevent DoS attacks
  - No relation to Web cookies
- To reduce DoS threats, Bob wants to remain **stateless** as long as possible
- But Bob must remember CP from message 1 (required for proof of identity in message 6)
- Bob must keep state from 1st message on
  - So, these “cookies” offer little DoS protection

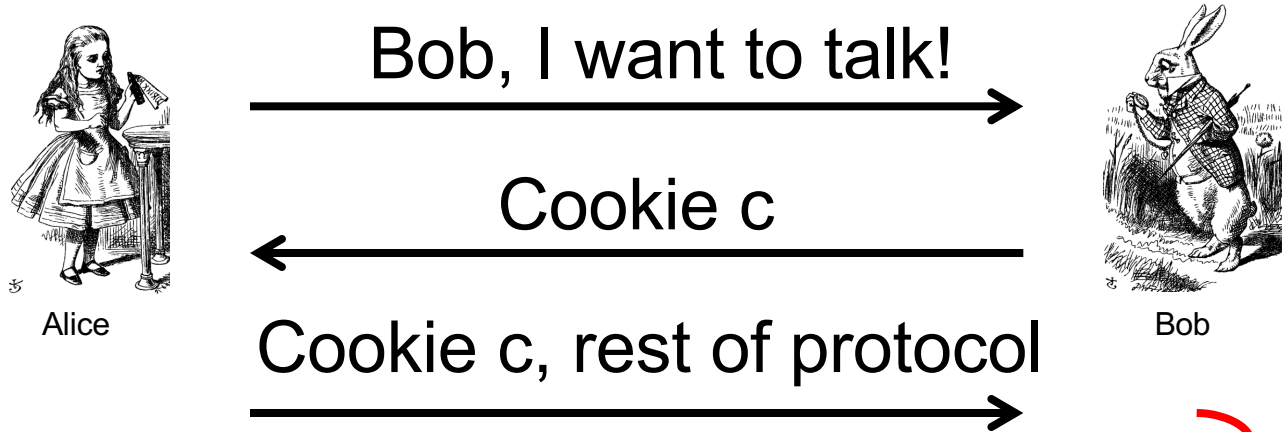
# IPSec cookies



- CP = crypto proposed, CS = crypto selected
- **IC = initiator “cookie”, RC = responder “cookie”**
- $K = h(\text{IC}, \text{RC}, g^{ab} \bmod p, R_A, R_B)$
- $\text{SKEYID} = h(R_A, R_B, g^{ab} \bmod p)$
- $\text{proof}_A = [h(\text{SKEYID}, g^a \bmod p, g^b \bmod p, \text{IC}, \text{RC}, \text{CP}, \text{"Alice"})]_{\text{Alice}}$



# Stateless Cookie Protocol



Does  $c = \text{hash}(\text{IP address, local secret})$ ?  
If so, continue protocol

$\text{Cookie } c = \text{hash}(\text{IP address, local secret})$

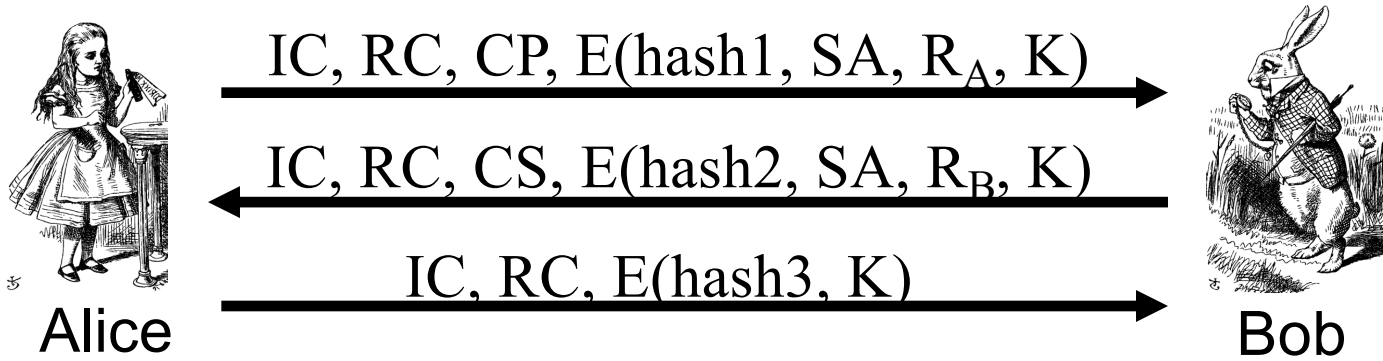
# IKE Phase 1 Summary

- Result of IKE phase 1 is
  - Mutual authentication
  - Shared symmetric key
  - **IKE Security Association (IKE SA)**
- But phase 1 is expensive
  - Especially in public key and/or main mode
- Developers of IKE thought it would be used for lots of things — not just IPsec
  - Partly explains the over-engineering...

# IKE Phase 2

- Phase 1 establishes **IKE SA**
- Phase 2 establishes **IPSec SA**
- Comparison to SSL
  - SSL session is comparable to IKE Phase 1
  - SSL connections are like IKE Phase 2
- IKE **could** be used for lots of things...
- ...but in practice, it's **not!**

# IKE Phase 2



- Key  $K$ ,  $IC$ ,  $RC$  and  $SA$  known from Phase 1 ( $SA$ : id),  $R_A$  &  $R_B$  new
- Proposal **CP** includes **ESP** and/or **AH**
- Hashes 1,2,3 depend on  $SKEYID$ ,  $SA$ ,  $R_A$  and  $R_B$
- Keys derived from  $KEYMAT = h(SKEYID, R_A, R_B, \text{junk})$
- Recall  $SKEYID$  depends on phase 1 key method

# IPSec

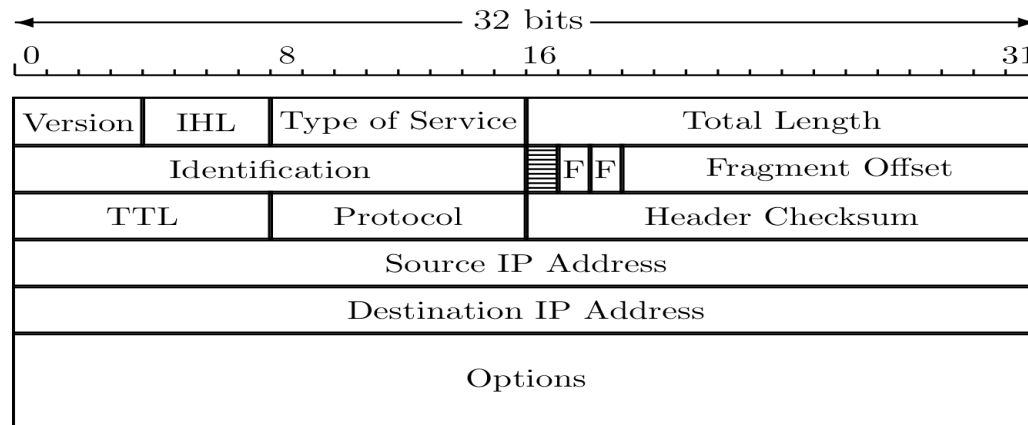
- After IKE Phase 1, we have an IKE SA
- After IKE Phase 2, we have an IPSec SA
- Both sides have a shared symmetric key
- Now what?
  - We want to protect **IP datagrams**
- But what is an IP datagram?
  - Considered from the perspective of IPSec...

# IP Review

- IP datagram is of the form




- Where IP header is



# IP and TCP


- Consider Web traffic
  - IP encapsulates TCP and...
  - ...TCP encapsulates HTTP



- IP **data** includes TCP header, etc.

# IPSec Transport Mode

- IPSec **Transport Mode**




- Transport mode designed for *host-to-host*
- Transport mode is efficient
  - Adds minimal amount of extra header
- The original header remains
  - Passive attacker can observe




# IPSec: Host-to-Host

- IPSec transport mode



# IPSec Tunnel Mode

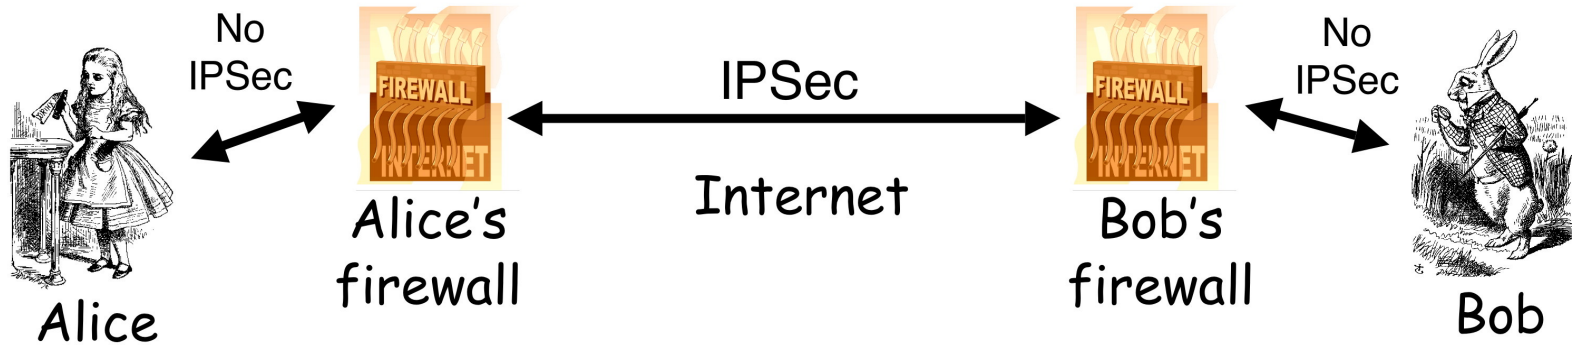
- IPSec Tunnel Mode



- Tunnel mode for *firewall-to-firewall* traffic
- Original IP packet encapsulated in IPSec
- Original IP header not visible to attacker
  - New IP header from firewall to firewall
  - Attacker does not know which hosts are talking

# IPSec: Firewall-to-Firewall


- **IPSec tunnel mode**



- Local networks not protected

# Comparison of IPSec Modes

- Transport Mode



- Tunnel Mode



- Transport Mode
  - Host-to-host
- Tunnel Mode
  - Firewall-to-firewall
- Transport Mode not necessary...
- ...but it's more efficient

# IPSec Security

- What kind of protection?
  - Confidentiality?
  - Integrity?
  - Both?
- What to protect?
  - Data?
  - Header?
  - Both?
- ESP/AH do some combinations of these

# AH vs ESP

- AH — Authentication Header
  - **Integrity only** (no confidentiality, no encryption)
  - Integrity-protect everything beyond IP header and some fields of header
- ESP — Encapsulating Security Payload
  - **Integrity and confidentiality** both **required**
  - Protects everything beyond IP header(data in IP)
  - Integrity-only by using NULL encryption

# ESP's NULL Encryption

- According to RFC 2410
  - NULL encryption “is a block cipher the origins of which appear to be lost in antiquity”
  - “Despite rumors”, there is no evidence that NSA “suppressed publication of this algorithm”
  - Evidence suggests it was developed in Roman times as exportable version of Caesar’s cipher
  - Can make use of keys of varying length
  - No IV is required
  - **$\text{Null}(P,K) = P$  for any  $P$  and any key  $K$**
- Bottom line: **Security people can be strange**

# Why Does AH Exist?

- Cannot encrypt IP header
  - Routers must look at the IP header
  - IP addresses, TTL, etc.
  - IP header exists to route packets!
- AH protects **immutable fields** in IP header
  - Cannot protect all header fields
  - TTL, for example, will change
- Why not use ESP with NULL encryption?
  - ESP provides no protection to the header



# IPSec and Complexity

- IPSec is a complex protocol
- **Over-engineered**
  - Lots of (generally useless) features
- Flawed
  - Some significant security issues
- Interoperability is serious challenge
  - Defeats the purpose of having a standard!